

Chapter 7 :: Microarchitecture

Digital Design and Computer Architecture

David Money Harris and Sarah L. Harris

Introduction

- Microarchitecture: how to implement an architecture in hardware
- Processor:
 - Datapath: functional blocks
 - Control: control signals

Application Software	programs
Operating Systems	device drivers
Architecture	instructions registers
Micro-architecture	datapaths controllers
Logic	adders memories
Digital Circuits	AND gates NOT gates
Analog Circuits	amplifiers filters
Devices	transistors diodes
Physics	electrons

Microarchitecture

- Multiple implementations for a single architecture:
 - Single-cycle
 - Each instruction executes in a single cycle
 - Multicycle
 - Each instruction is broken up into a series of shorter steps
 - Pipelined
 - Each instruction is broken up into a series of steps
 - Multiple instructions execute at once.

Processor Performance

- Program execution time

Execution Time = (# instructions)(cycles/instruction)(seconds/cycle)

- Definitions:
 - Cycles/instruction = CPI
 - Seconds/cycle = clock period
 - $1/\text{CPI} = \text{Instructions/cycle} = \text{IPC}$
- Challenge is to satisfy constraints of:
 - Cost
 - Power
 - Performance

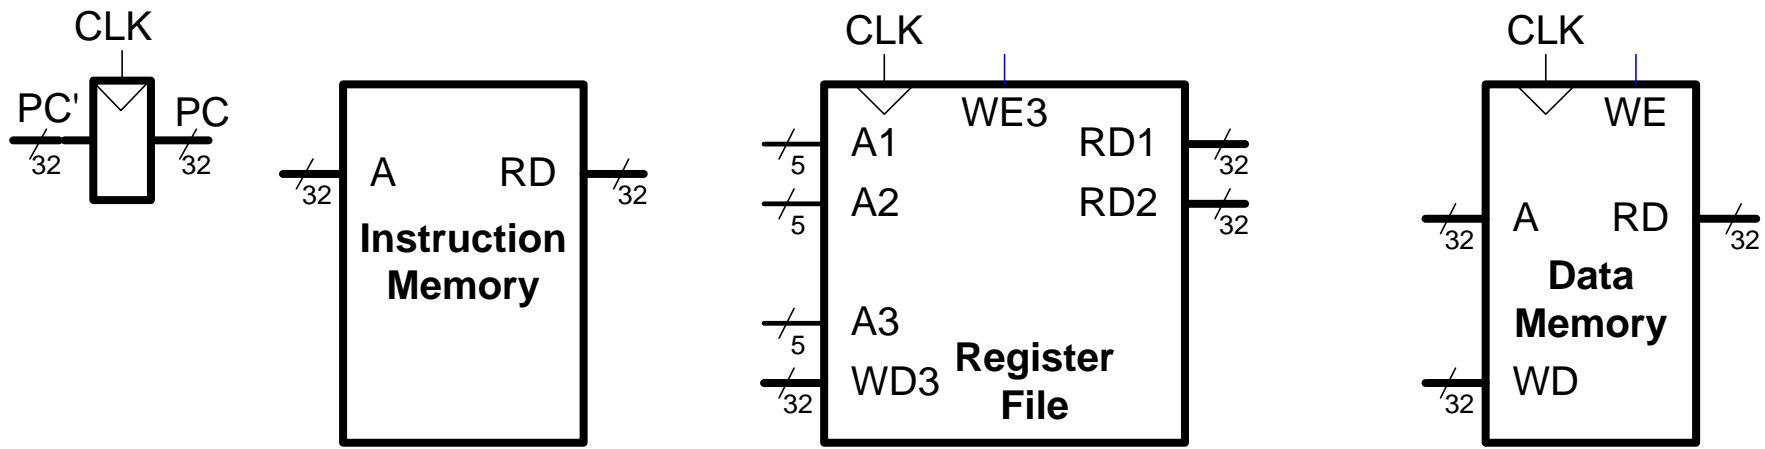
MIPS Processor

- We consider a subset of MIPS instructions:
 - R-type instructions: and, or, add, sub, slt
 - Memory instructions: lw, sw
 - Branch instructions: beq
- Later consider adding addi and j

Architectural State


- Determines everything about a processor:
 - PC
 - 32 registers
 - Memory

MIPS State Elements




Data Memory

```
library IEEE;
use IEEE.STD_LOGIC_1164.all; use STD.TEXTIO.all;
use IEEE.STD_LOGIC_UNSIGNED.all; use IEEE.STD_LOGIC_ARITH.all;
entity dmem is -- data memory
port (clk, we: in STD_LOGIC;
a, wd: in STD_LOGIC_VECTOR (31 downto 0);
rd: out STD_LOGIC_VECTOR (31 downto 0));
end;
architecture behave of dmem is
begin
process is
type ramtype is array (63 downto 0) of STD_LOGIC_VECTOR (31 downto 0);
variable mem: ramtype;
begin
-- read or write memory
loop
if clk'event and clk = '1' then
if (we = '1') then mem (CONV_INTEGER (a(7 downto 2))) := wd;
end if;
end if;
rd <= mem (CONV_INTEGER (a (7 downto 2)));
wait on clk, a;
end loop;
end process;
end;
```



RESETTABLE FLIP-FLOP

```
library IEEE; use IEEE.STD_LOGIC_1164.all; use IEEE.STD_LOGIC_ARITH.all;
entity flopr is -- flip-flop with synchronous reset
generic (width: integer);
port (  clk, reset: in STD_LOGIC;
       d:          in STD_LOGIC_VECTOR(width-1 downto 0);
       q:          out STD_LOGIC_VECTOR(width-1 downto 0));
end;
architecture asynchronous of flopr is
begin
process (clk, reset) begin
if reset = '1' then q <= CONV_STD_LOGIC_VECTOR(0, width);
elsif clk'event and clk = '1' then
    q <= d;
end if;
end process;
end;
```

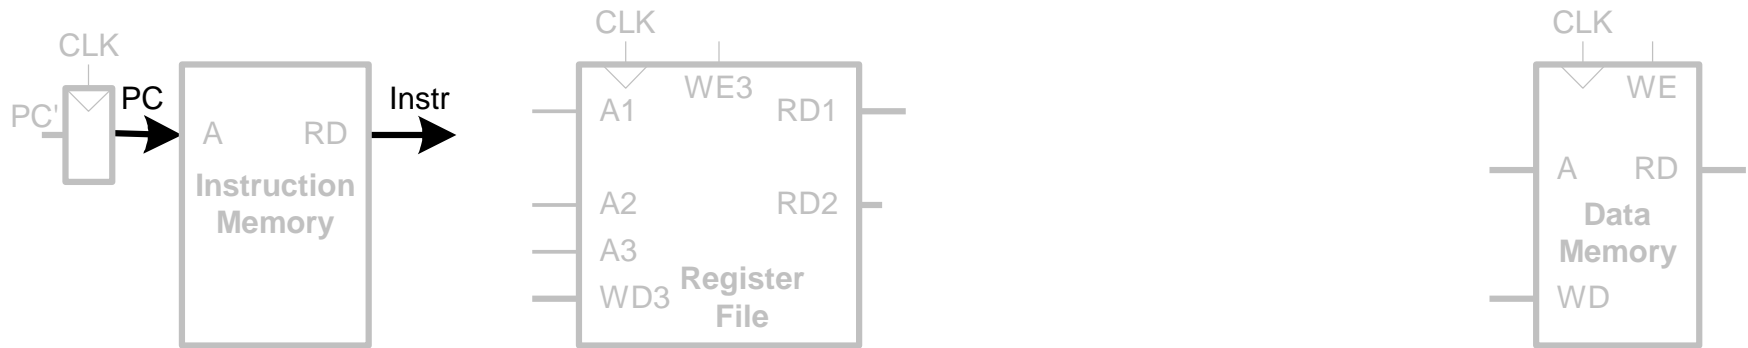


Single-Cycle MIPS Processor

- Datapath
- Control

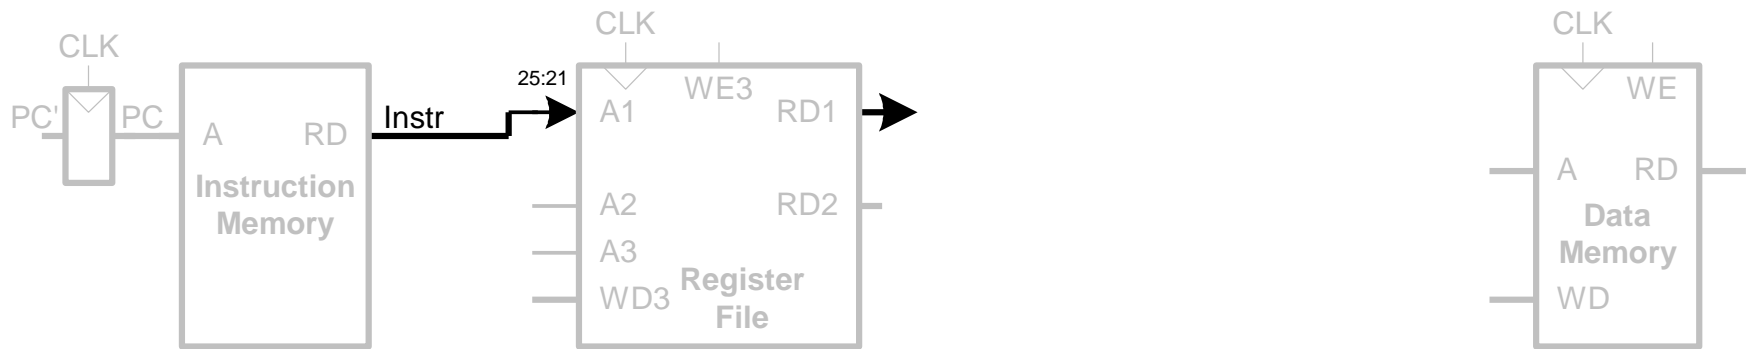
Single-Cycle Datapath: lw fetch

- First consider executing lw
- **STEP 1:** Fetch instruction



Single-Cycle Datapath: l_w register read

- **STEP 2:** Read source operands from register file

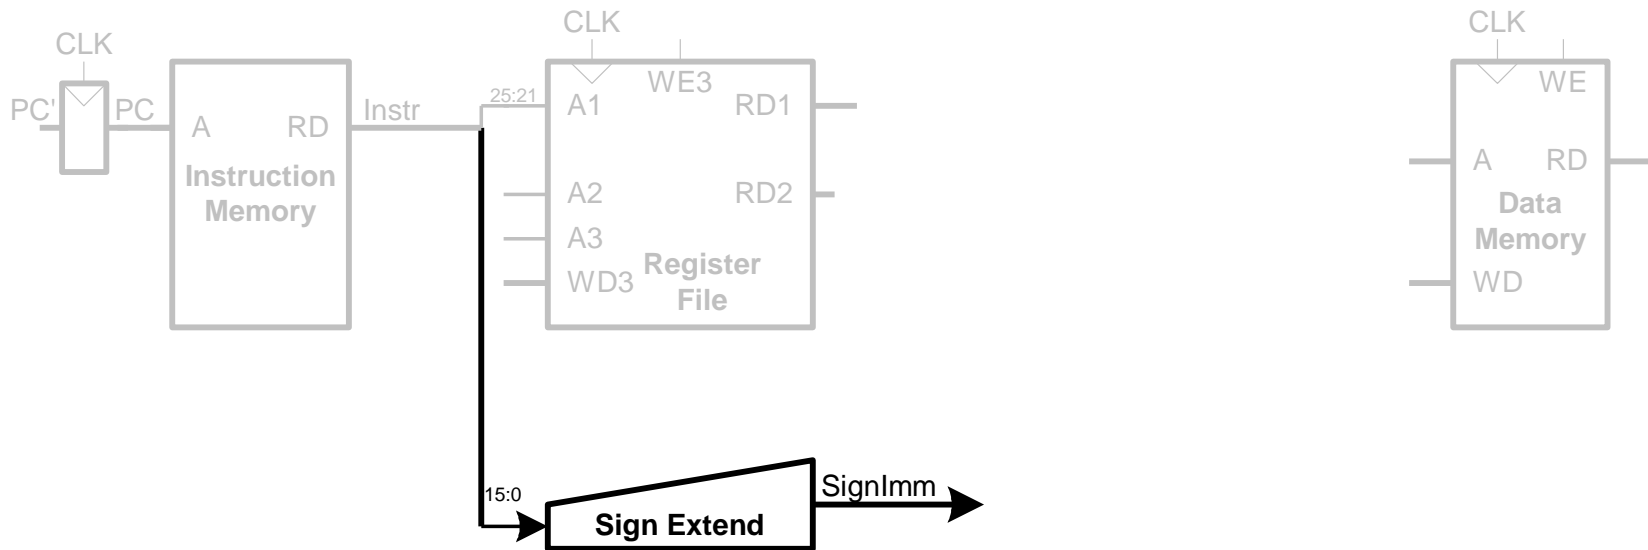


I-Type



Single-Cycle Datapath: l_w immediate

- **STEP 3:** Sign-extend the immediate

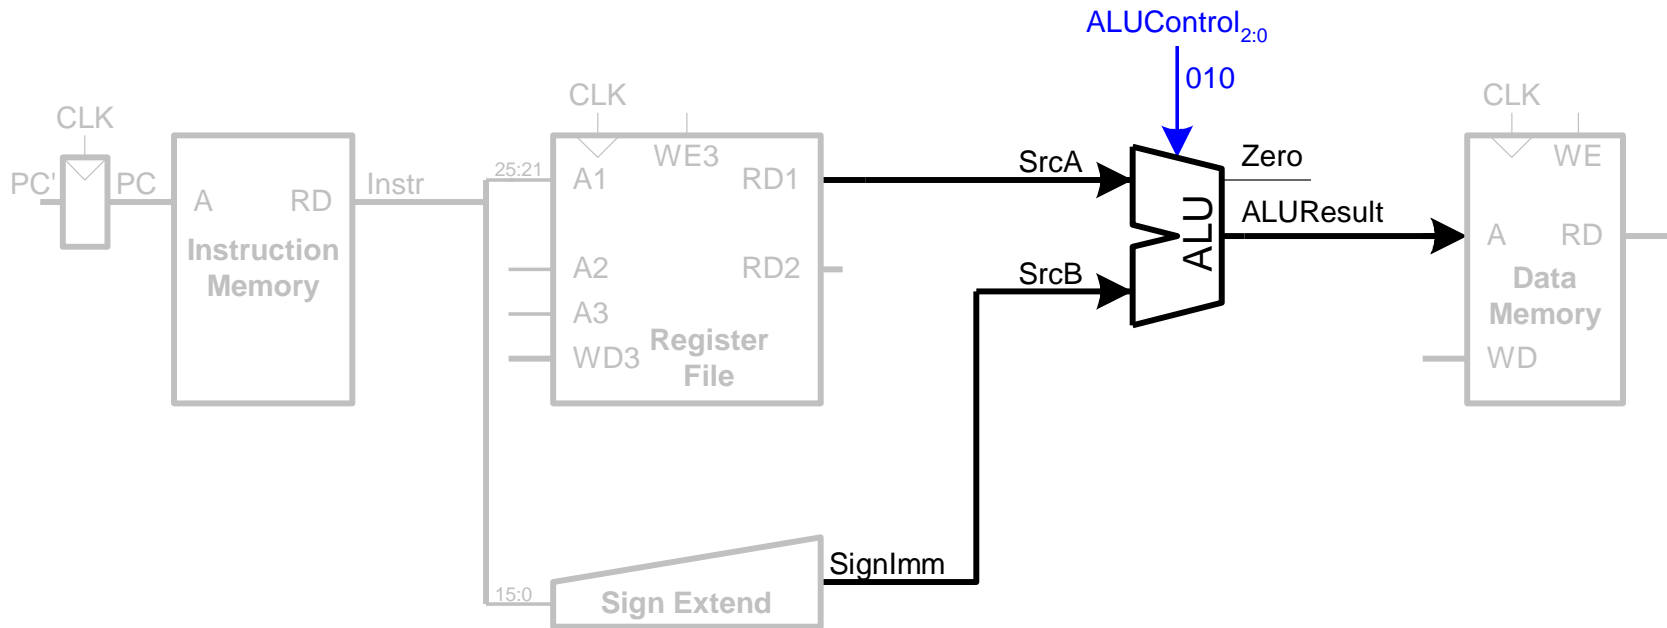


I-Type



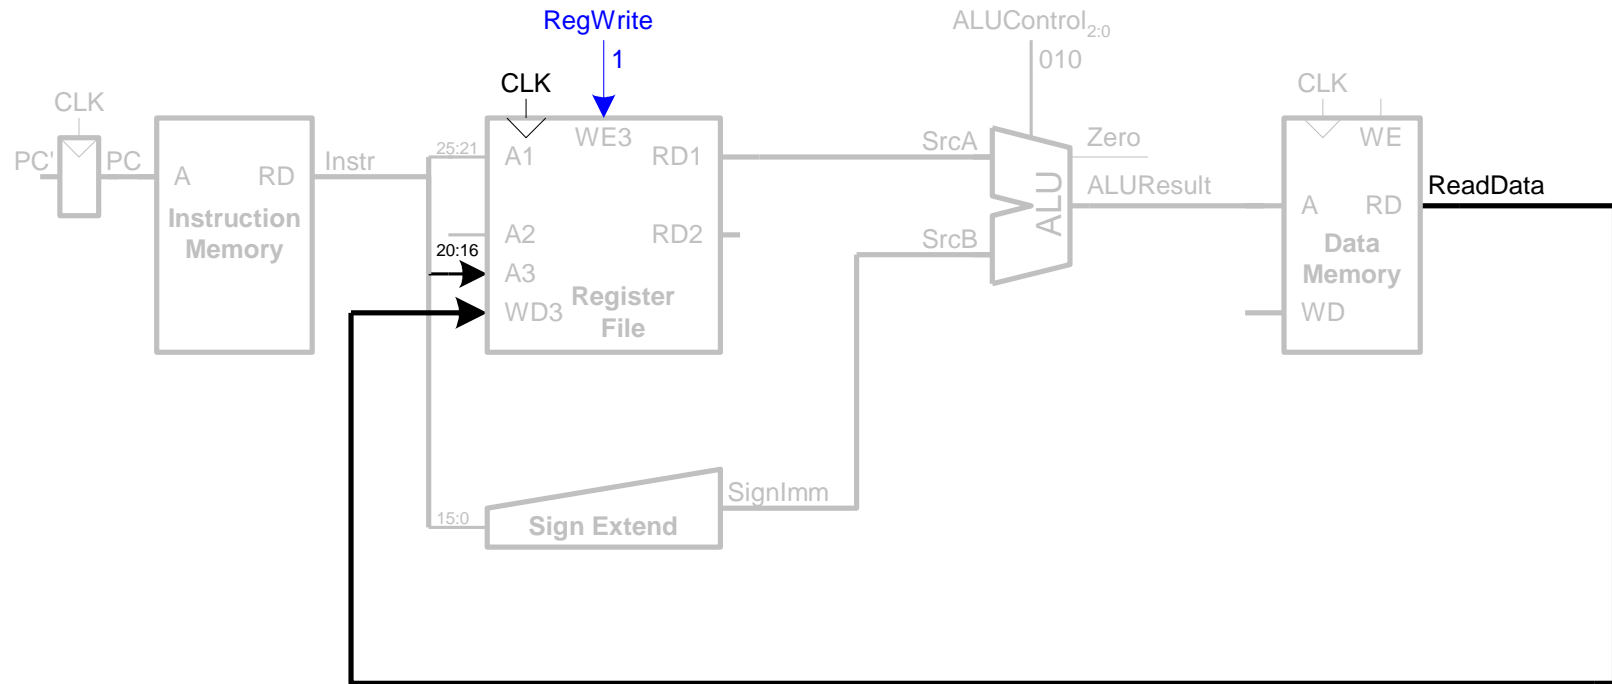
Single-Cycle Datapath: l_w address

- STEP 4:** Compute the memory address



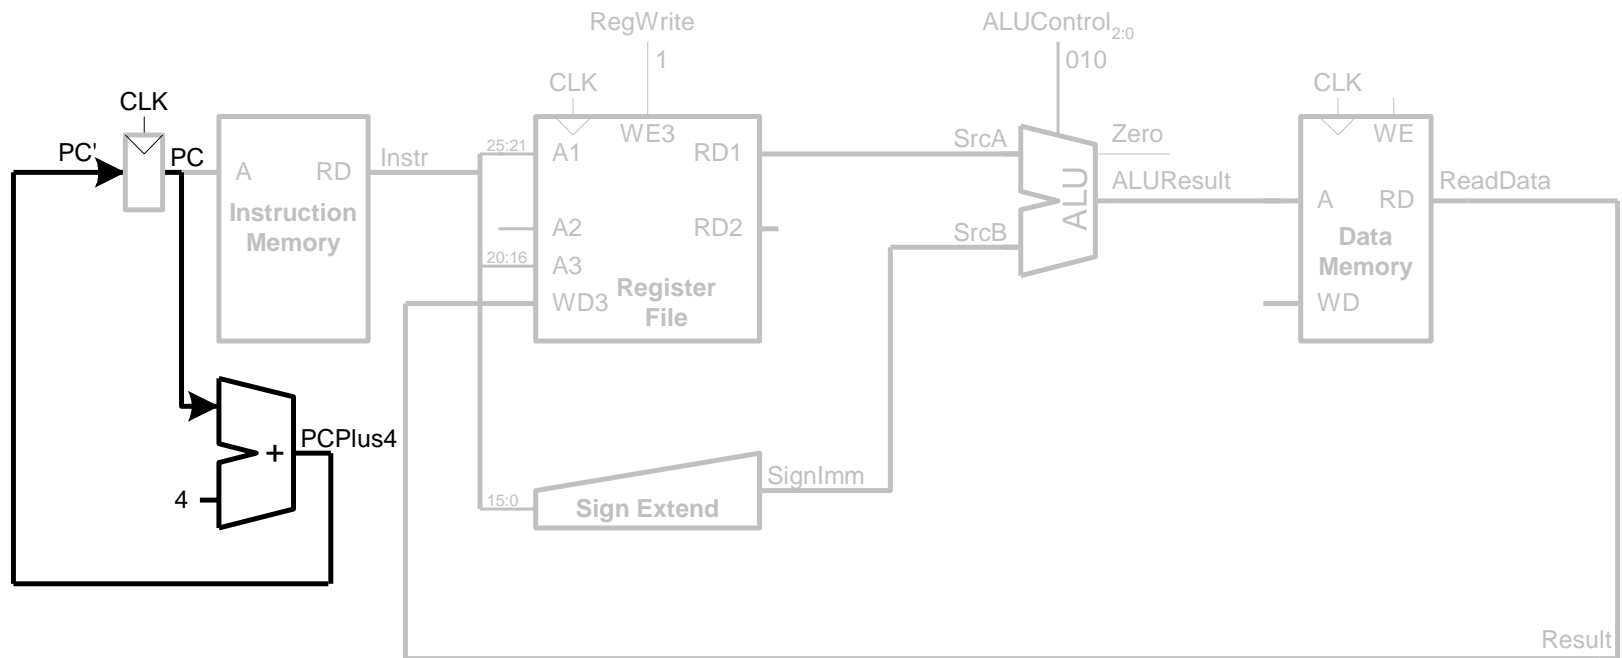
Single-Cycle Datapath: l_w memory read

- STEP 5:** Read data from memory and write it back to register file



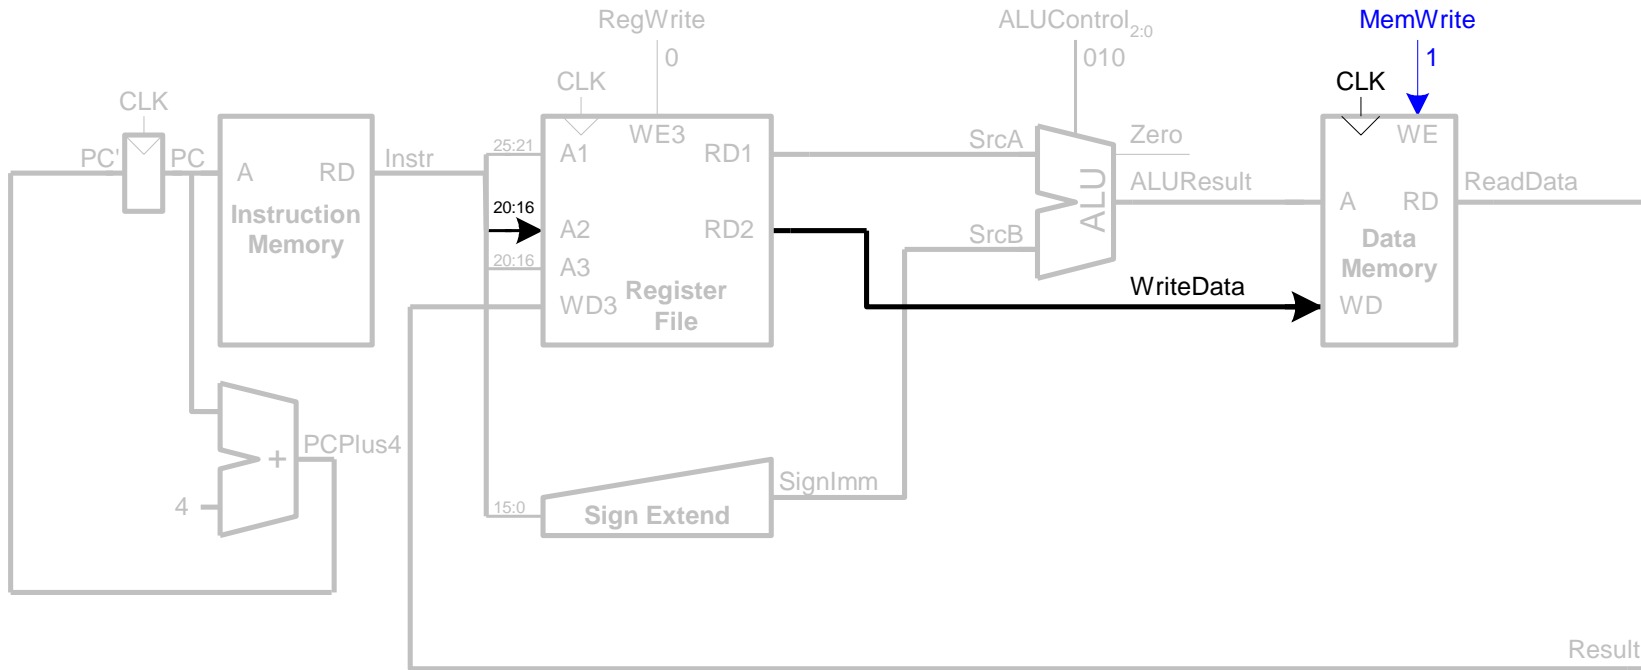
Single-Cycle Datapath: 1_W PC increment

- STEP 6:** Determine the address of the next instruction



Single-Cycle Datapath: SW

- Write data in rt to memory



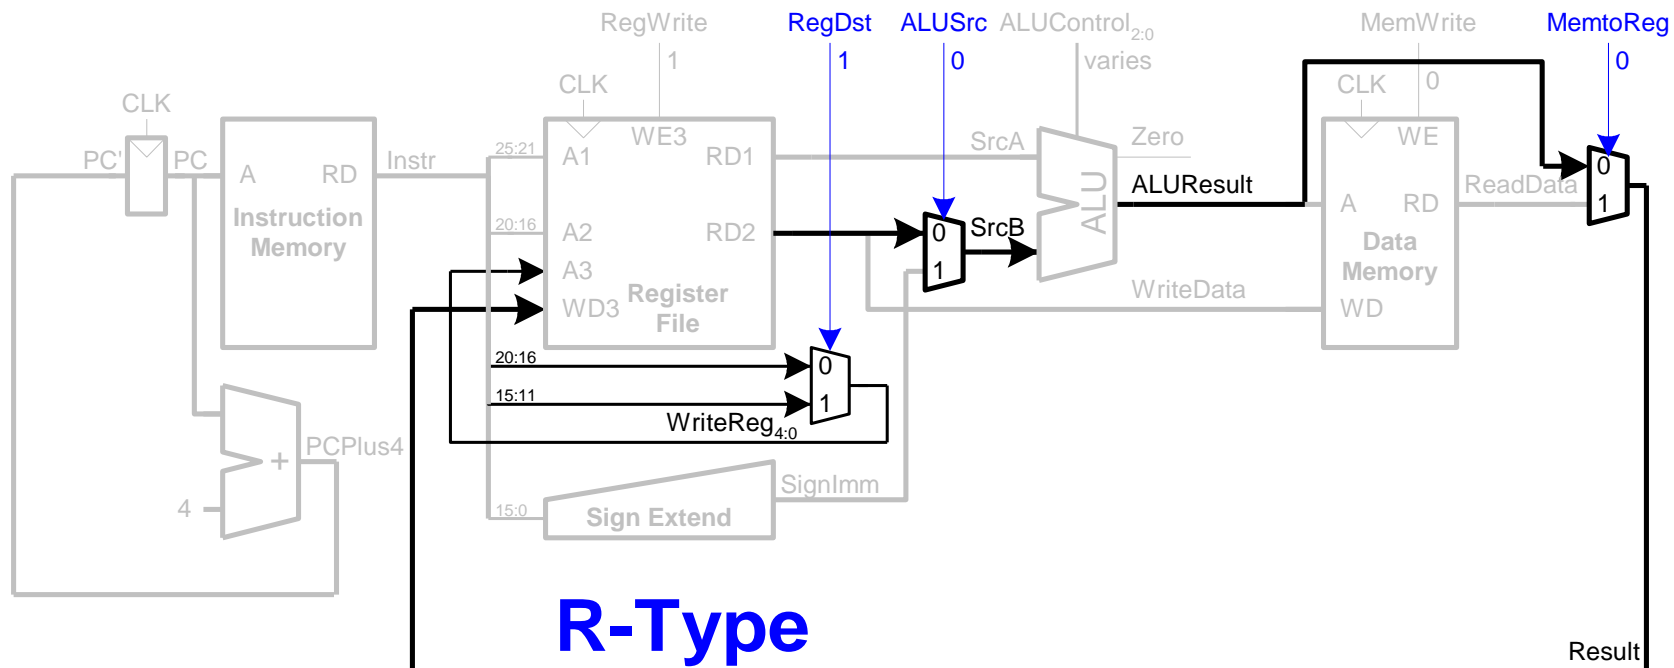
I-Type

op	rs	rt	imm
6 bits	5 bits	5 bits	16 bits

7-<17>

Single-Cycle Datapath: R-type instructions

- Read from *rs* and *rt*
- Write *ALUResult* to register file
- Write to *rd* (instead of *rt*)

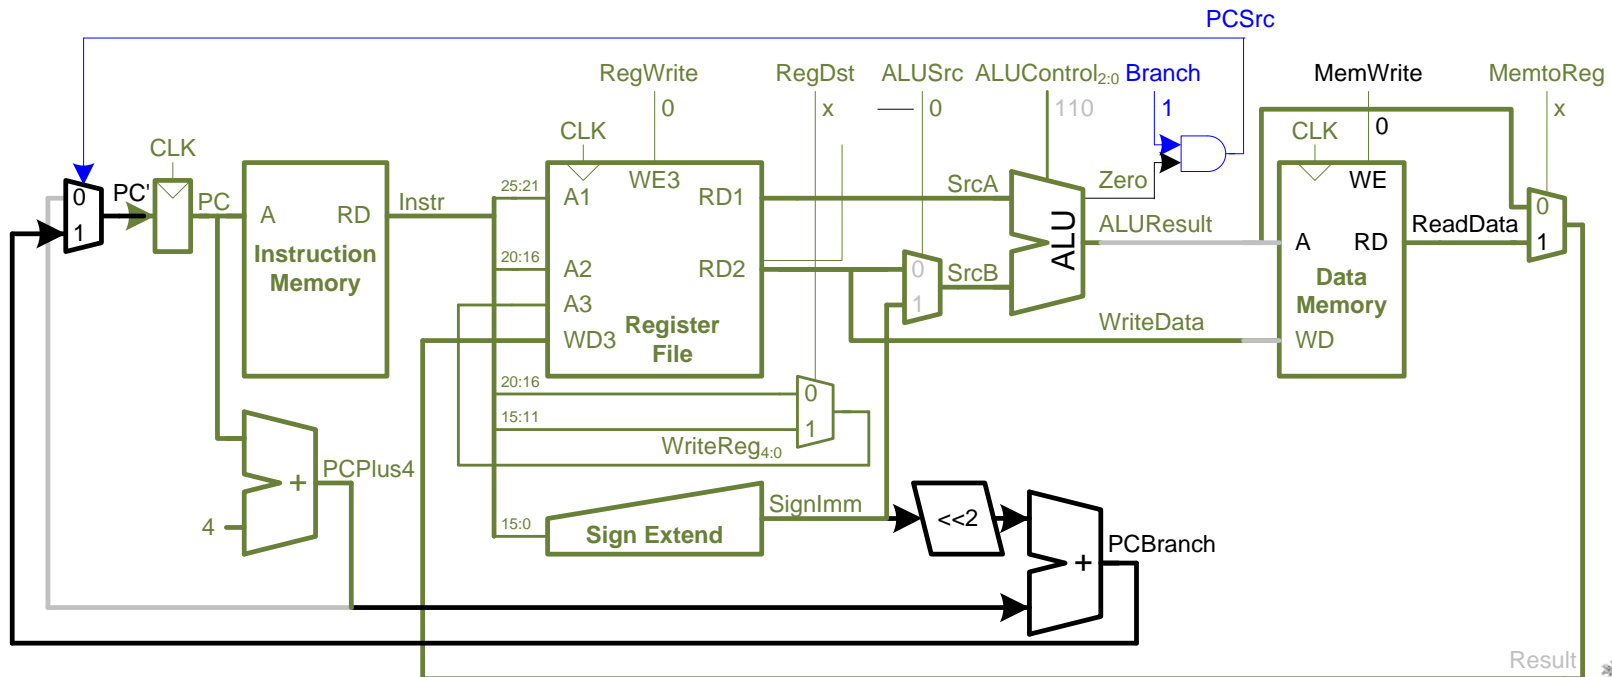


op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

Single-Cycle Datapath: beq

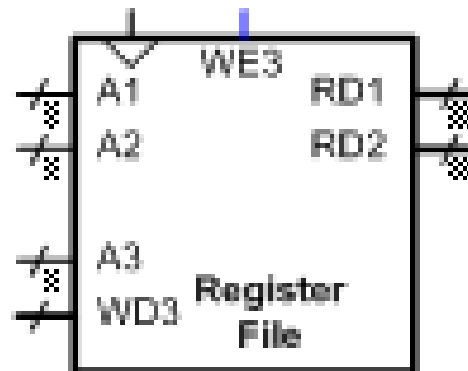
- Determine whether values in `rs` and `rt` are equal
- Calculate branch target address:

$$\text{BTA} = (\text{sign-extended immediate} \ll 2) + (\text{PC} + 4)$$



Register File in VHDL (1)

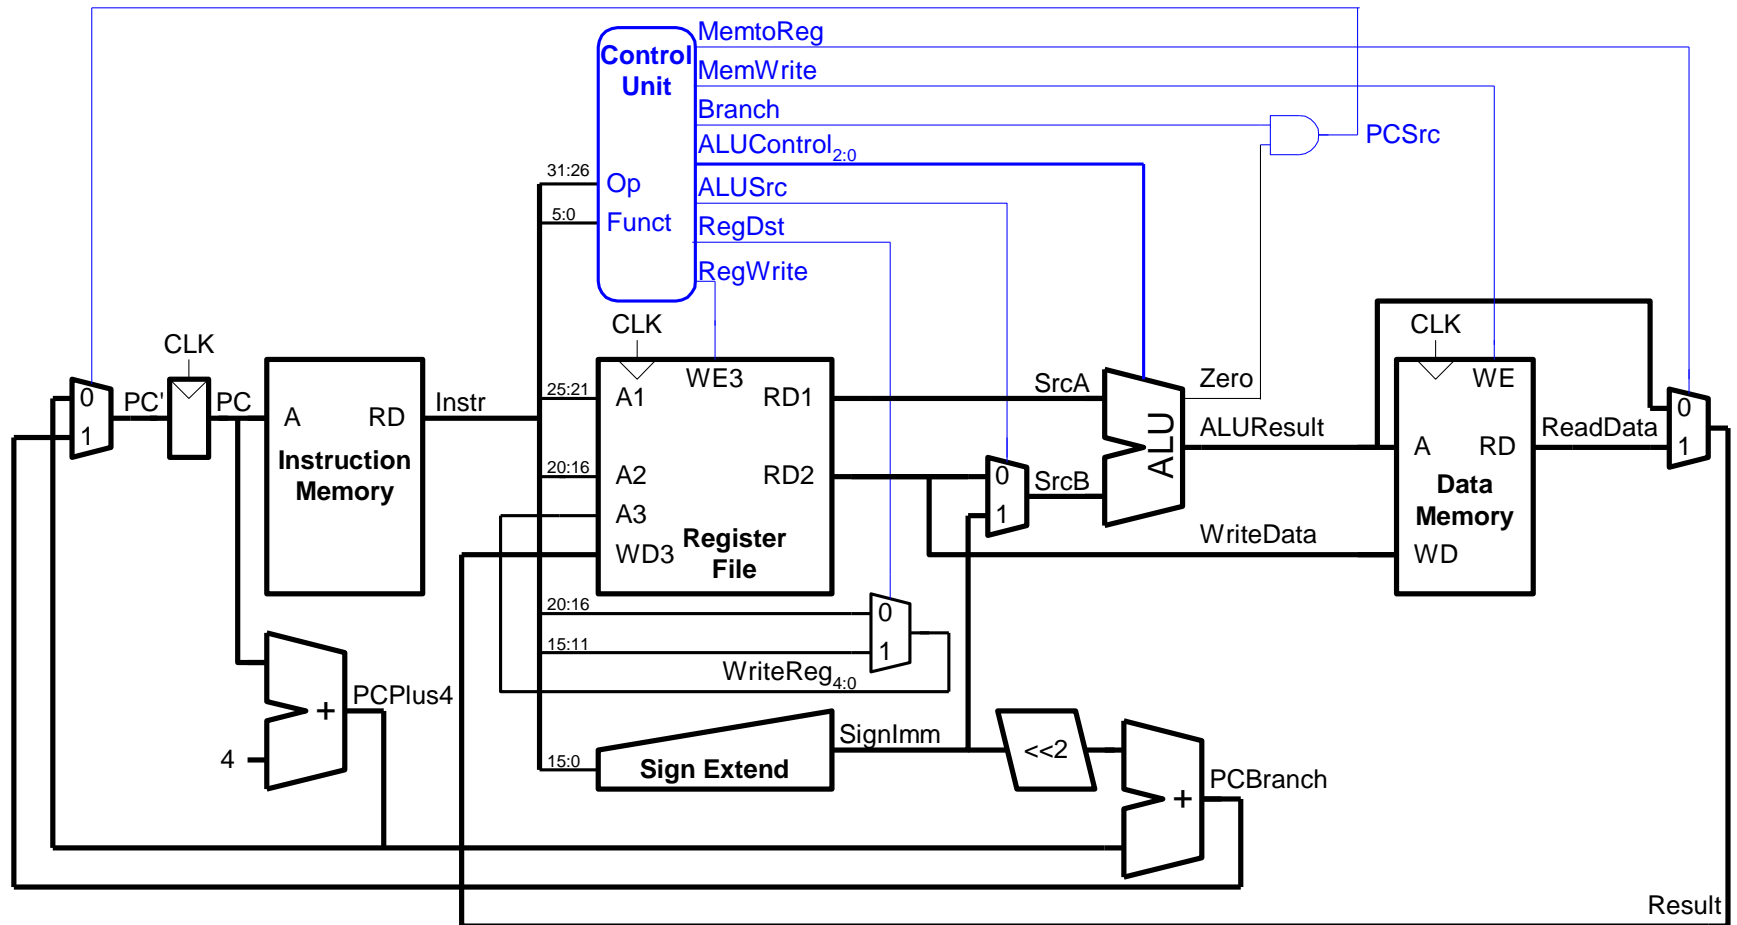
```
library IEEE; use IEEE.STD_LOGIC_1164.all;
use IEEE.STD_LOGIC_UNSIGNED.all;
entity regfile is -- three-port register file
port(  clk: in STD_LOGIC;
      we3: in STD_LOGIC;
      ra1, ra2, wa3: in STD_LOGIC_VECTOR(4 downto 0);
      wd3: in STD_LOGIC_VECTOR(31 downto 0);
      rd1, rd2: out STD_LOGIC_VECTOR(31 downto 0));
end;
architecture behave of regfile is
type ramtype is array (31 downto 0) of STD_LOGIC_VECTOR (31 downto 0);
signal mem: ramtype;
begin
```




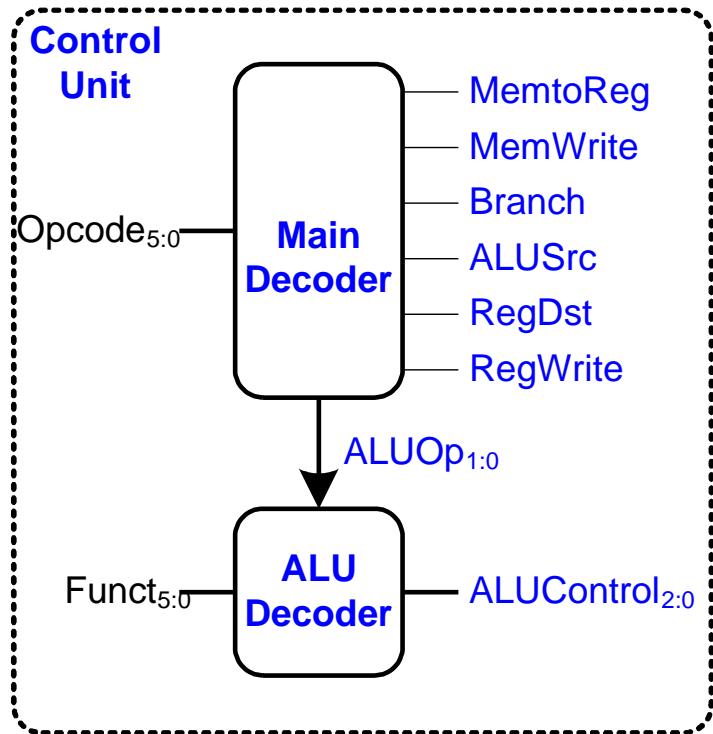
Register File in VHDL (2)

```
begin
-- three-ported register file
-- read two ports combinatorially
-- write third port on rising edge of clock
process(clk) begin
    if clk'event and clk = '1' then
        if we3 = '1' then mem(CONV_INTEGER(wa3)) <= wd3;
        end if;
    end if;
end process;
process (ra1, ra2) begin
    if (conv_integer (ra1) = 0) then rd1 <= X"00000000";
-- register 0 holds 0
    else rd1 <= mem(CONV_INTEGER (ra1));
    end if;
    if (conv_integer(ra2) = 0) then rd2 <= X"00000000";
    else rd2 <= mem(CONV_INTEGER(ra2));
    end if;
end process;
end;
```

Complete Single-Cycle Processor



Control Unit

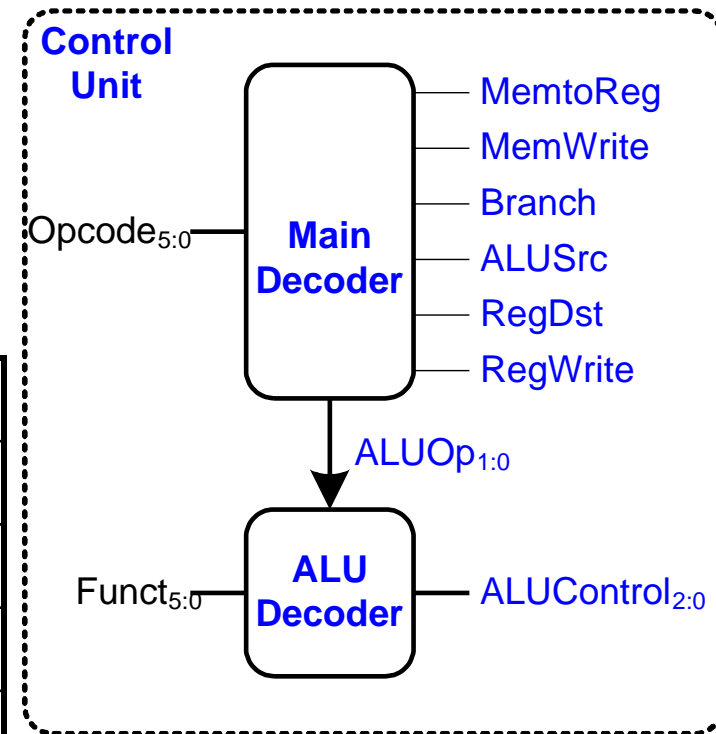


F_{2:0}	Function
000	A & B
001	A B
010	A + B
011	not used
100	A & ~B
101	A ~B
110	A - B
111	SLT

Control Unit: ALU Decoder

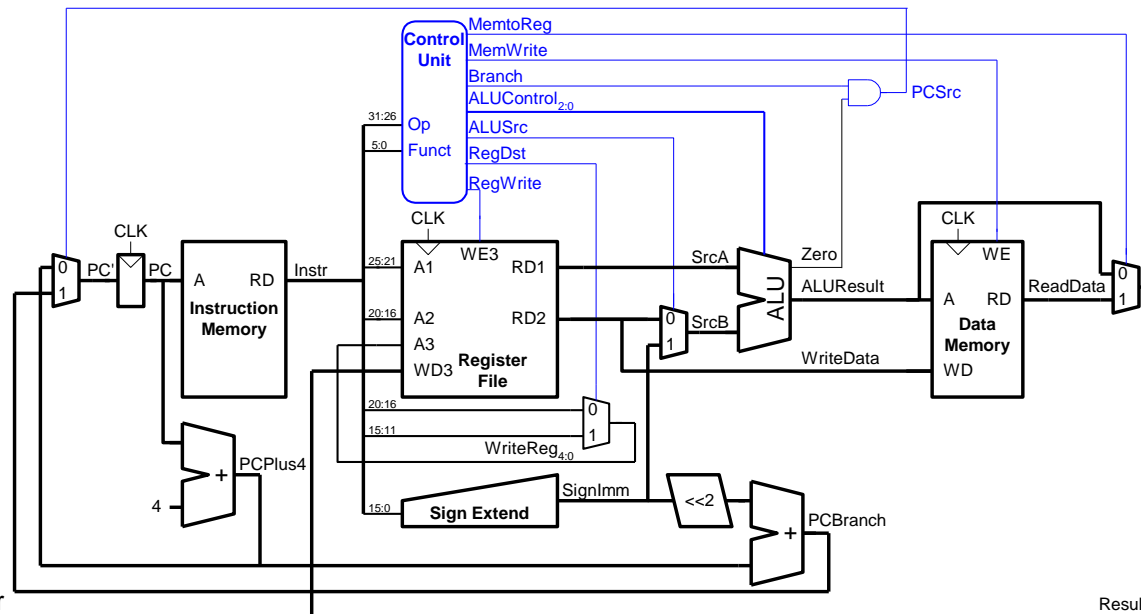
ALUOp _{1:0}	Meaning
00	Add
01	Subtract
10	Look at Funct
11	Not Used

ALUOp _{1:0}	Funct	ALUControl _{2:0}
00	X	010 (Add)
X1	X	110 (Subtract)
1X	100000 (add)	010 (Add)
1X	100010 (sub)	110 (Subtract)
1X	100100 (and)	000 (And)
1X	100101 (or)	001 (Or)
1X	101010 (slt)	111 (SLT)



Control Unit: Main Decoder

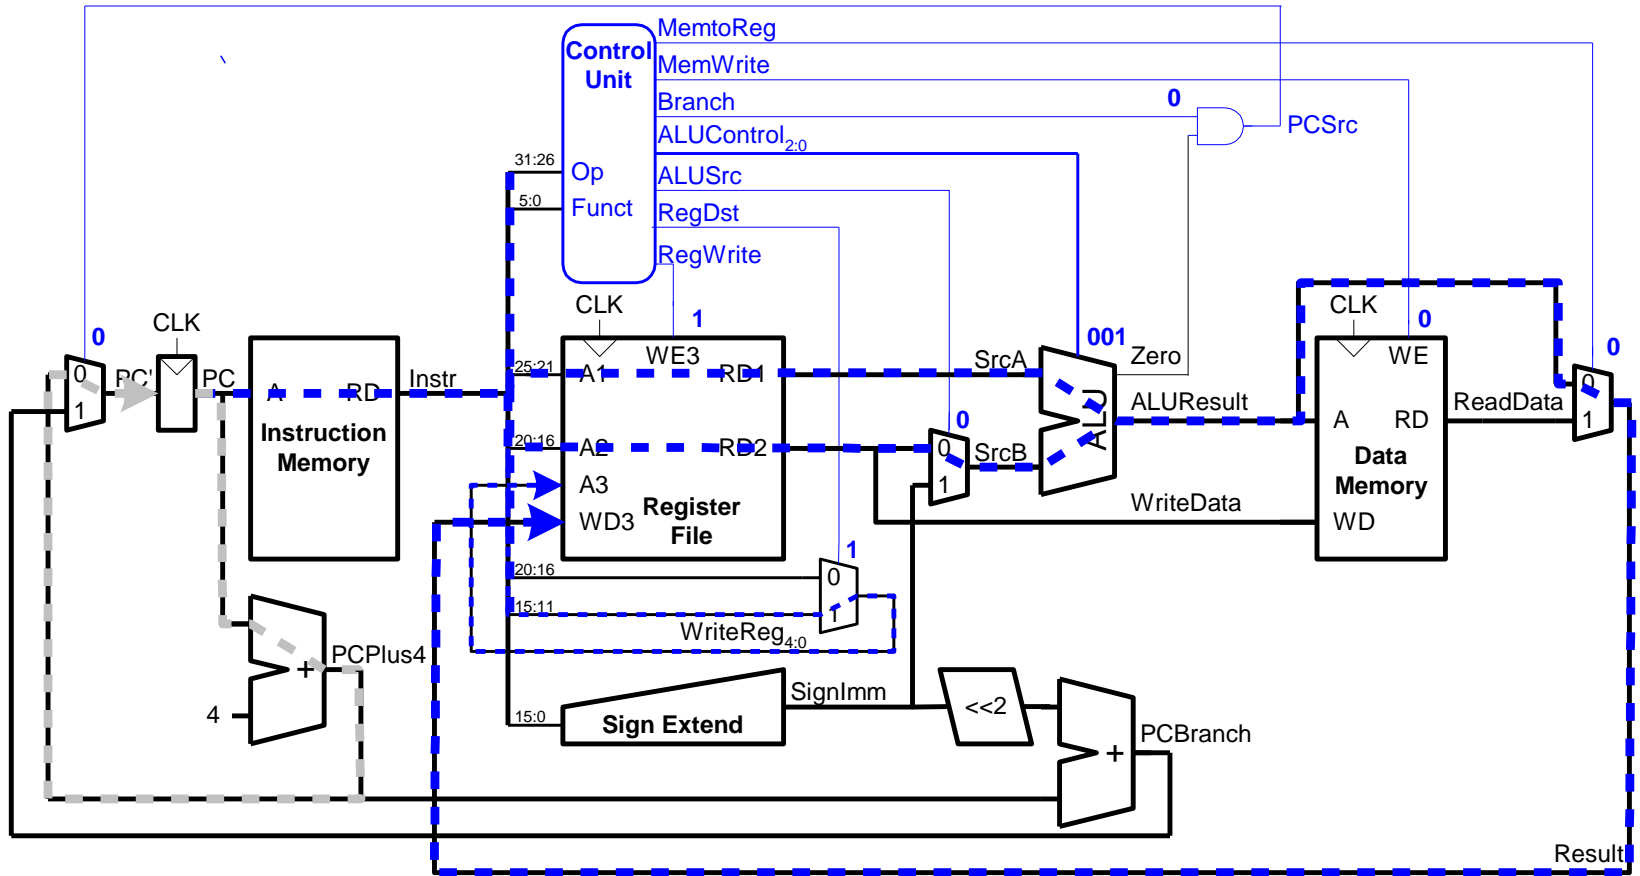
Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000							
lw	100011							
sw	101011							
beq	000100							



Control Unit: Main Decoder

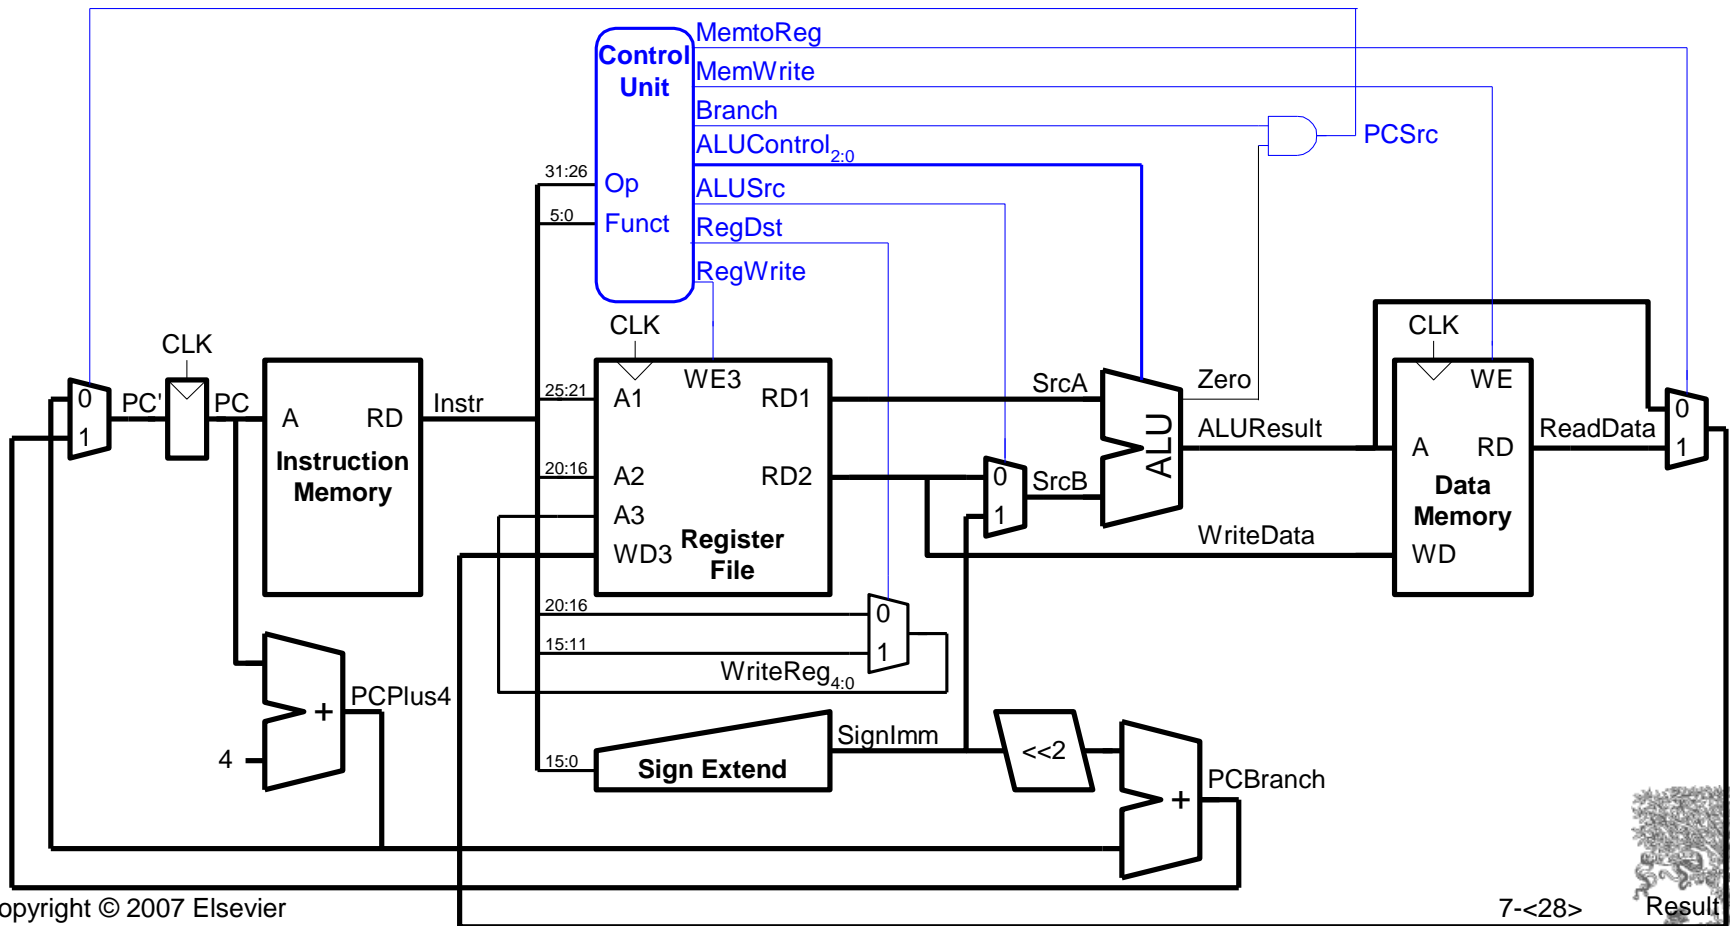
Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01

Single-Cycle Datapath Example: or



Extended Functionality: addi

- No change to datapath



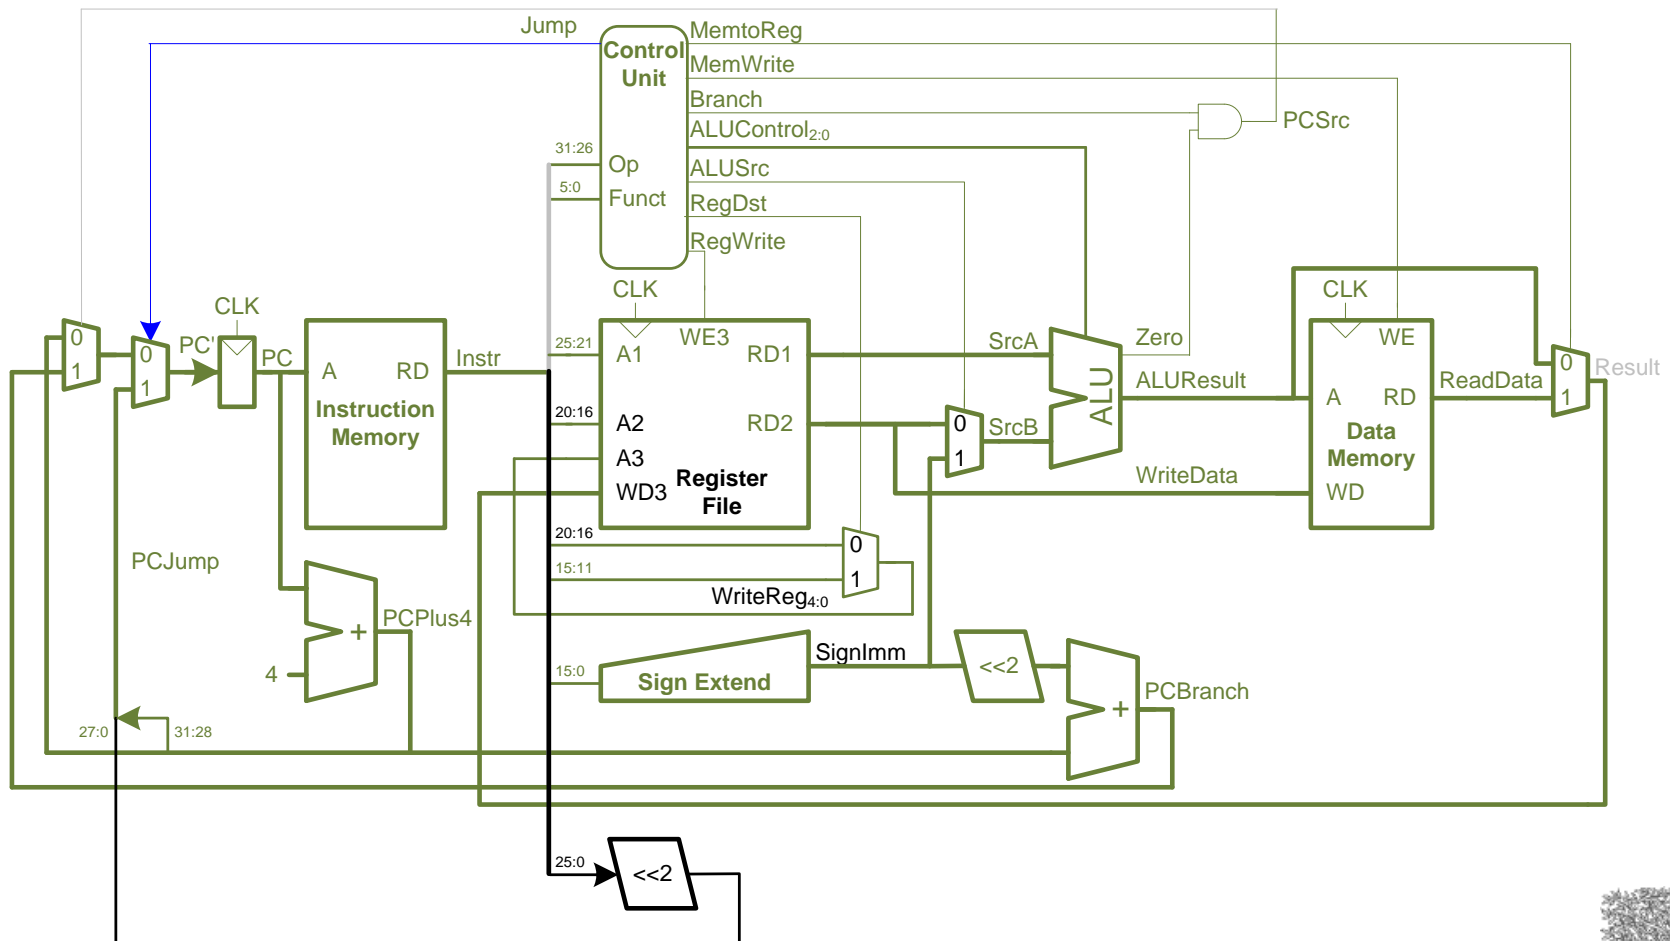
Control Unit: `addi`

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
<code>addi</code>	001000							

Control Unit: `addi`

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
<code>addi</code>	001000	1	0	1	0	0	0	00

Extended Functionality: j



Control Unit: Main Decoder

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}	Jump
R-type	000000	1	1	0	0	0	0	10	0
lw	100011	1	0	1	0	0	1	00	0
sw	101011	0	X	1	0	1	X	00	0
beq	000100	0	X	0	1	0	X	01	0
j	000100								

Control Unit: Main Decoder

Instruction	Op _{5:0}	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp _{1:0}	Jump
R-type	000000	1	1	0	0	0	0	10	0
lw	100011	1	0	1	0	0	1	00	0
sw	101011	0	X	1	0	1	X	00	0
beq	000100	0	X	0	1	0	X	01	0
j	000100	0	X	X	0	0	X	XX	1

Review: Processor Performance

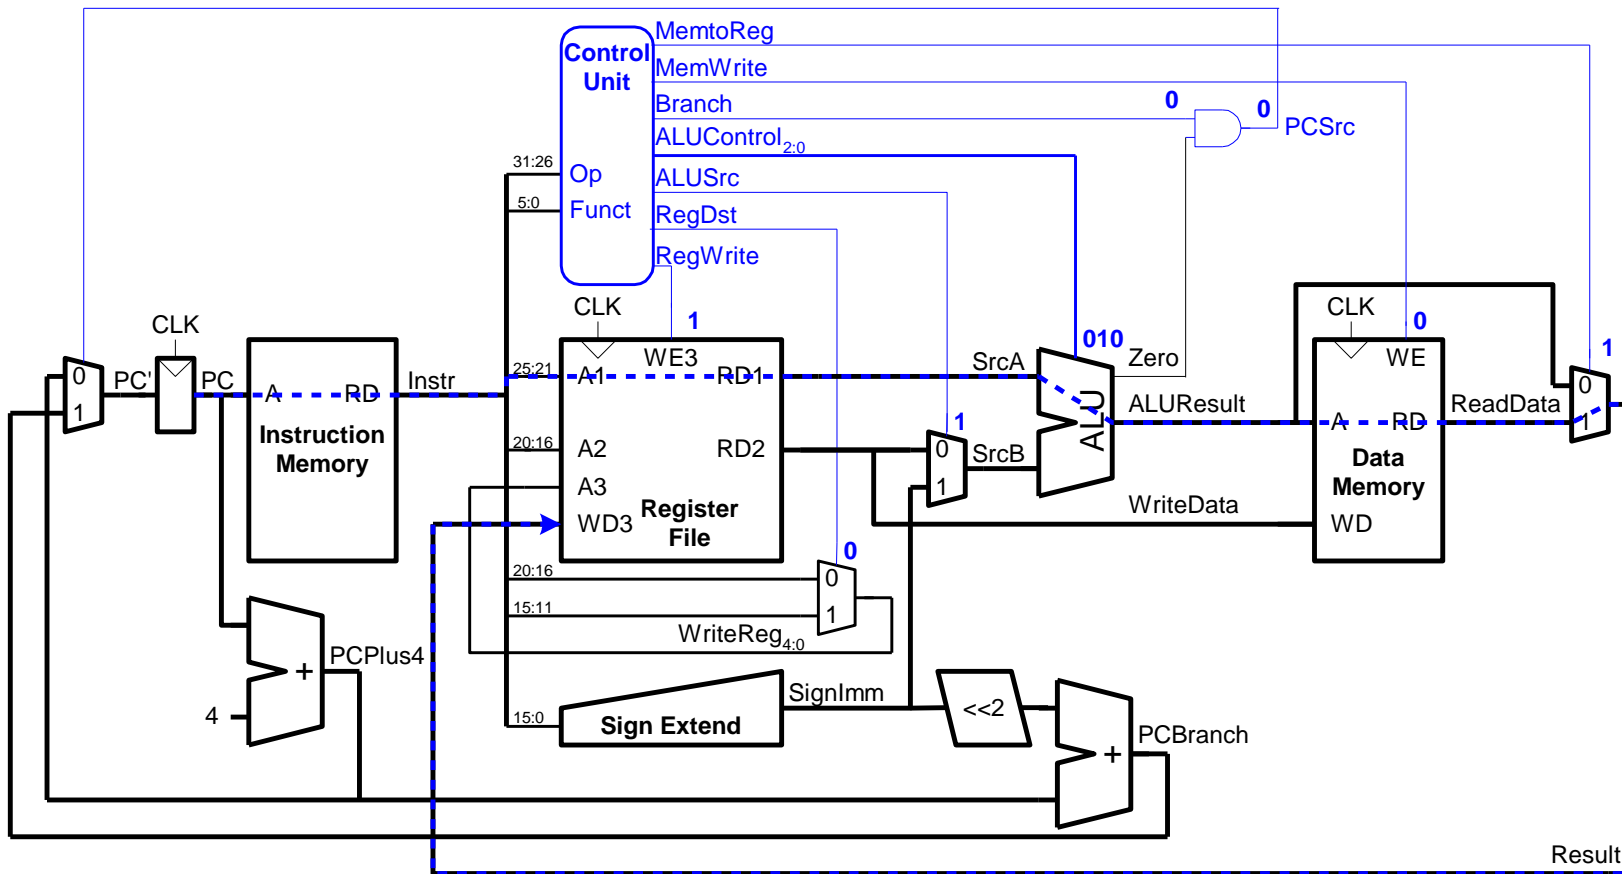
Program Execution Time

$$= (\# \text{ instructions})(\text{cycles/instruction})(\text{seconds/cycle})$$

$$= \# \text{ instructions} \times \text{CPI} \times T_c$$

Single-Cycle Performance

- T_C is limited by the critical path (1w)



Single-Cycle Performance

- Single-cycle critical path:

$$T_c = t_{pcq_PC} + t_{mem} + \max(t_{RFread}, t_{sext} + t_{mux}) + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

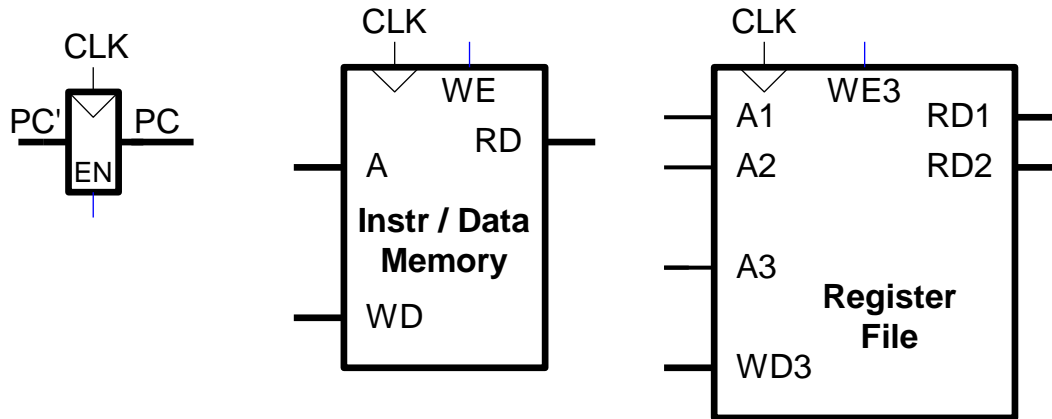
- In most implementations, limiting paths are:
 - memory, ALU, register file.
 - $T_c = t_{pcq_PC} + 2t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{RFsetup}$

Multicycle MIPS Processor

- Single-cycle microarchitecture:
 - + simple
 - cycle time limited by longest instruction (1_w)
 - two adders/ALUs and two memories
- Multicycle microarchitecture:
 - + higher clock speed
 - + simpler instructions run faster
 - + reuse expensive hardware on multiple cycles
 - sequencing overhead paid many times
- Same design steps: datapath & control

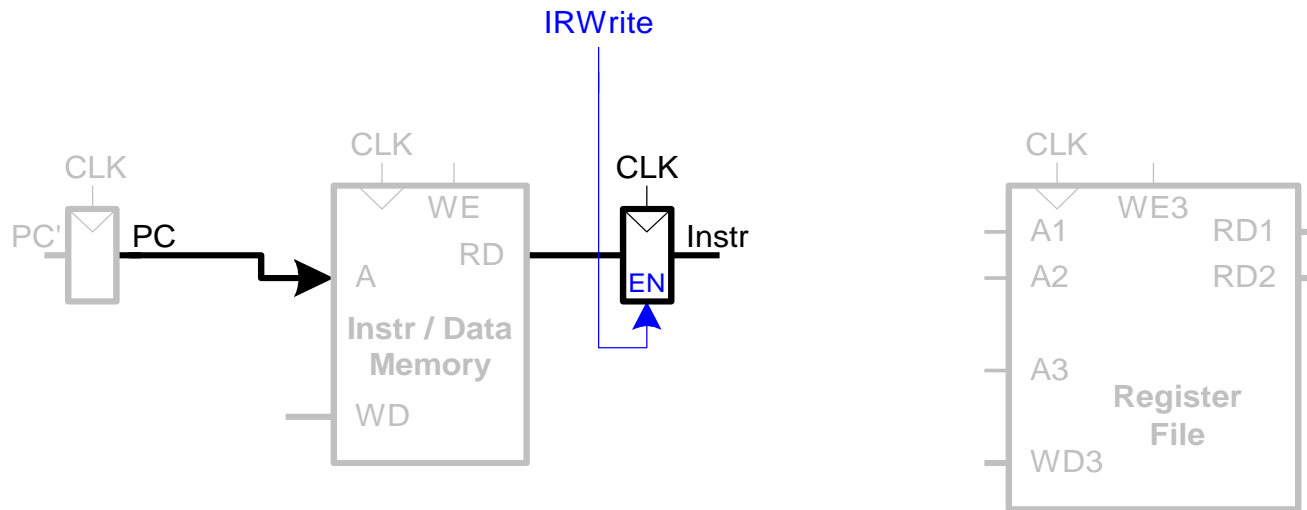
Multicycle State Elements

- Replace Instruction and Data memories with a single unified memory
 - More realistic

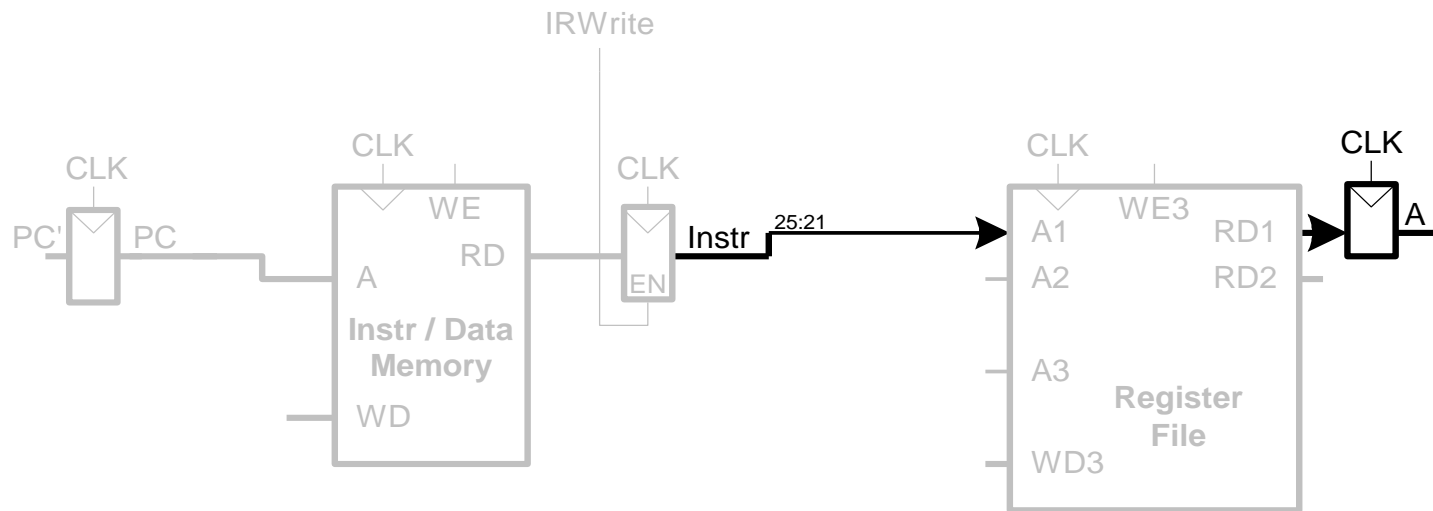


Multicycle Datapath: instruction fetch

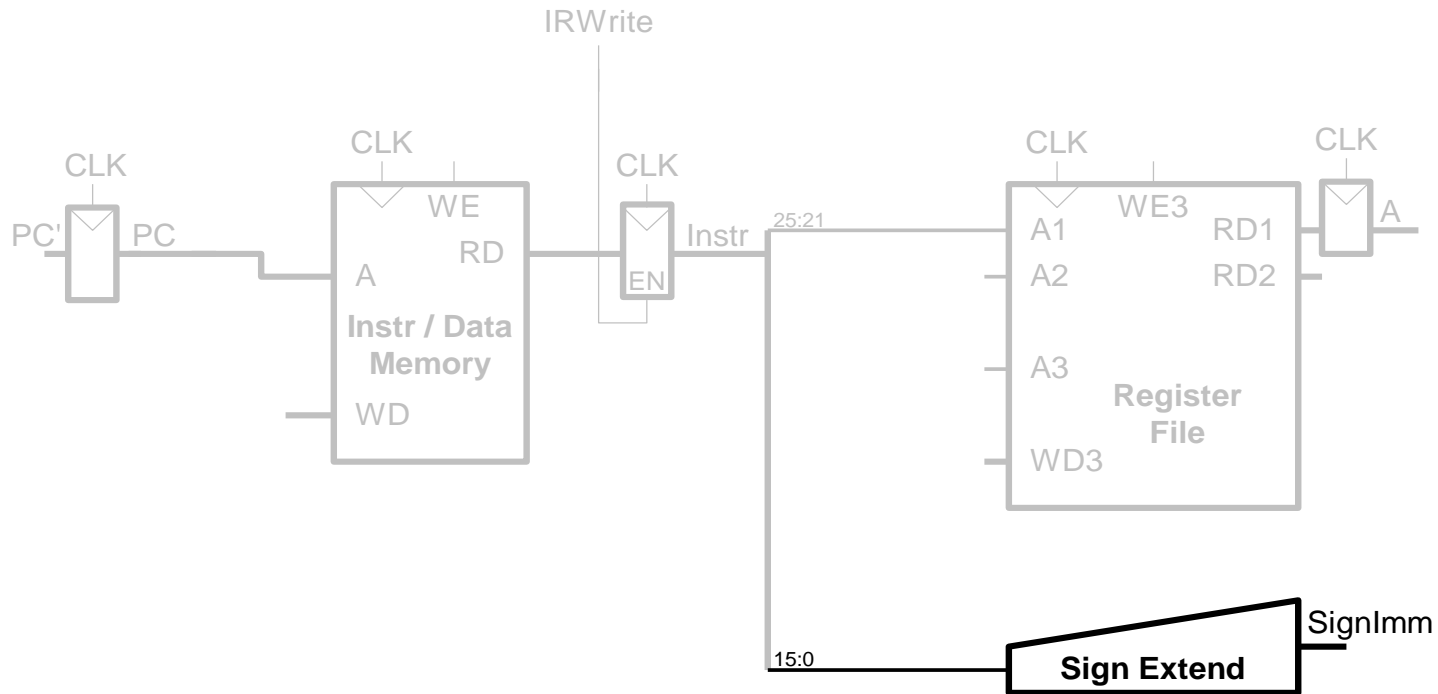
- First consider executing lw
- **STEP 1:** Fetch instruction



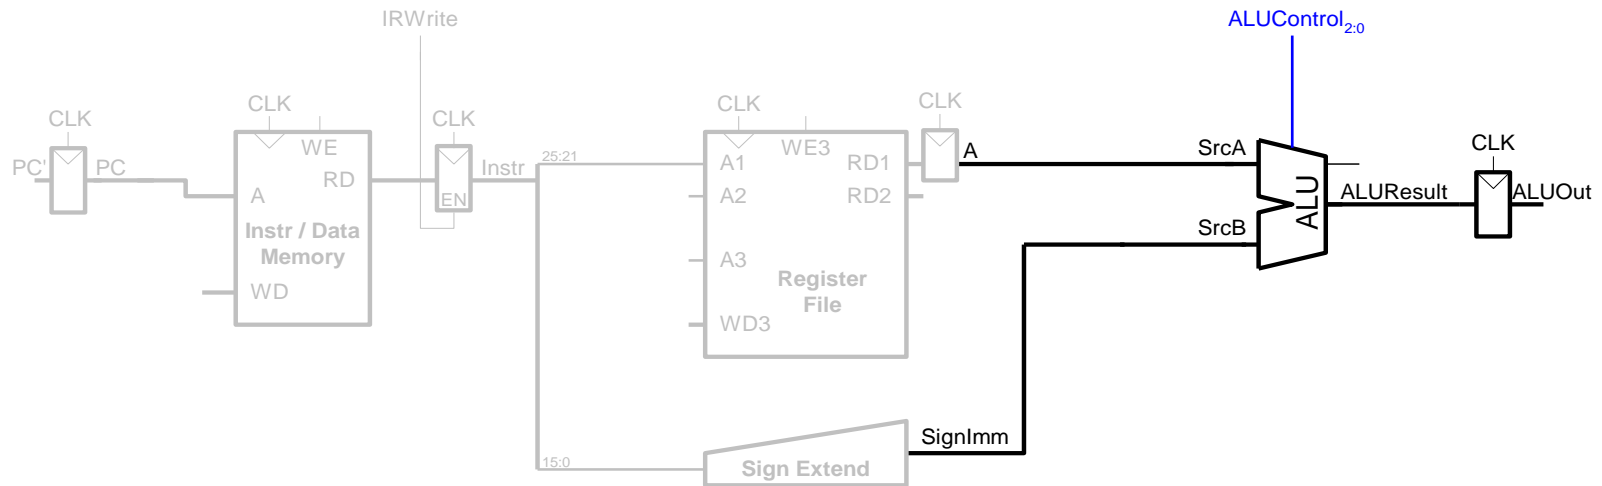
Multicycle Datapath: lw register read



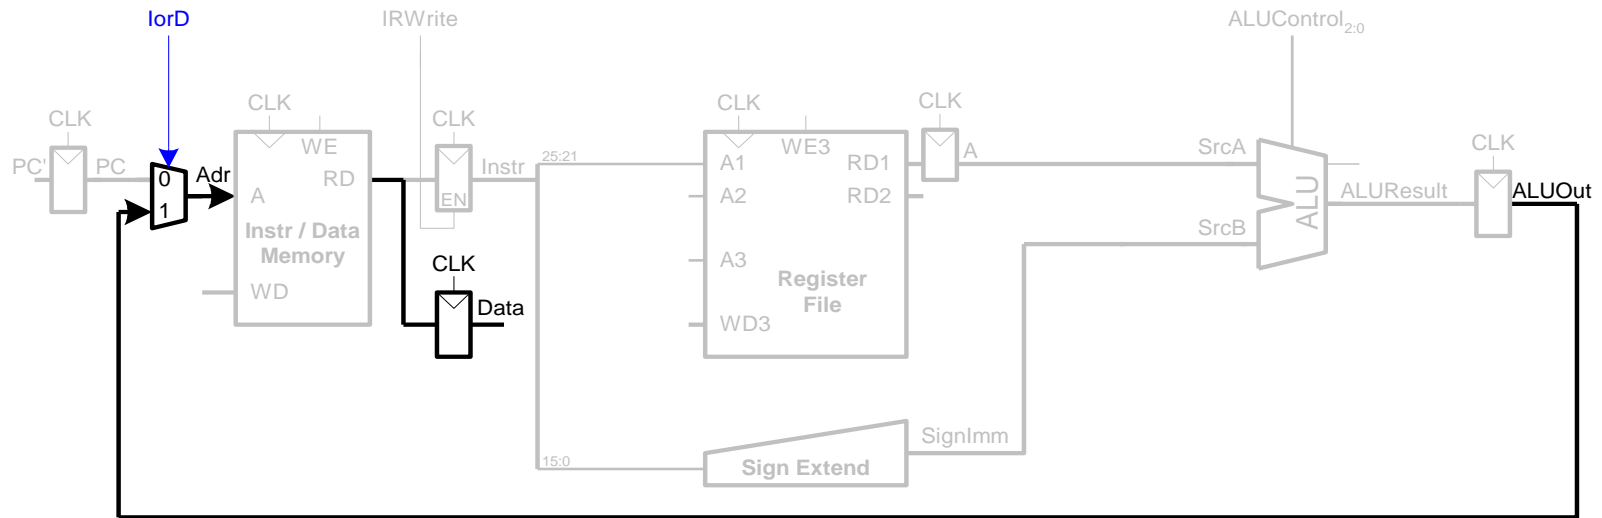
Multicycle Datapath: l_w immediate



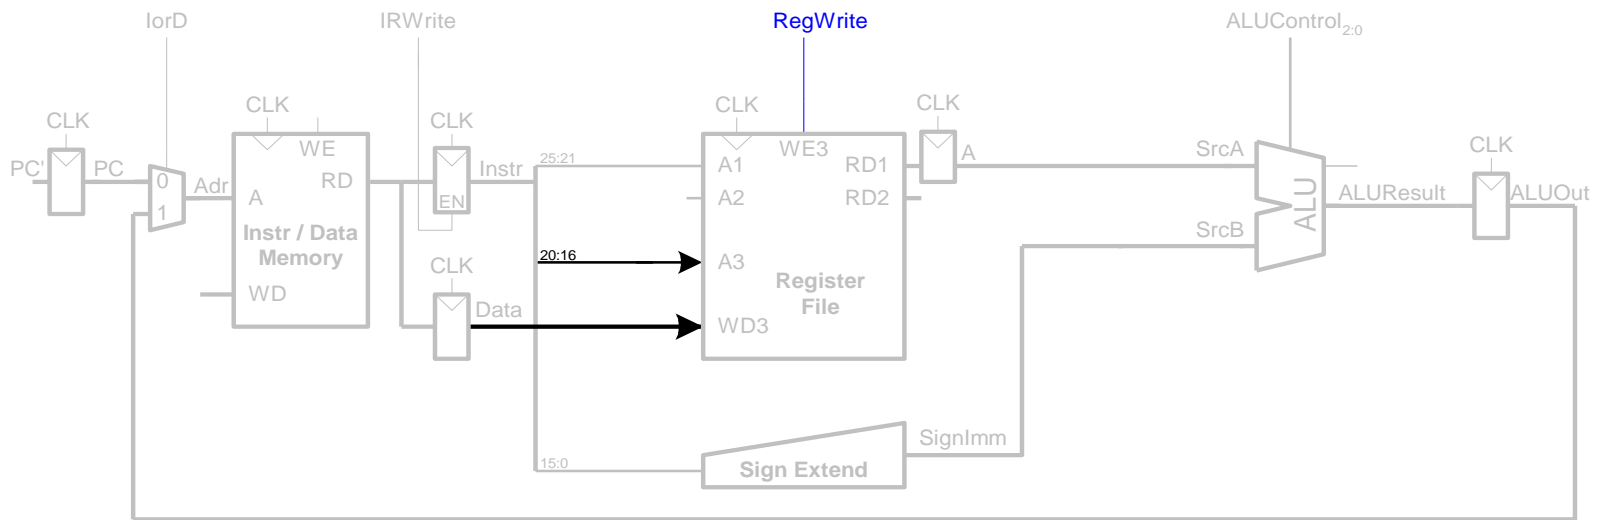
Multicycle Datapath: l_w address



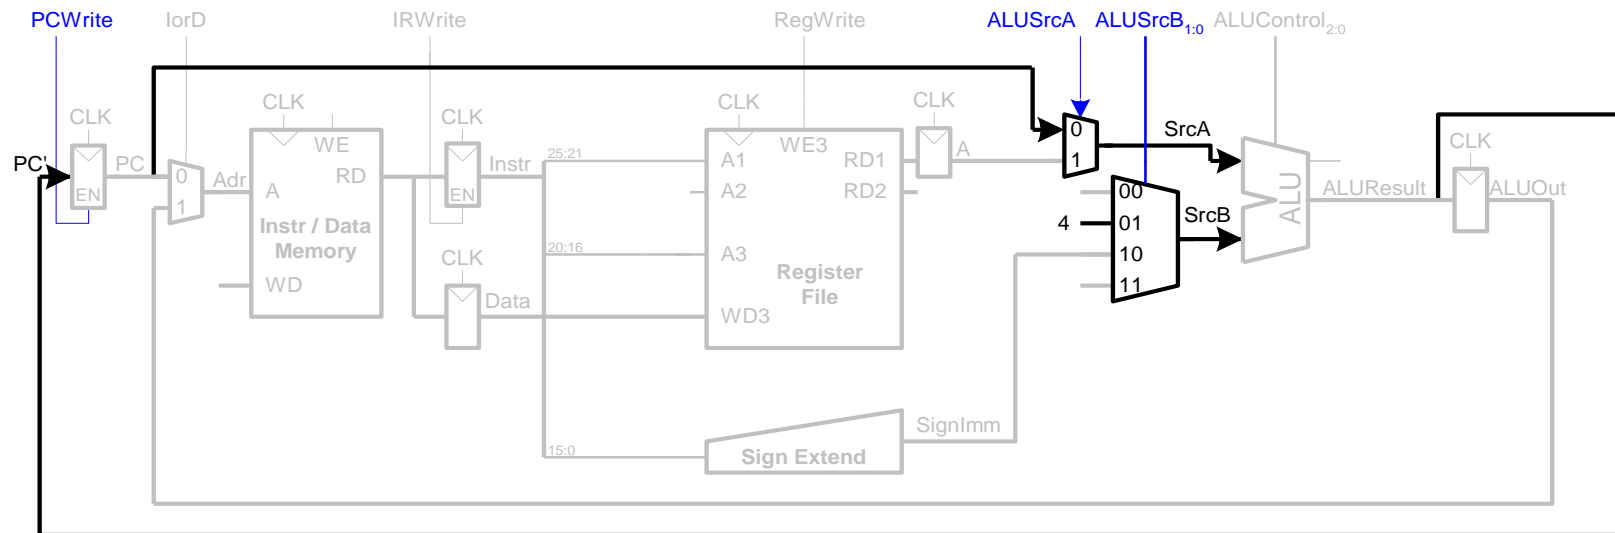
Multicycle Datapath: I_W memory read



Multicycle Datapath: lw write register

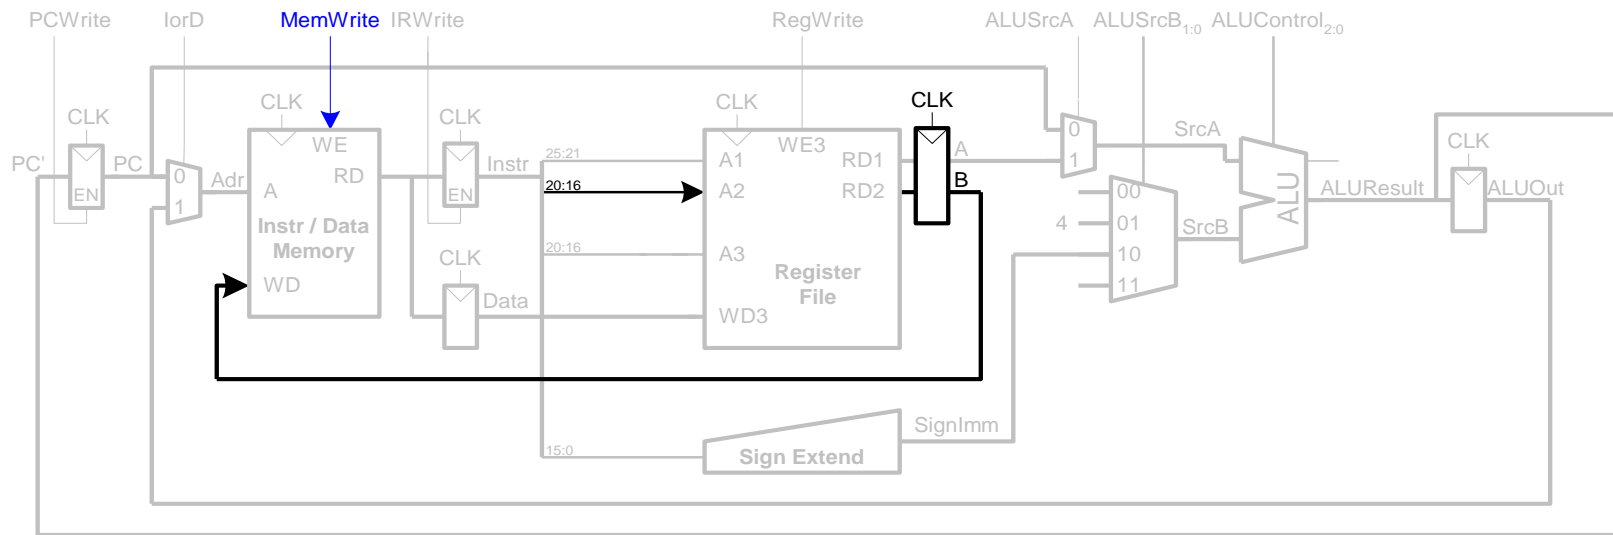


Multicycle Datapath: increment PC



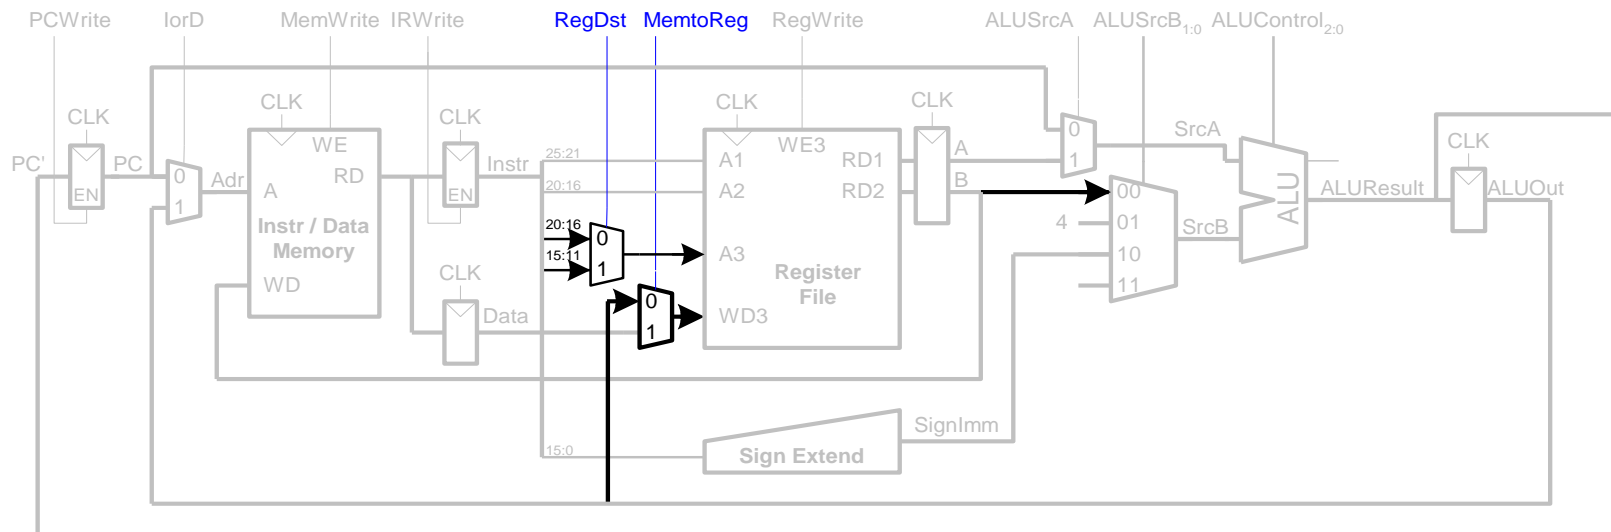
Multicycle Datapath: SW

- Write data in rt to memory



Multicycle Datapath: R-type Instructions

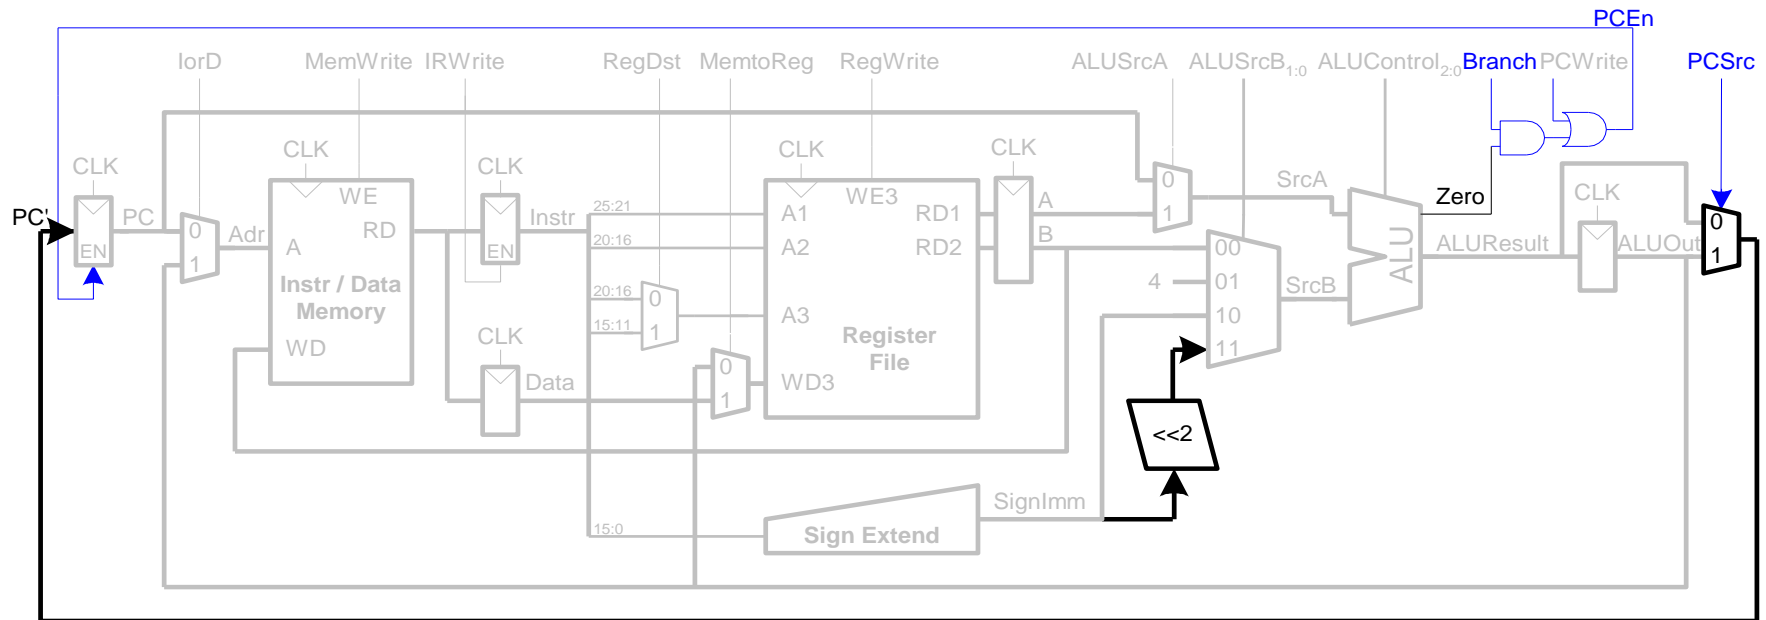
- Read from `rs` and `rt`
- Write *ALUResult* to register file
- Write to `rd` (instead of `rt`)



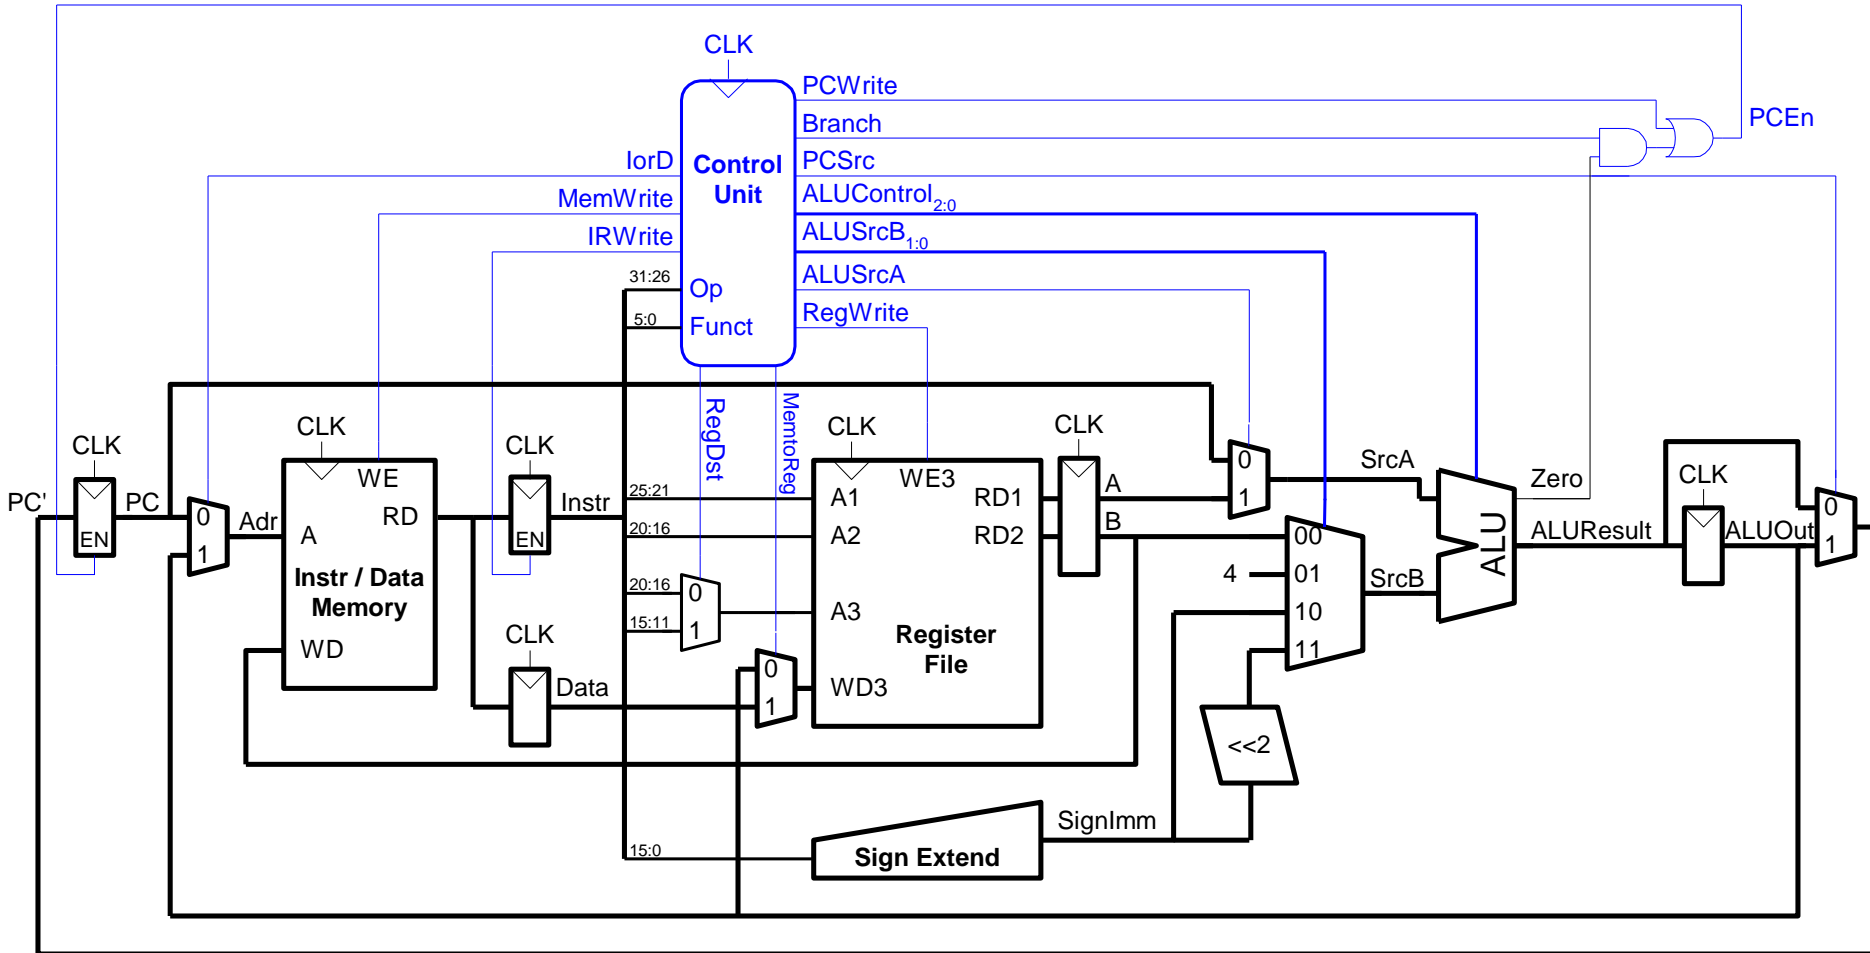
Multicycle Datapath: beq

- Determine whether values in `rs` and `rt` are equal
- Calculate branch target address:

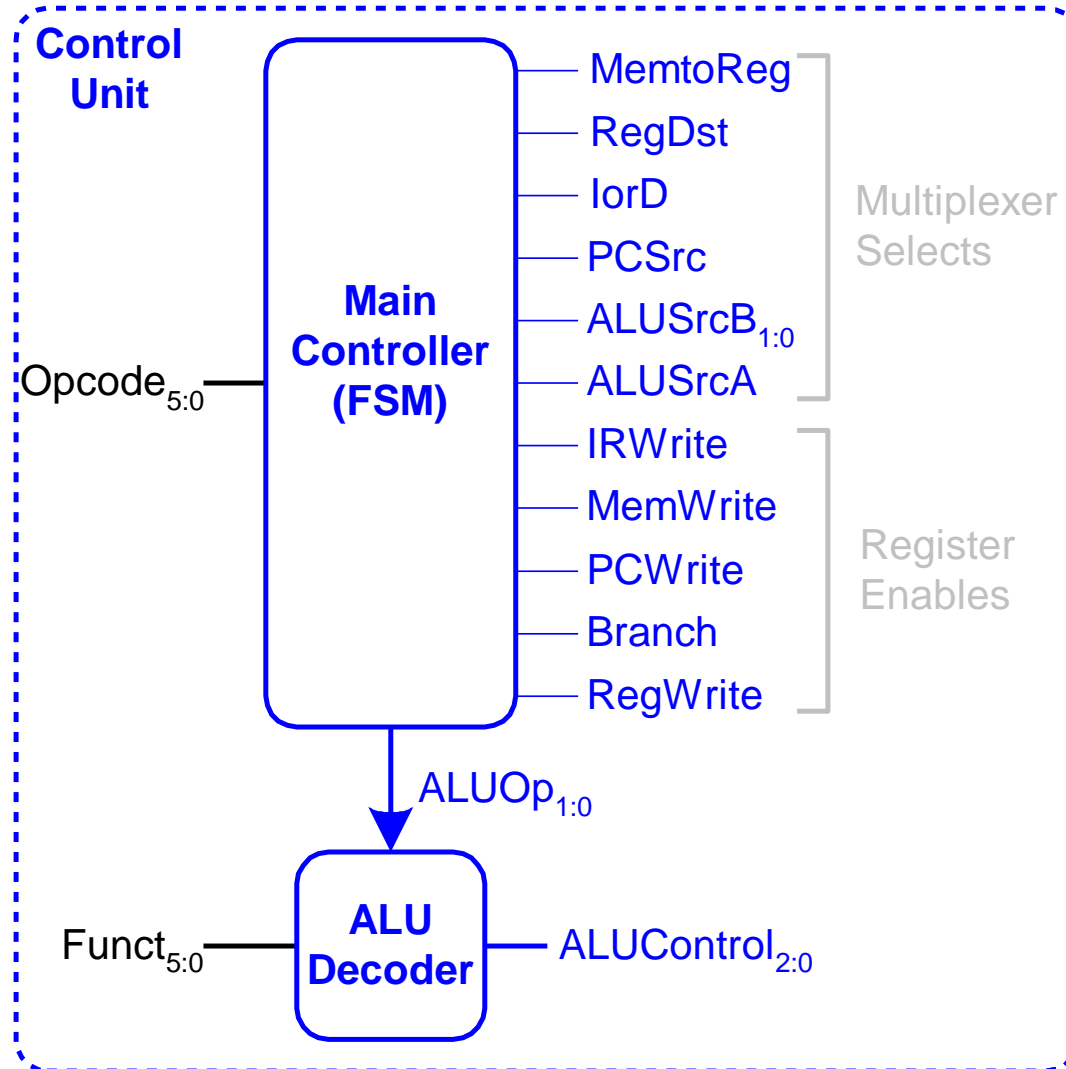
$$\text{BTA} = (\text{sign-extended immediate} \ll 2) + (\text{PC} + 4)$$



Complete Multicycle Processor



Control Unit



© 2004 Blackwell Publishing Ltd
Journal of Internal Medicine 255: 103–110

Reset

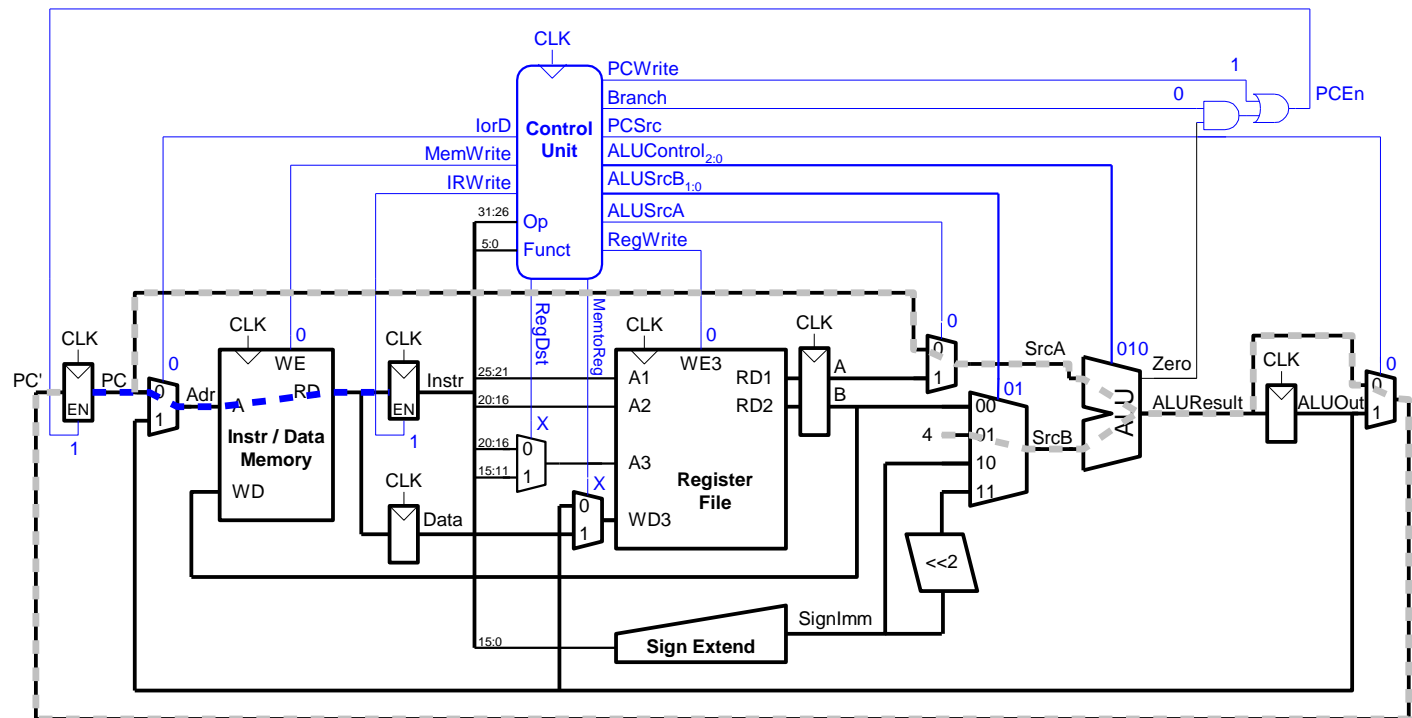


Main Controller FSM: Fetch

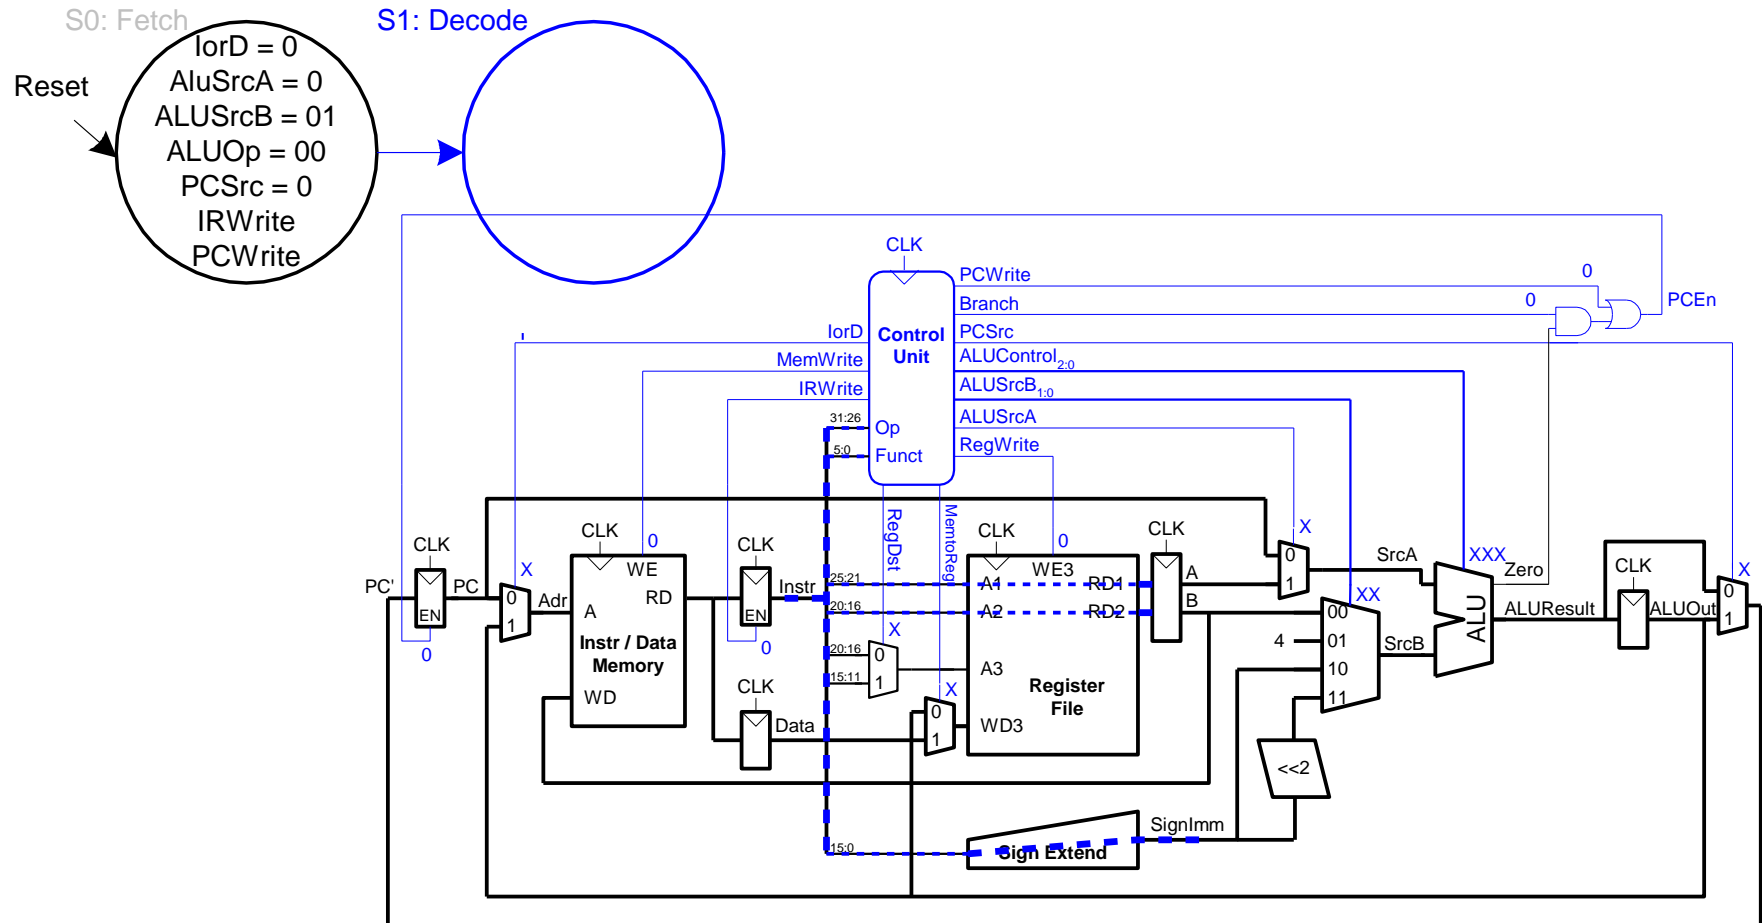
S0: Fetch

Reset

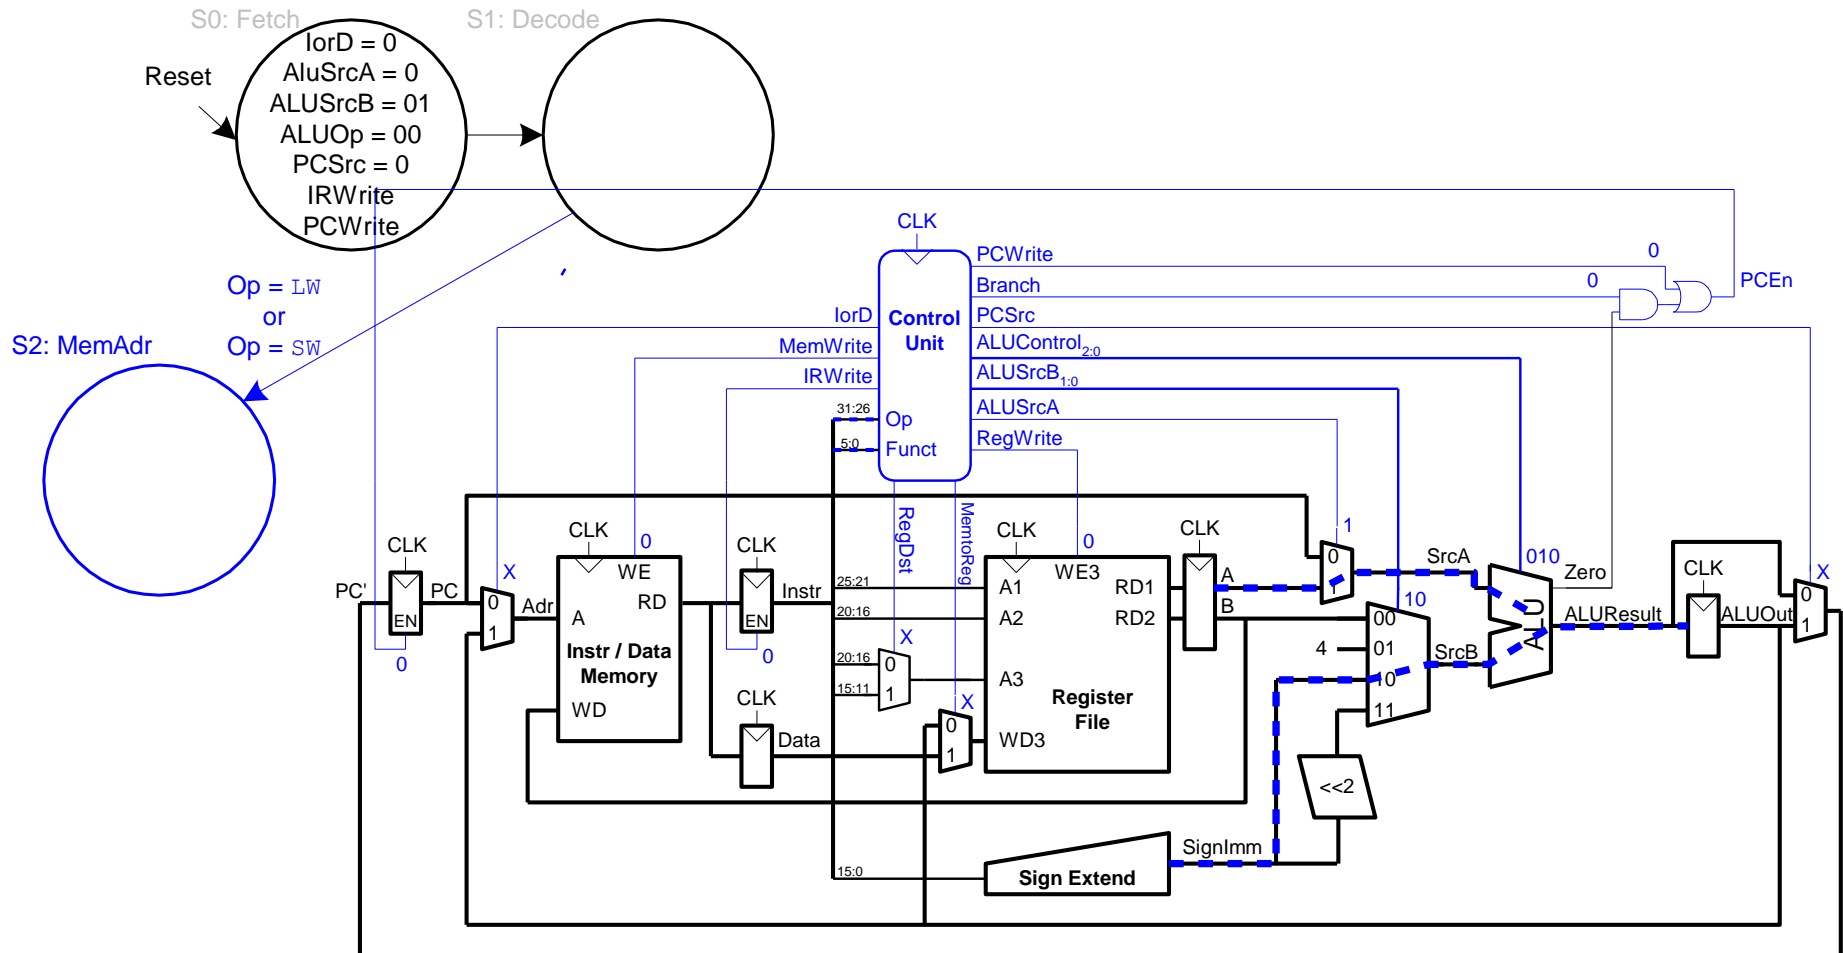
lorD = 0
AluSrcA = 0
ALUSrcB = 01
ALUOp = 00
PCSrc = 0
IRWrite
PCWrite



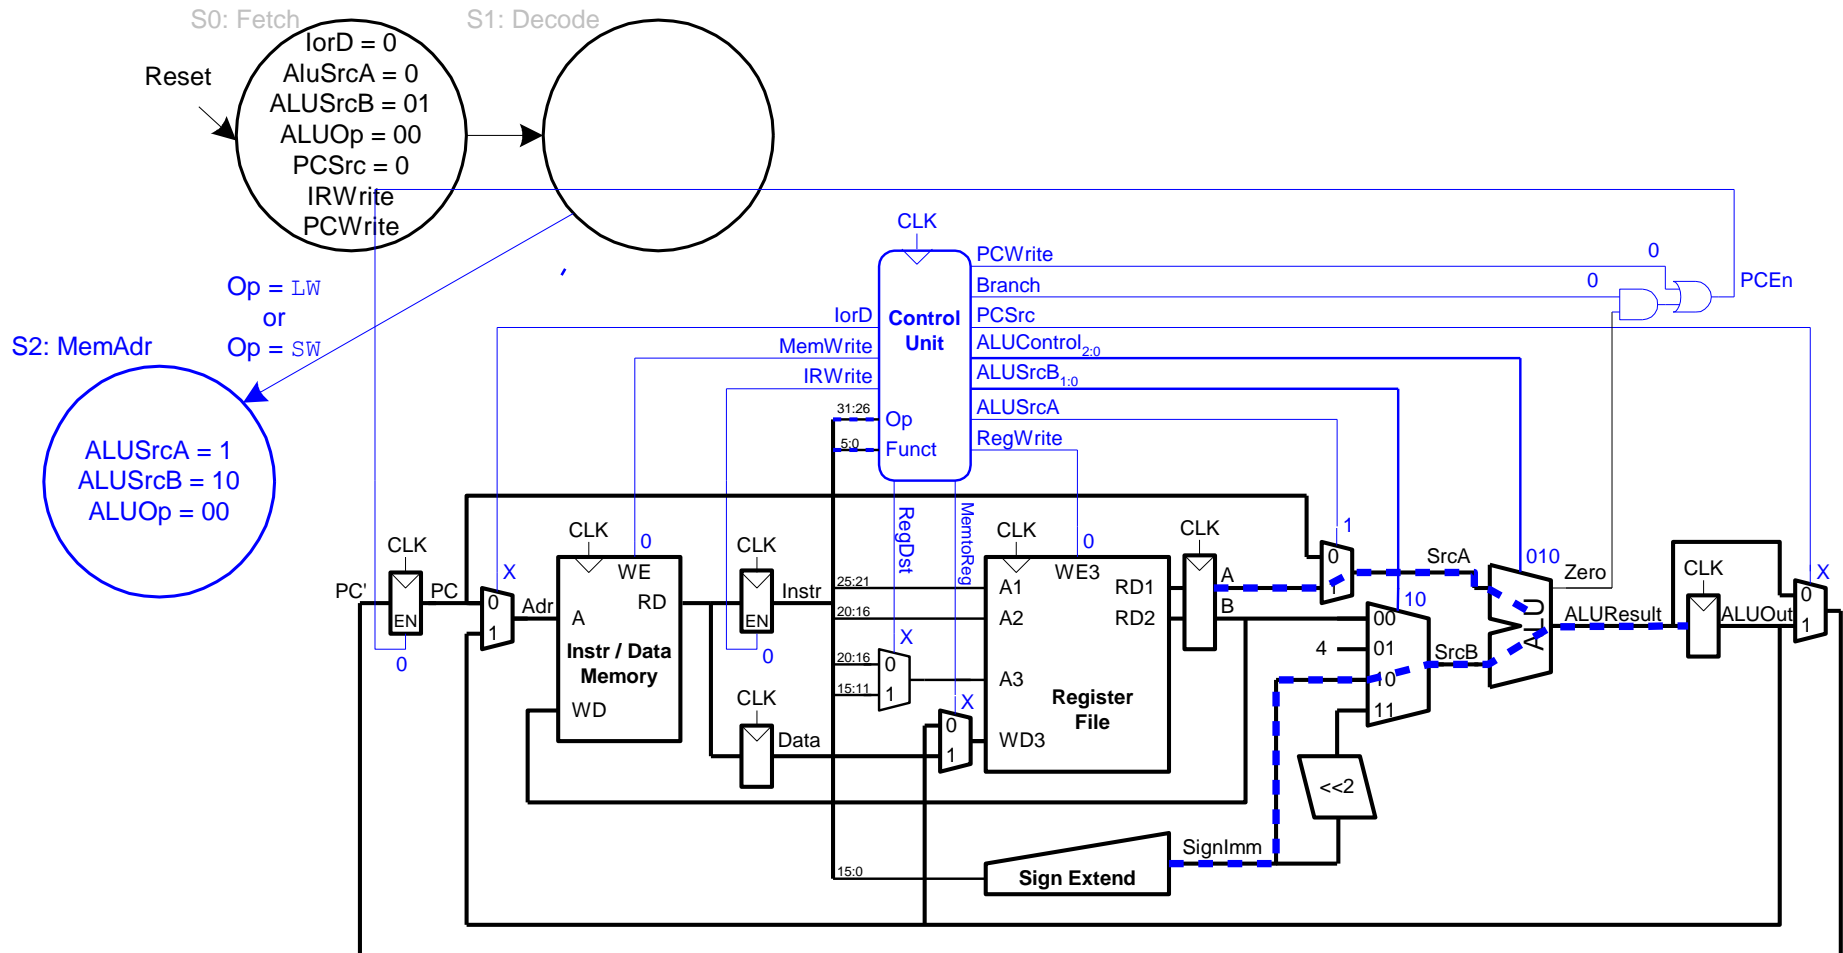
Main Controller FSM: Decode



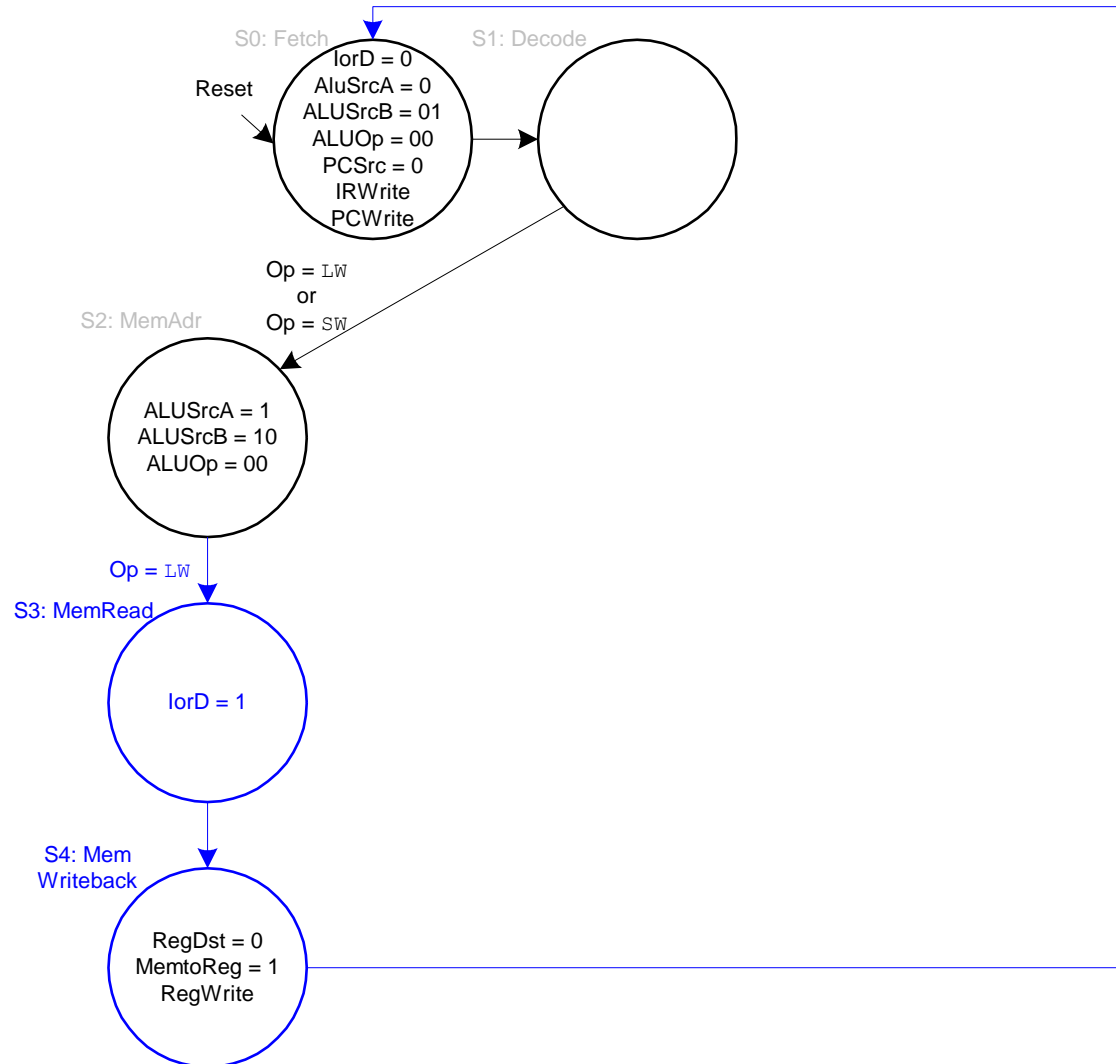
Main Controller FSM: Address Calculation



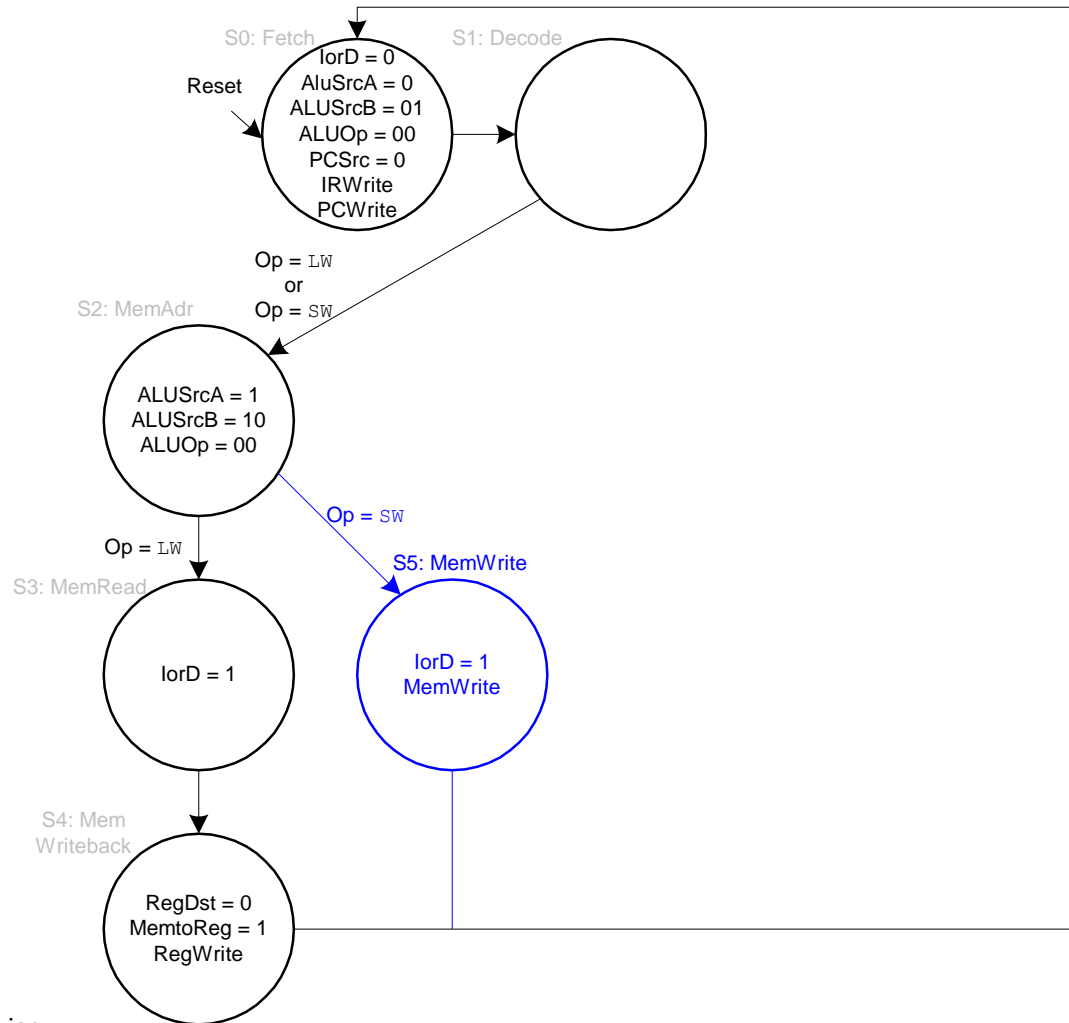
Main Controller FSM: Address Calculation



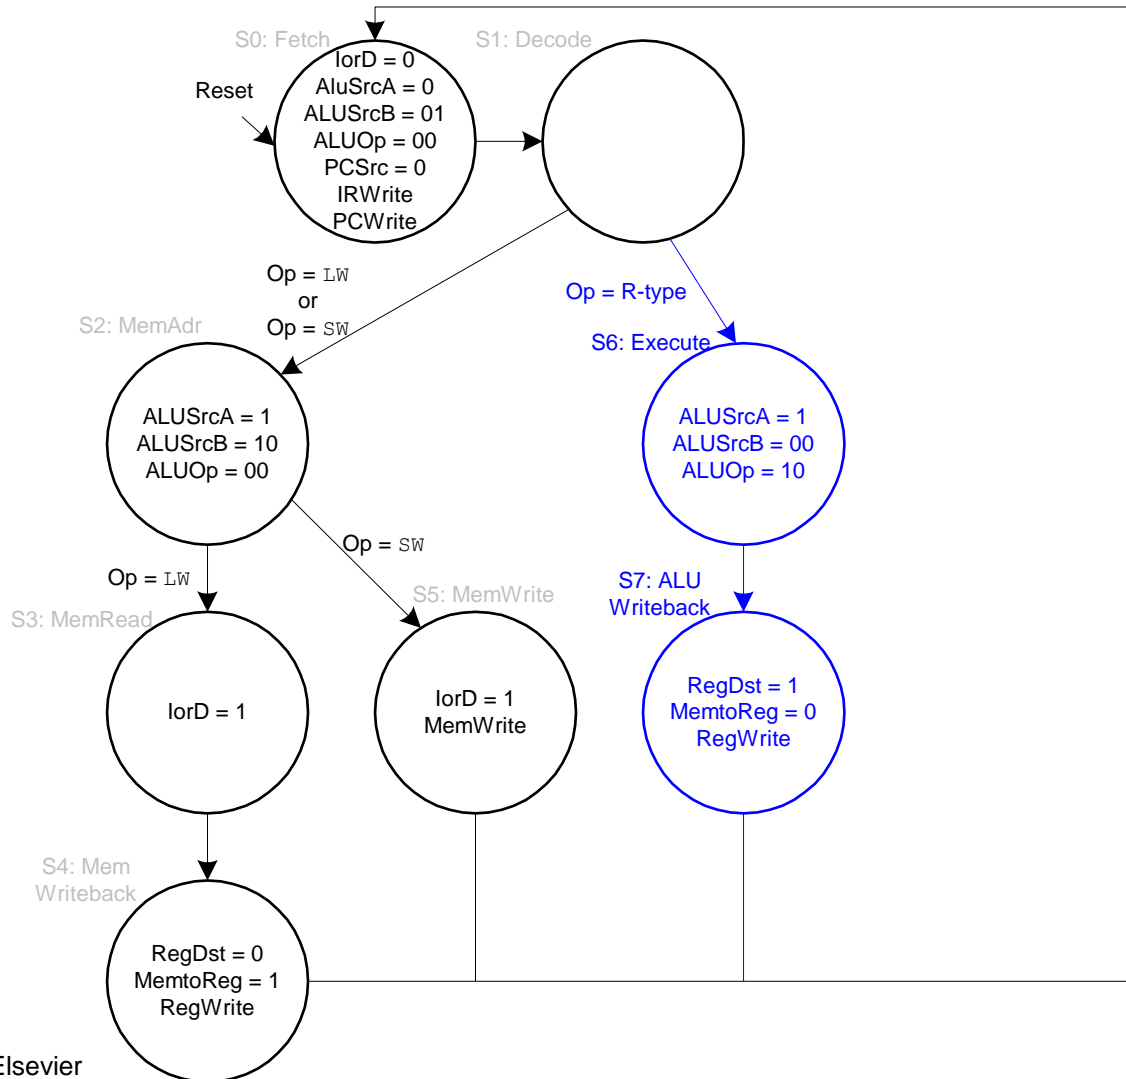
Main Controller FSM: lw



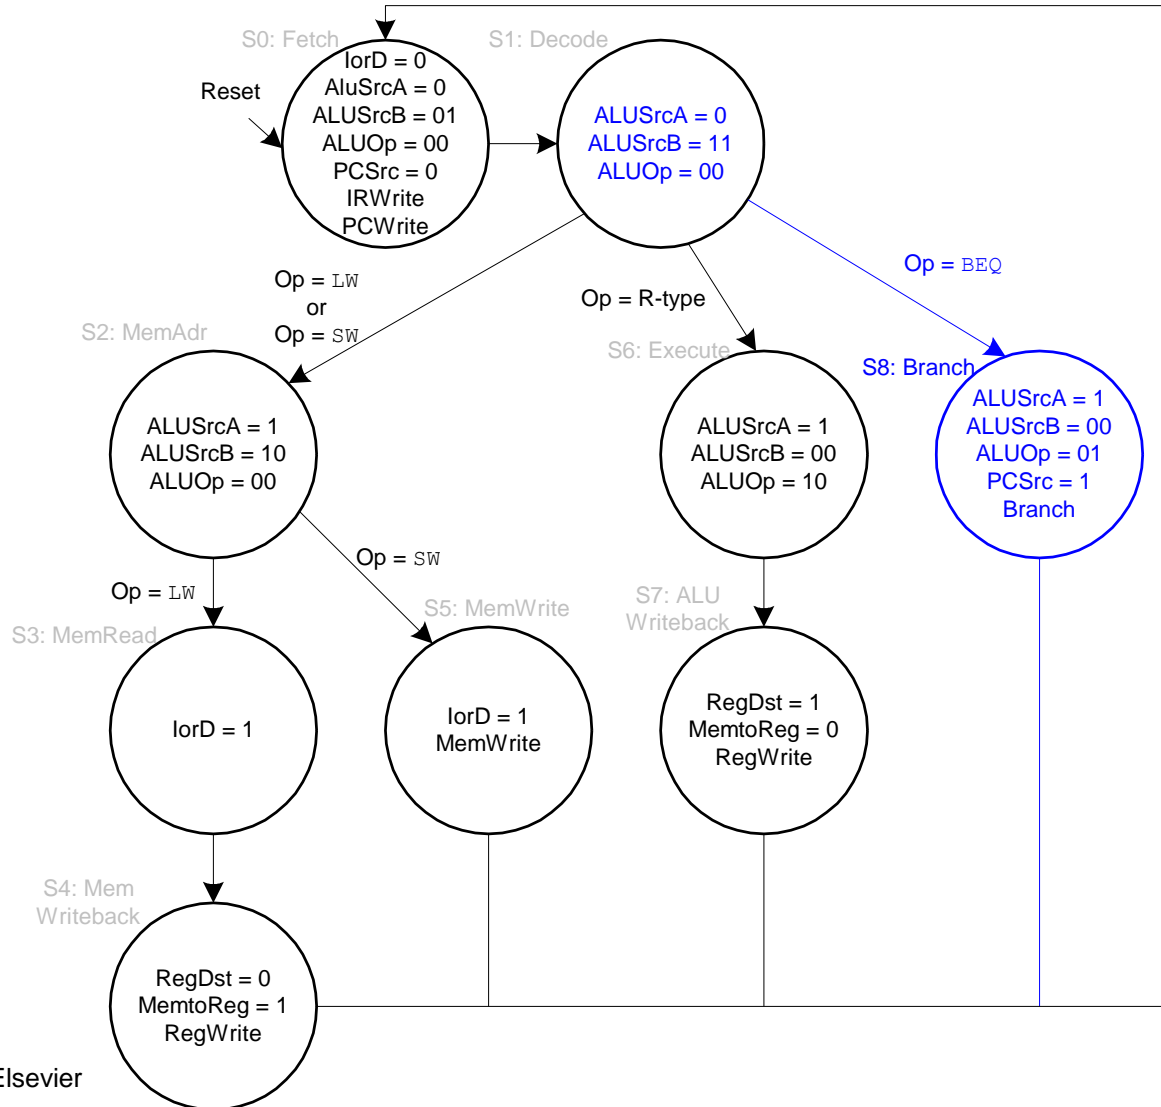
Main Controller FSM: SW



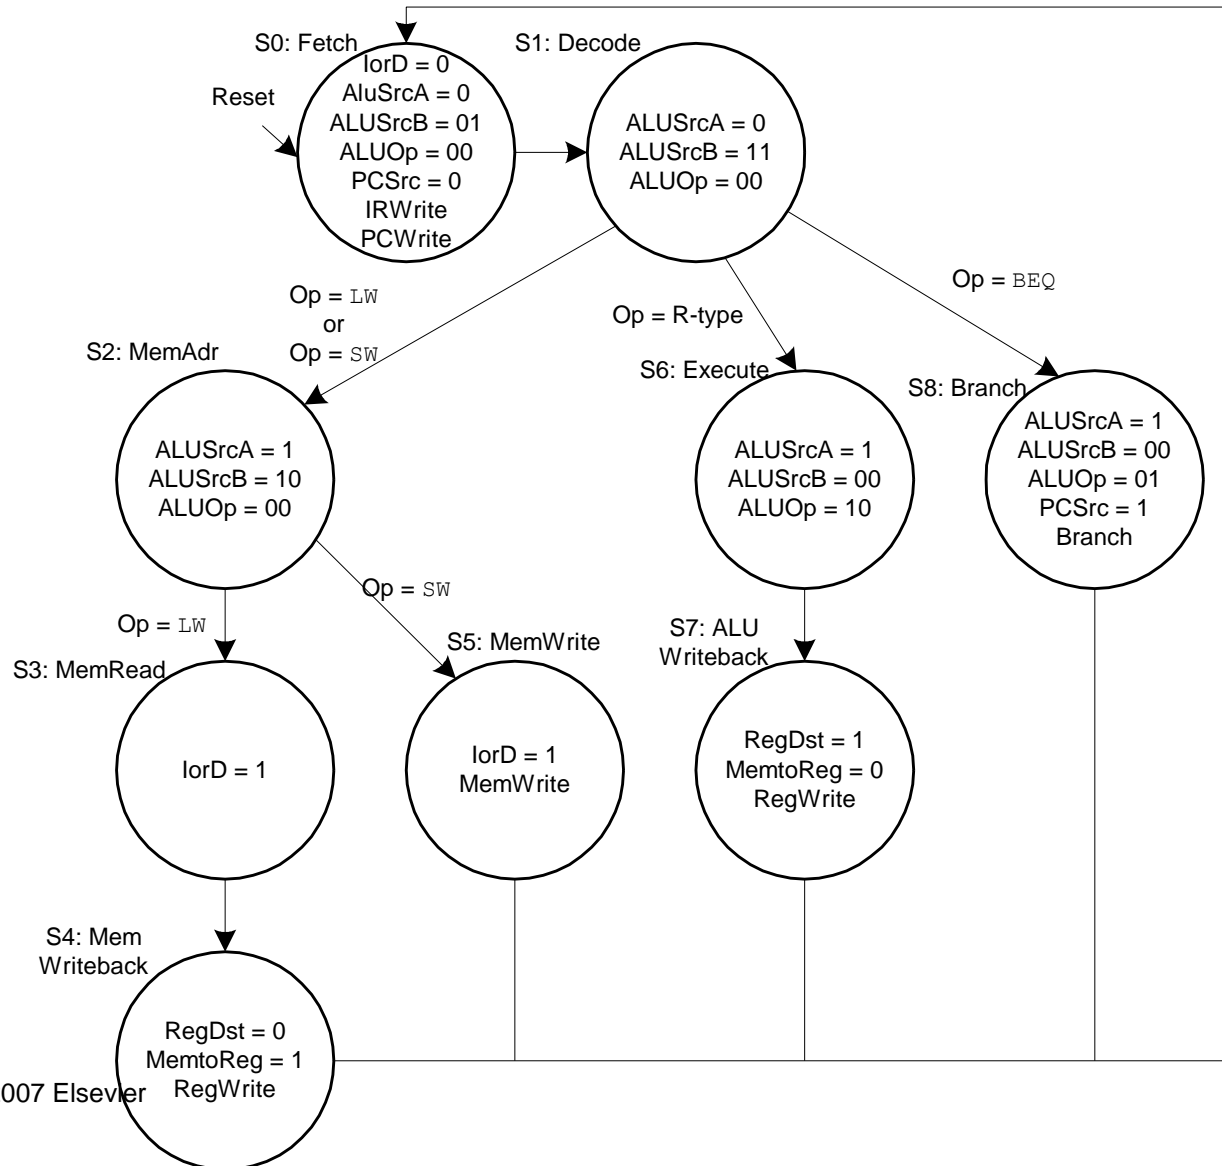
Main Controller FSM: R-Type



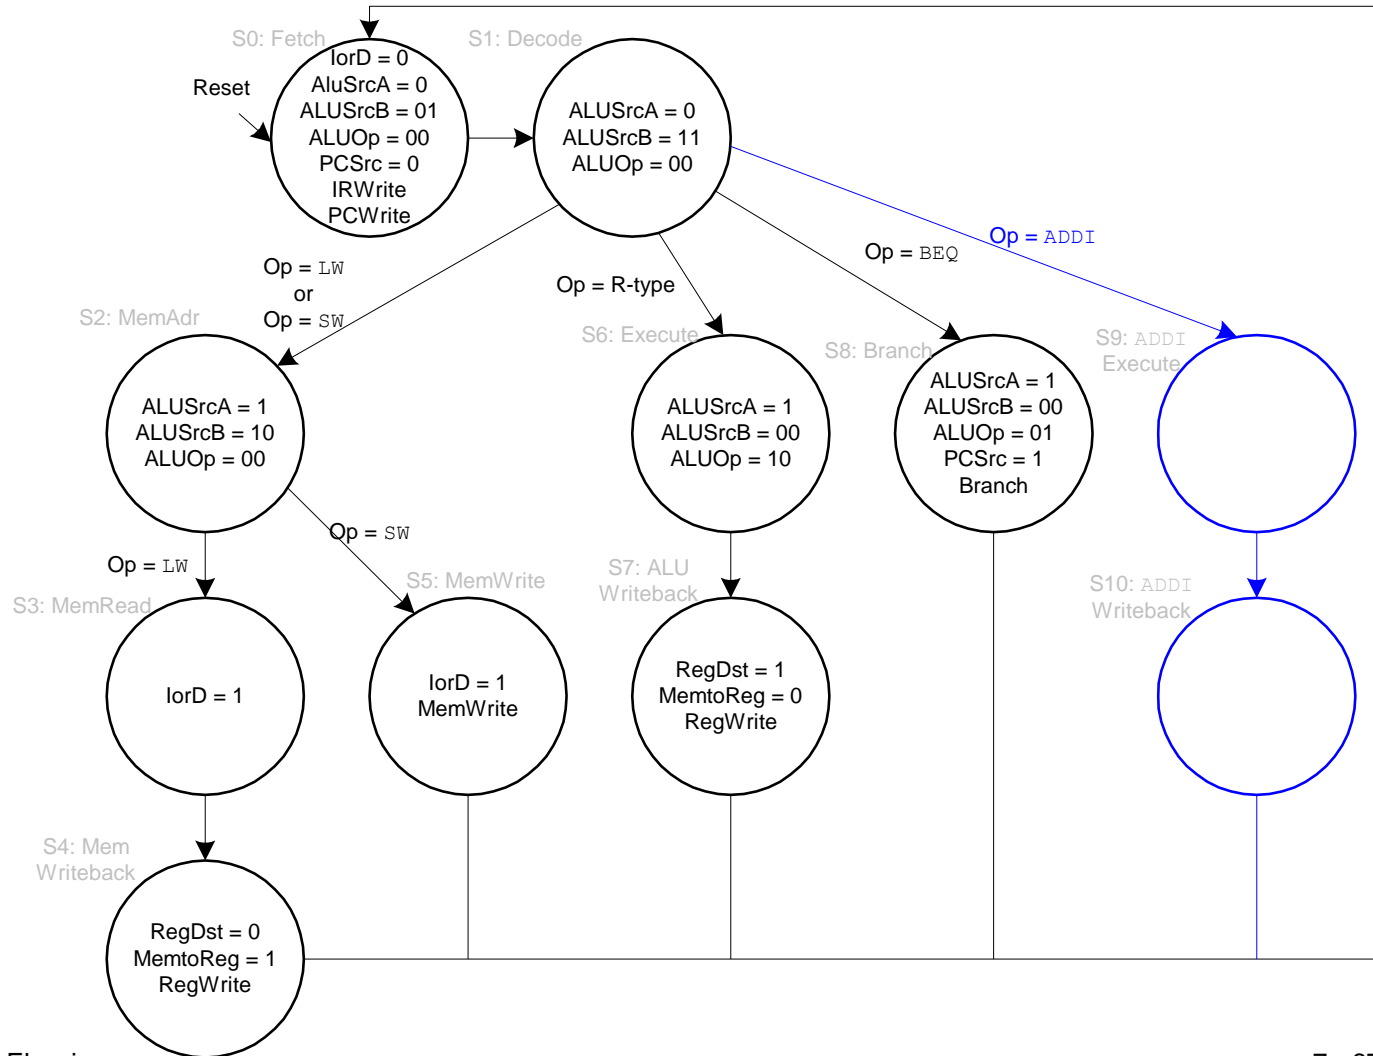
Main Controller FSM: beq



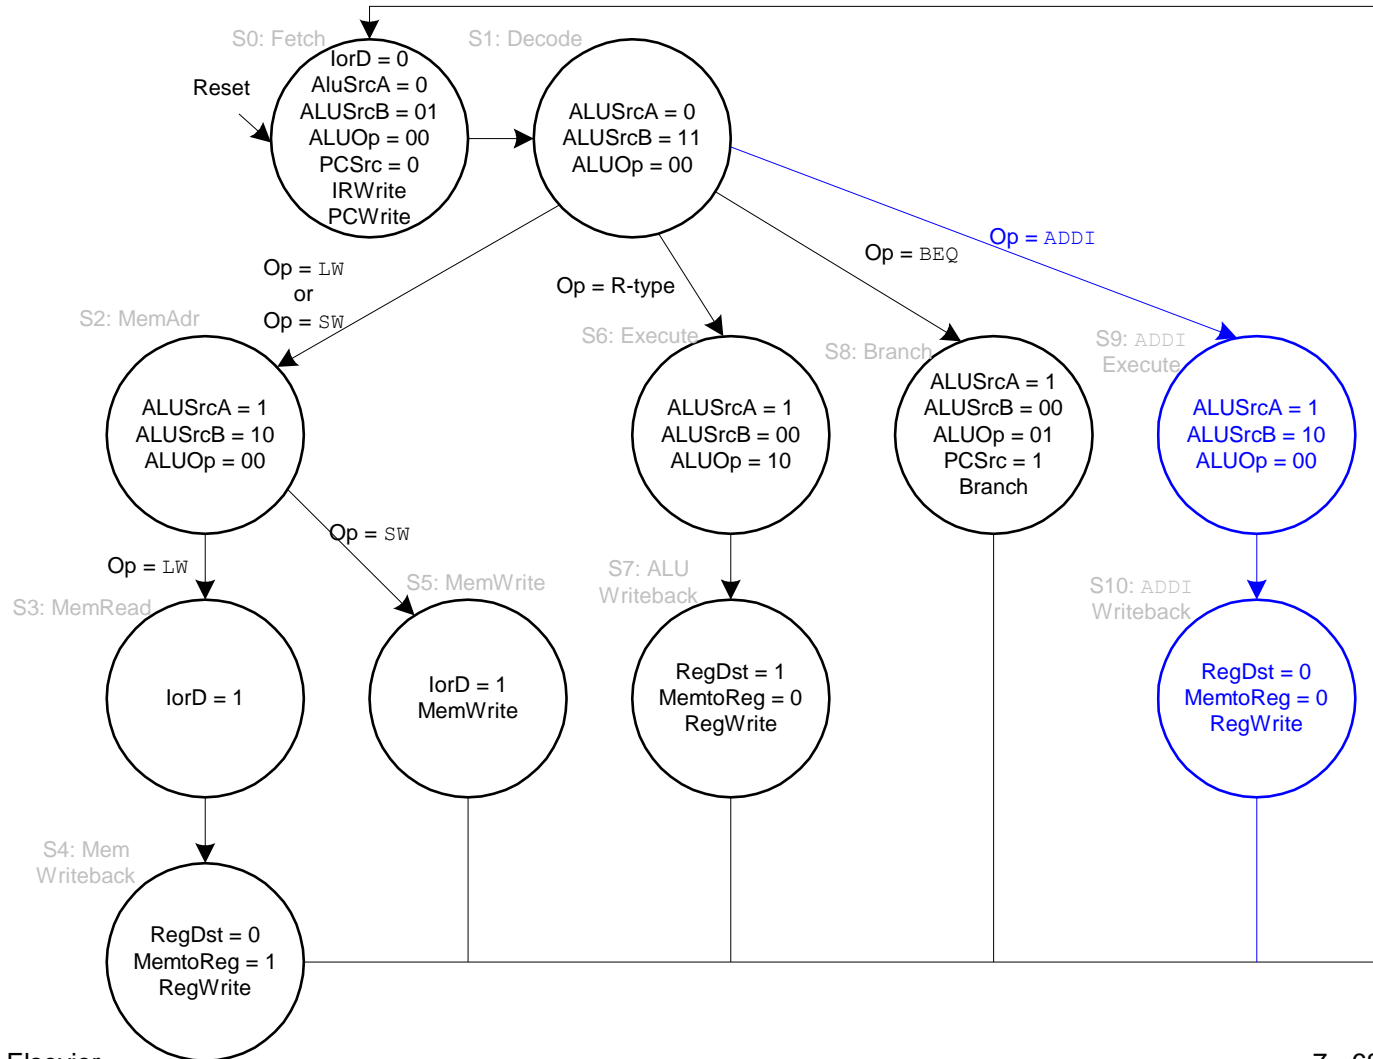
Complete Multicycle Controller FSM



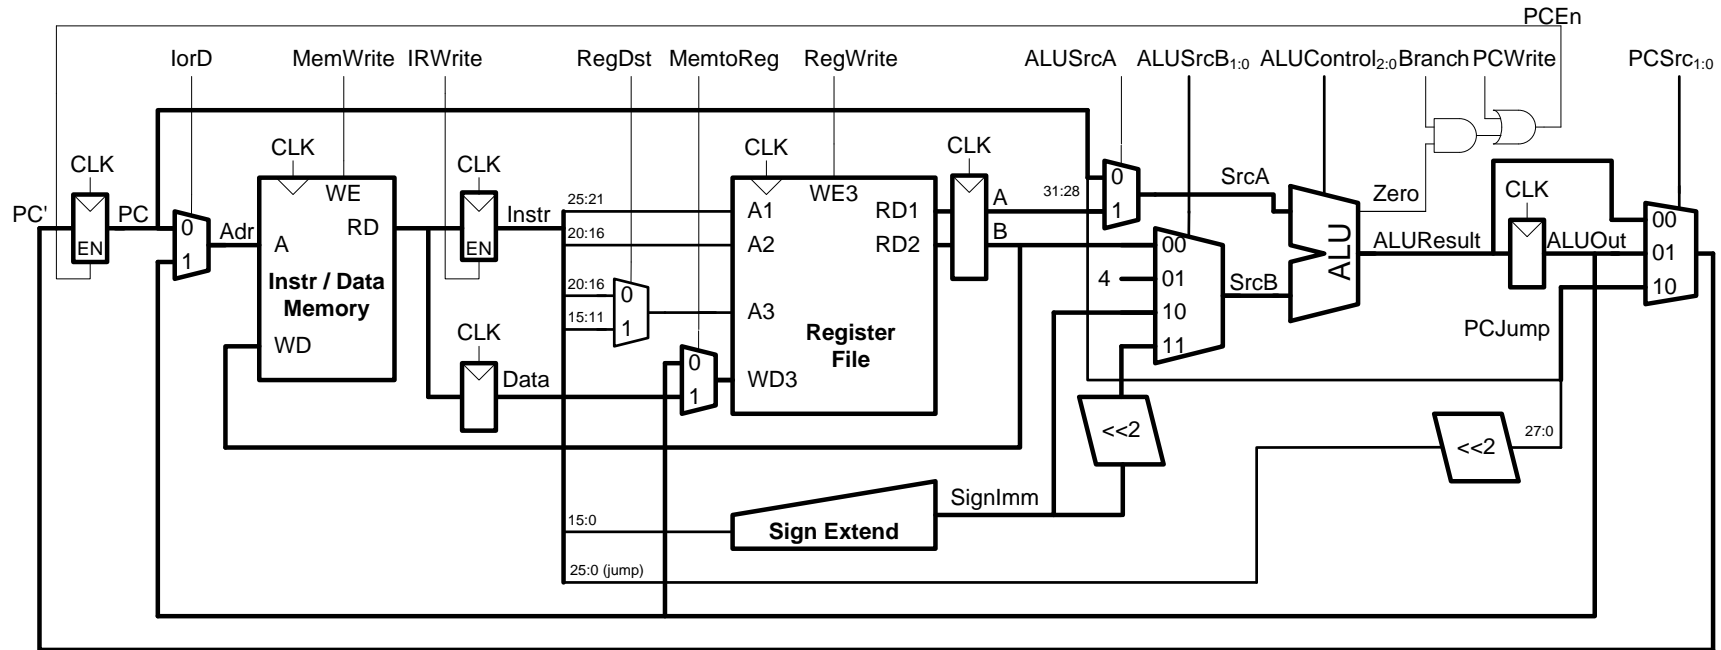
Main Controller FSM: addi



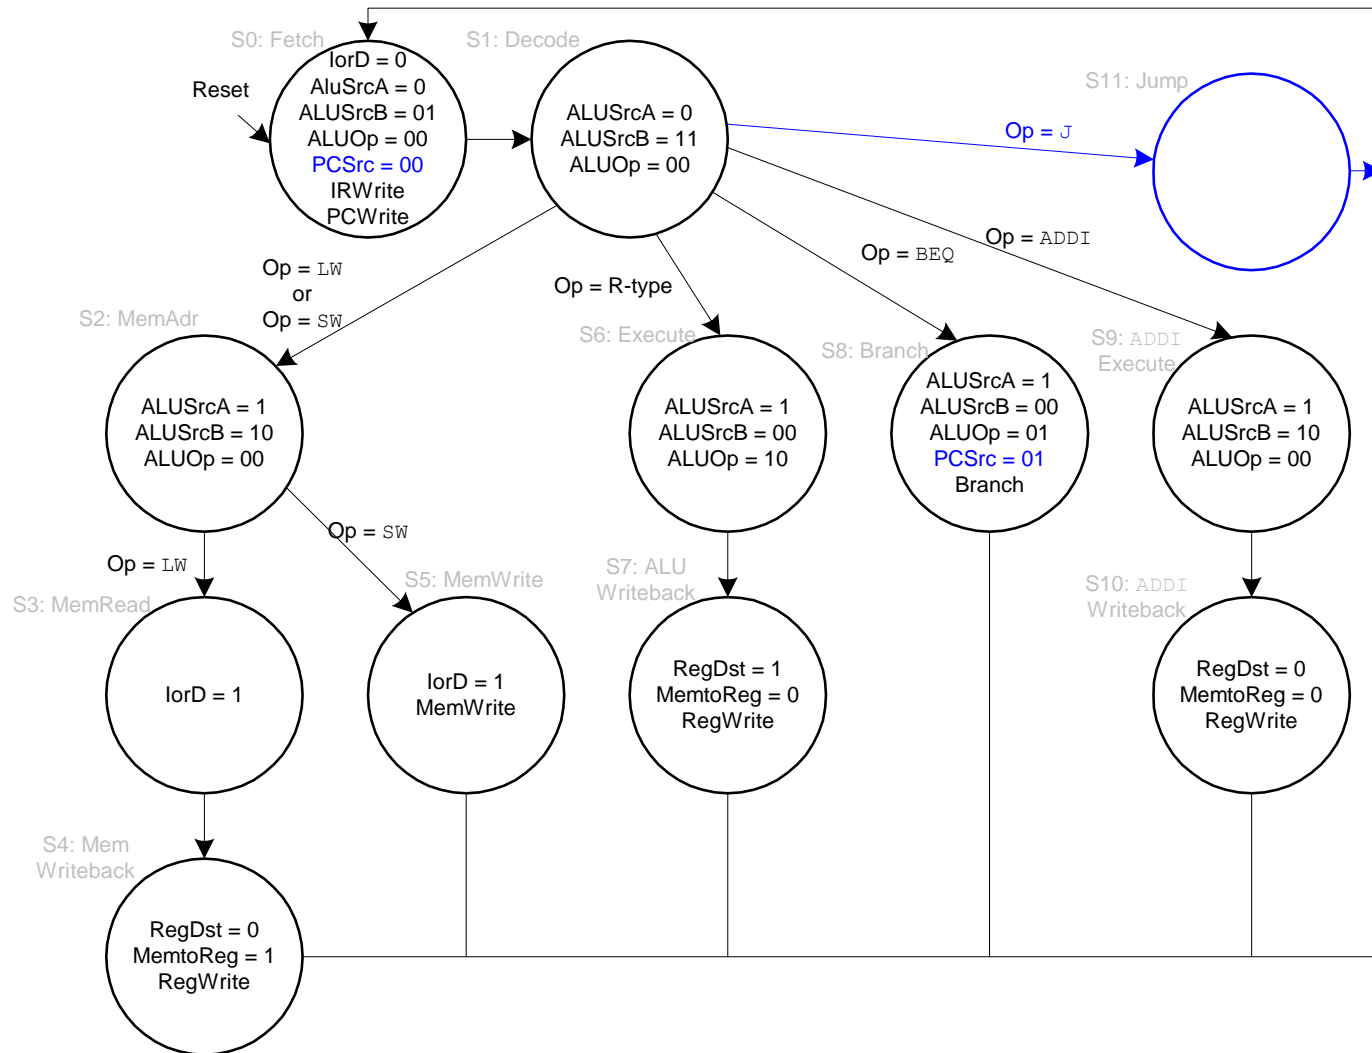
Main Controller FSM: addi



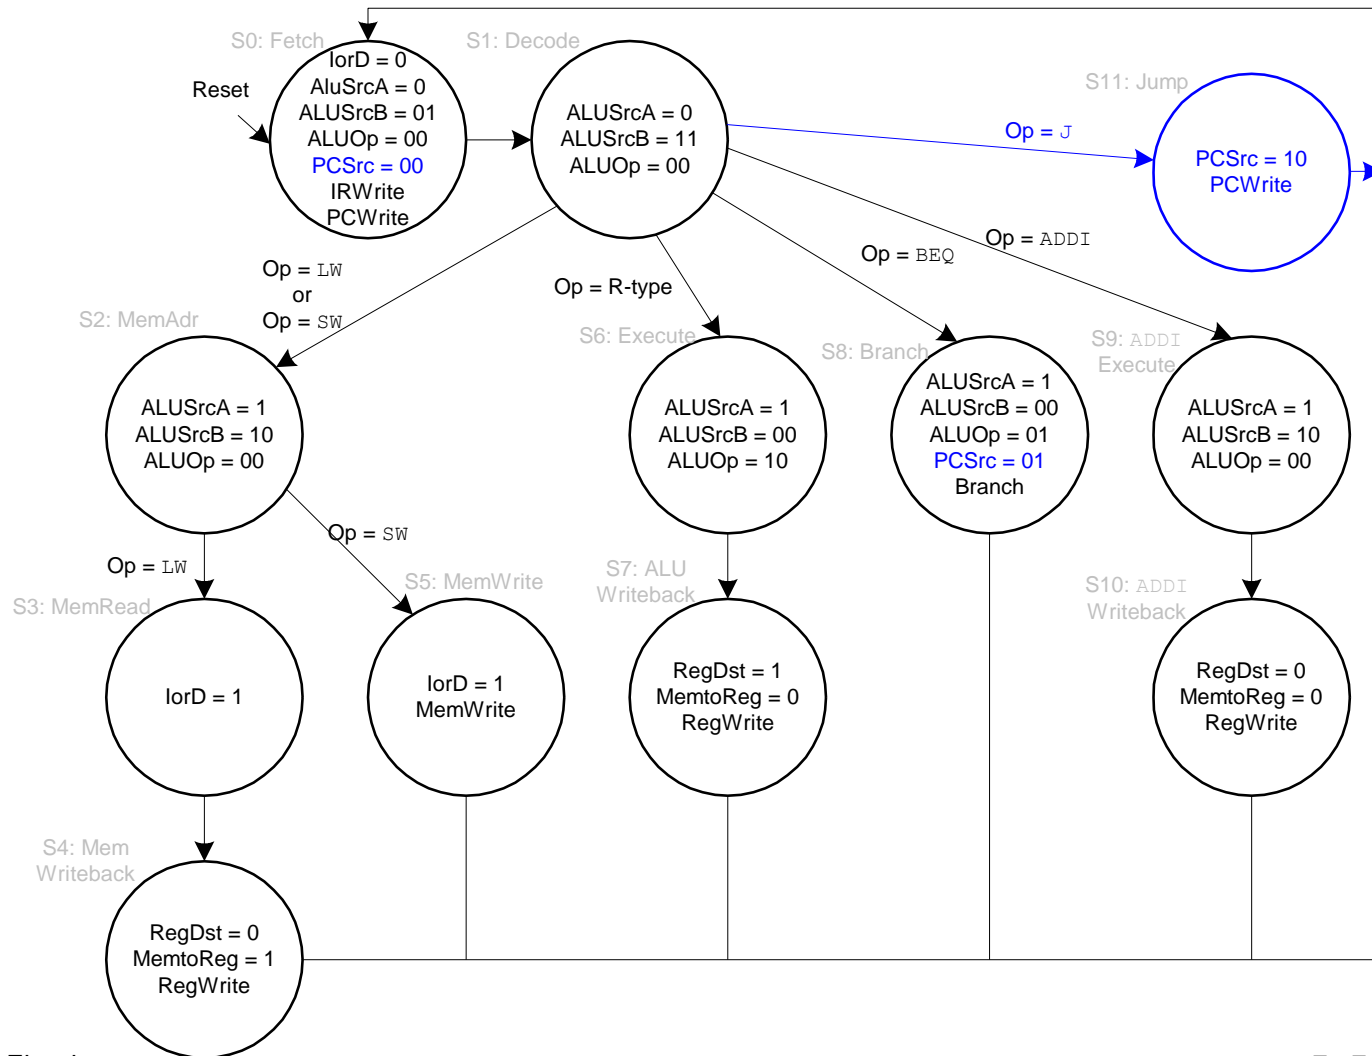
Extended Functionality: j



Control FSM: j



Control FSM: j



Multicycle Performance

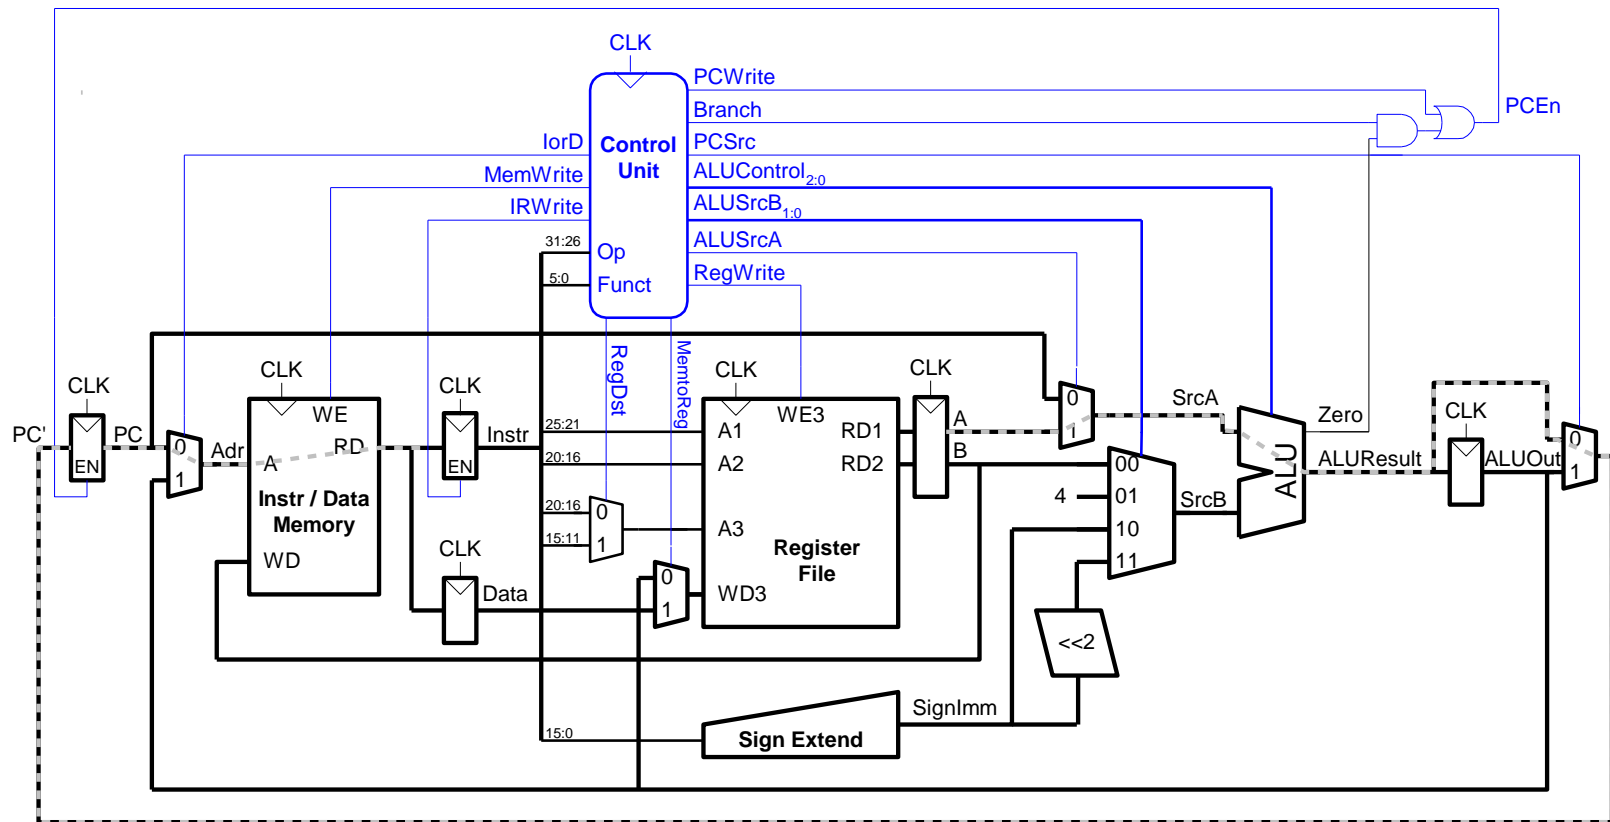
- Instructions take different number of cycles:
 - 3 cycles: beq, j
 - 4 cycles: R-Type, sw, addi
 - 5 cycles: lw
- CPI is weighted average
- SPECINT2000 benchmark:
 - 25% loads
 - 10% stores
 - 11% branches
 - 2% jumps
 - 52% R-type

$$\text{Average CPI} = (0.11 + 0.02)(3) + (0.52 + 0.10)(4) + (0.25)(5) = 4.12$$

Multicycle Performance

- Multicycle critical path:

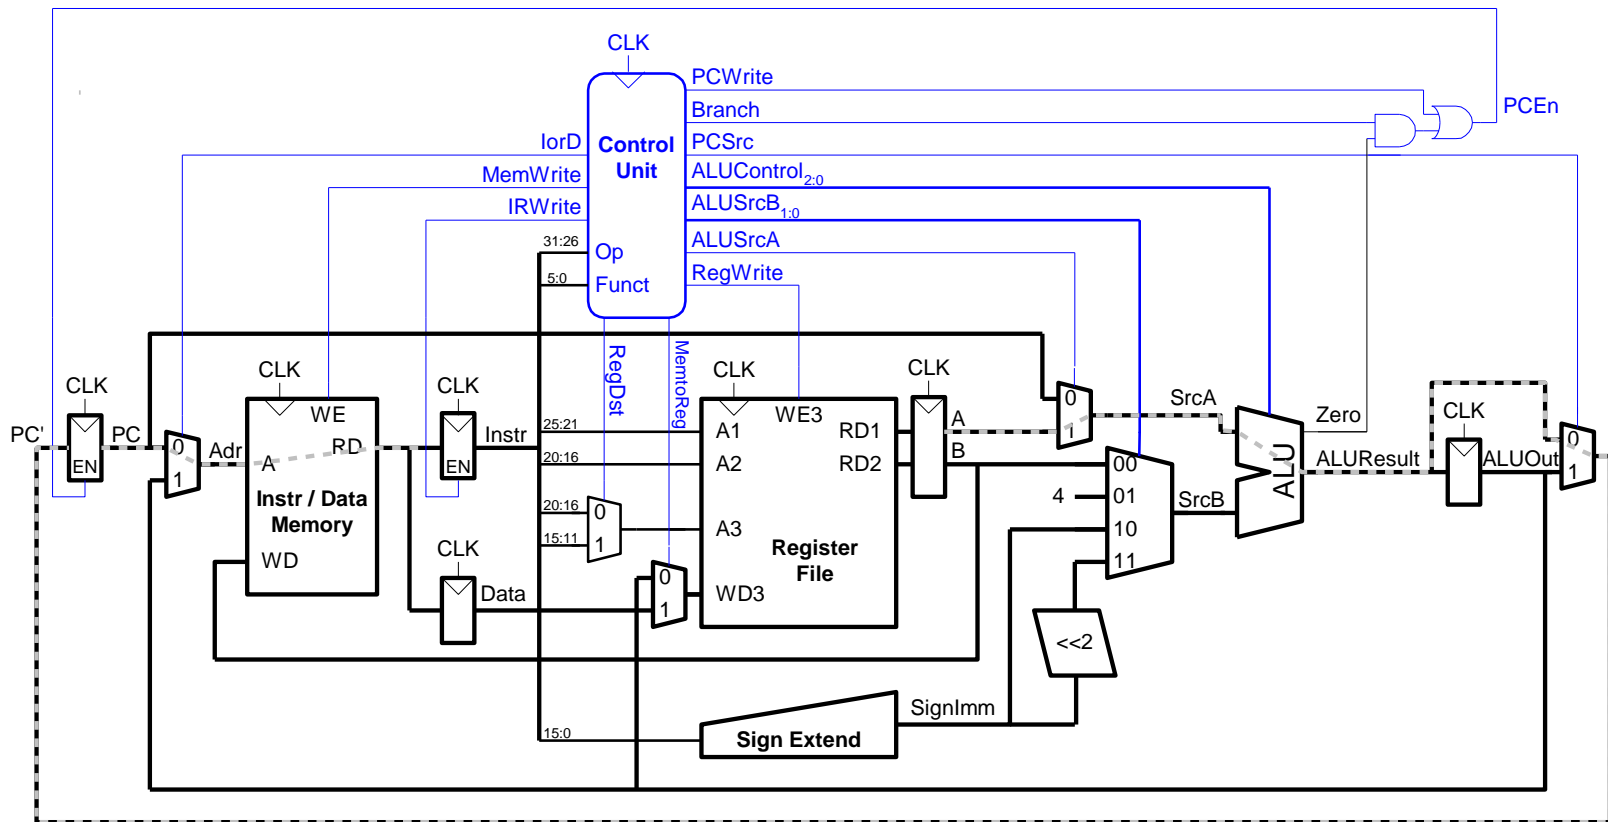
$$T_c =$$



Multicycle Performance

- Multicycle critical path:

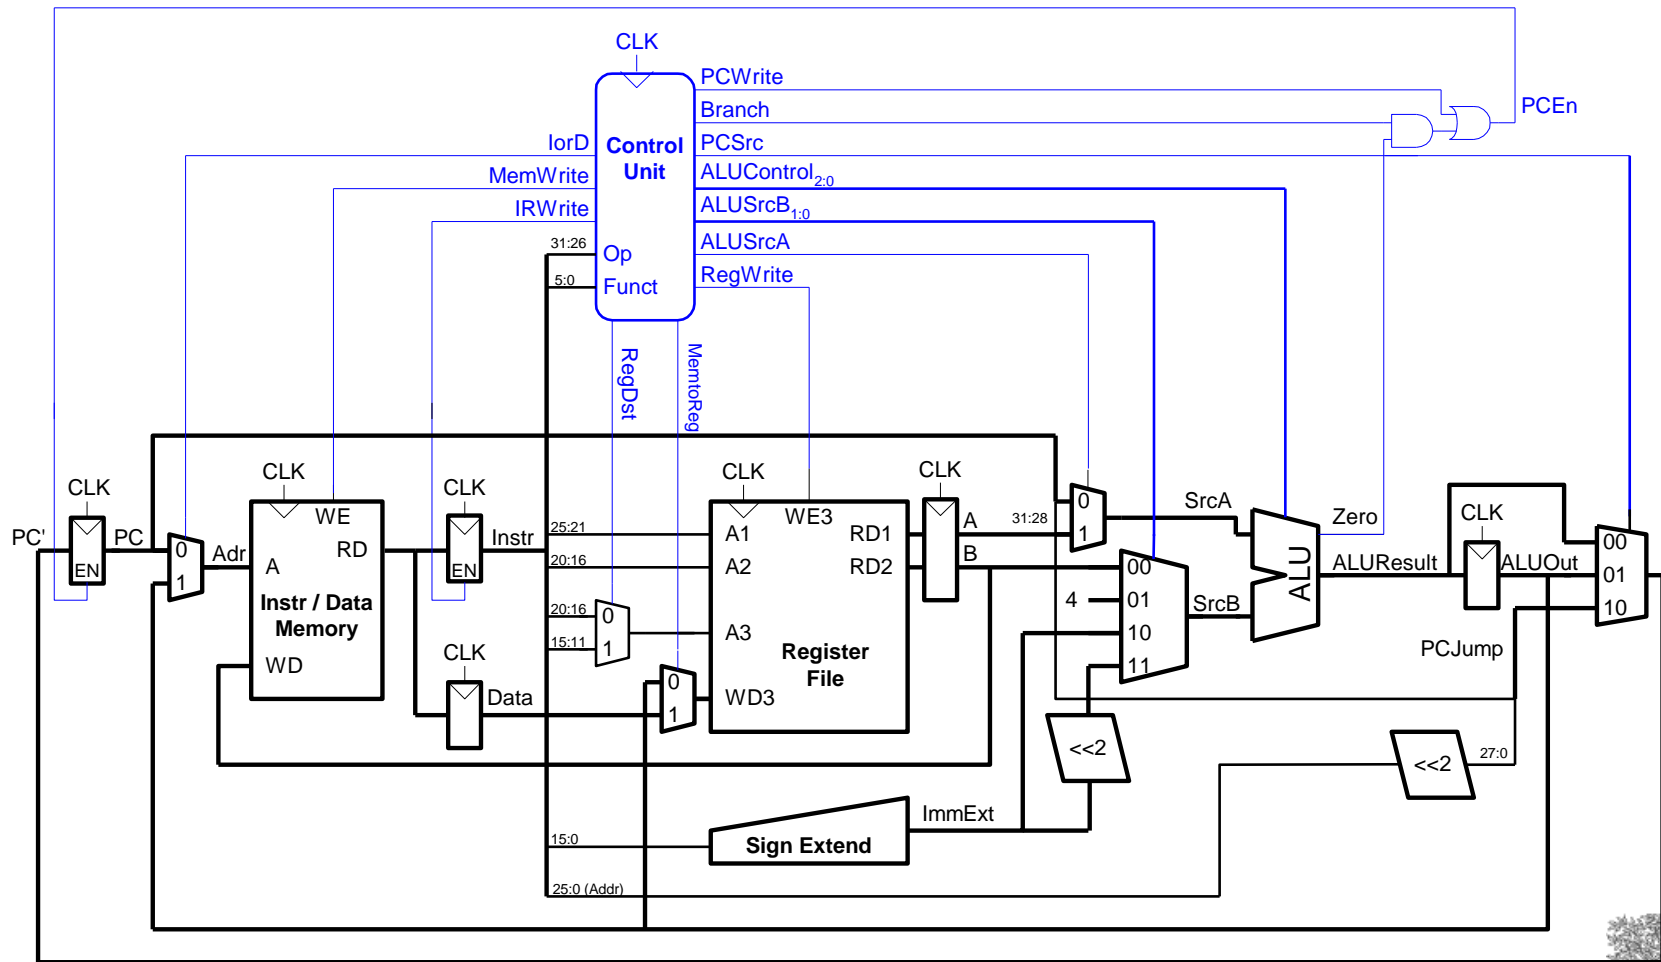
$$T_c = t_{pcq} + t_{mux} + \max(t_{ALU} + t_{mux}, t_{mem}) + t_{setup}$$



the 1990s, the number of people in the world who are under 15 years of age is expected to increase by 1.5 billion. This increase is expected to be concentrated in the developing countries, where the population is expected to increase by 2.5 billion.



Review: Multicycle MIPS Processor

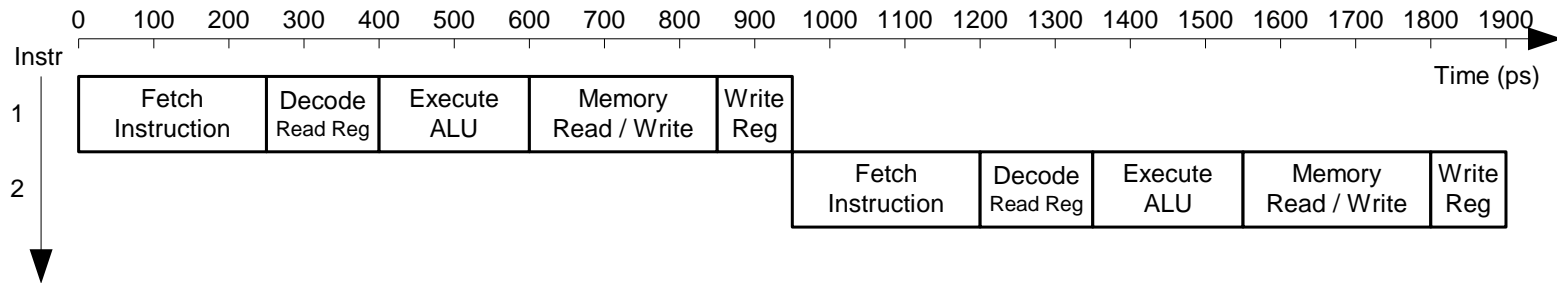


Pipelined MIPS Processor

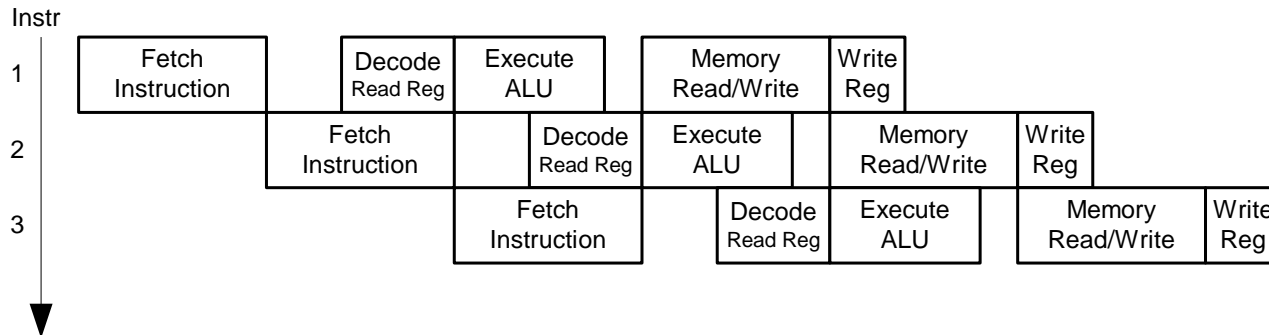
- Temporal parallelism
- Divide single-cycle processor into 5 stages:
 - Fetch
 - Decode
 - Execute
 - Memory
 - Writeback
- Add pipeline registers between stages

Single-Cycle vs. Pipelined Performance

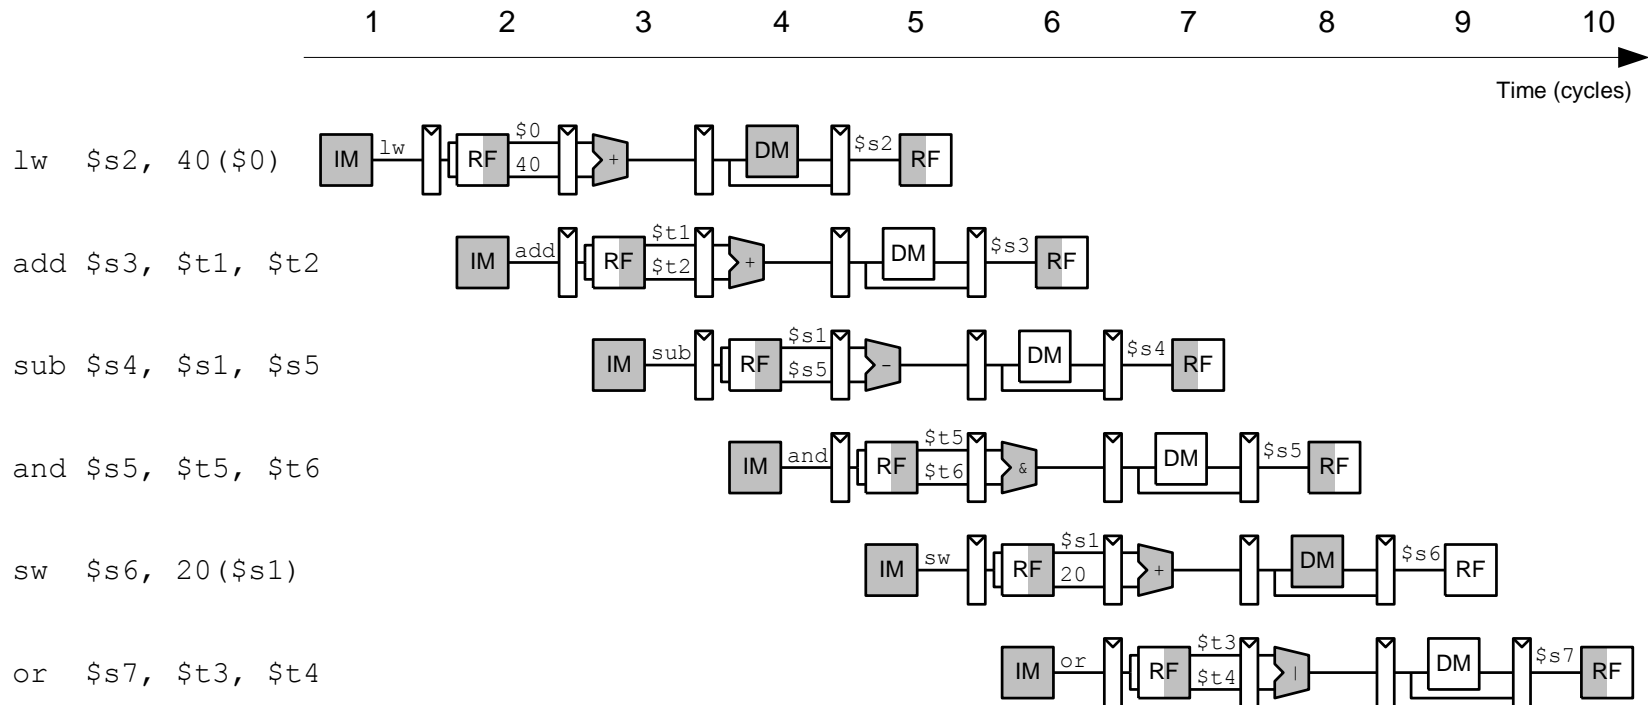
Single-Cycle



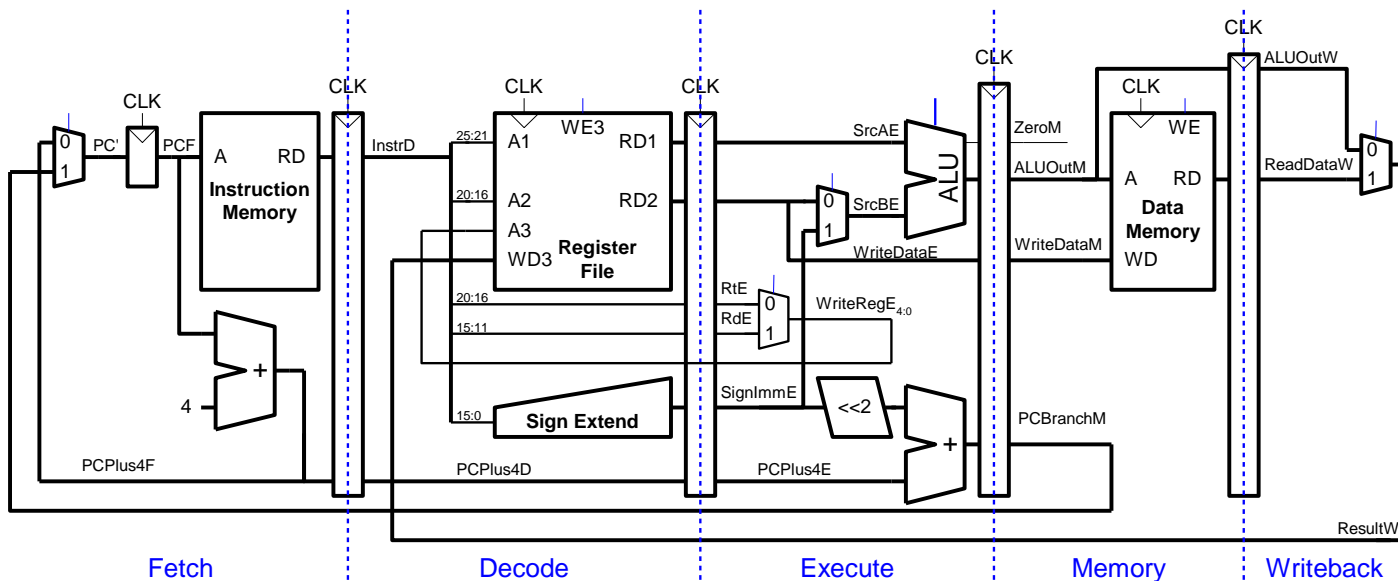
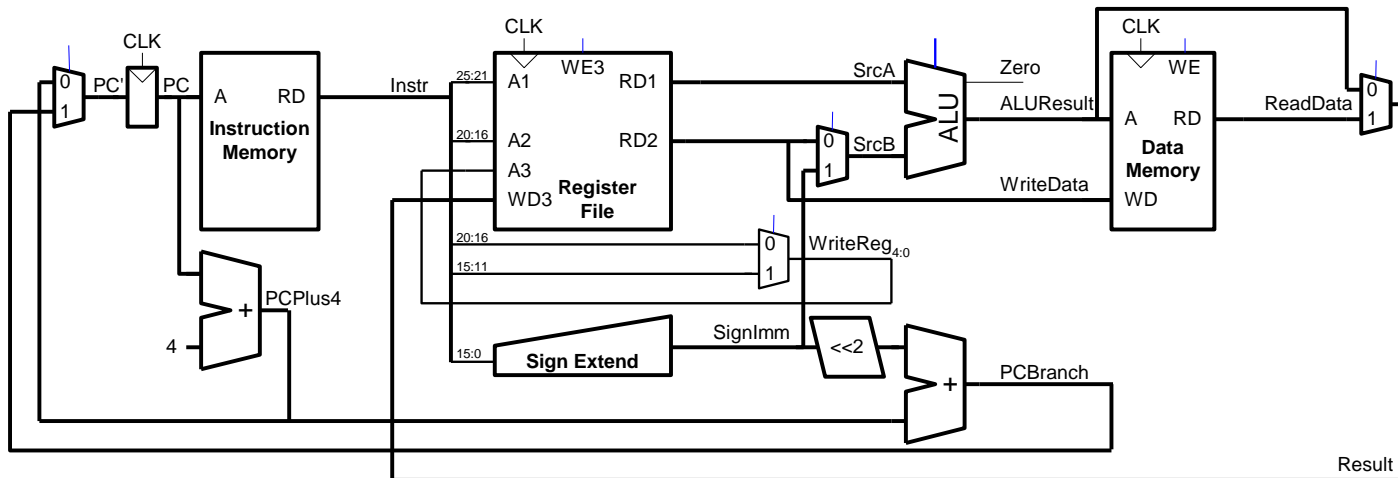
Pipelined



Pipelining Abstraction

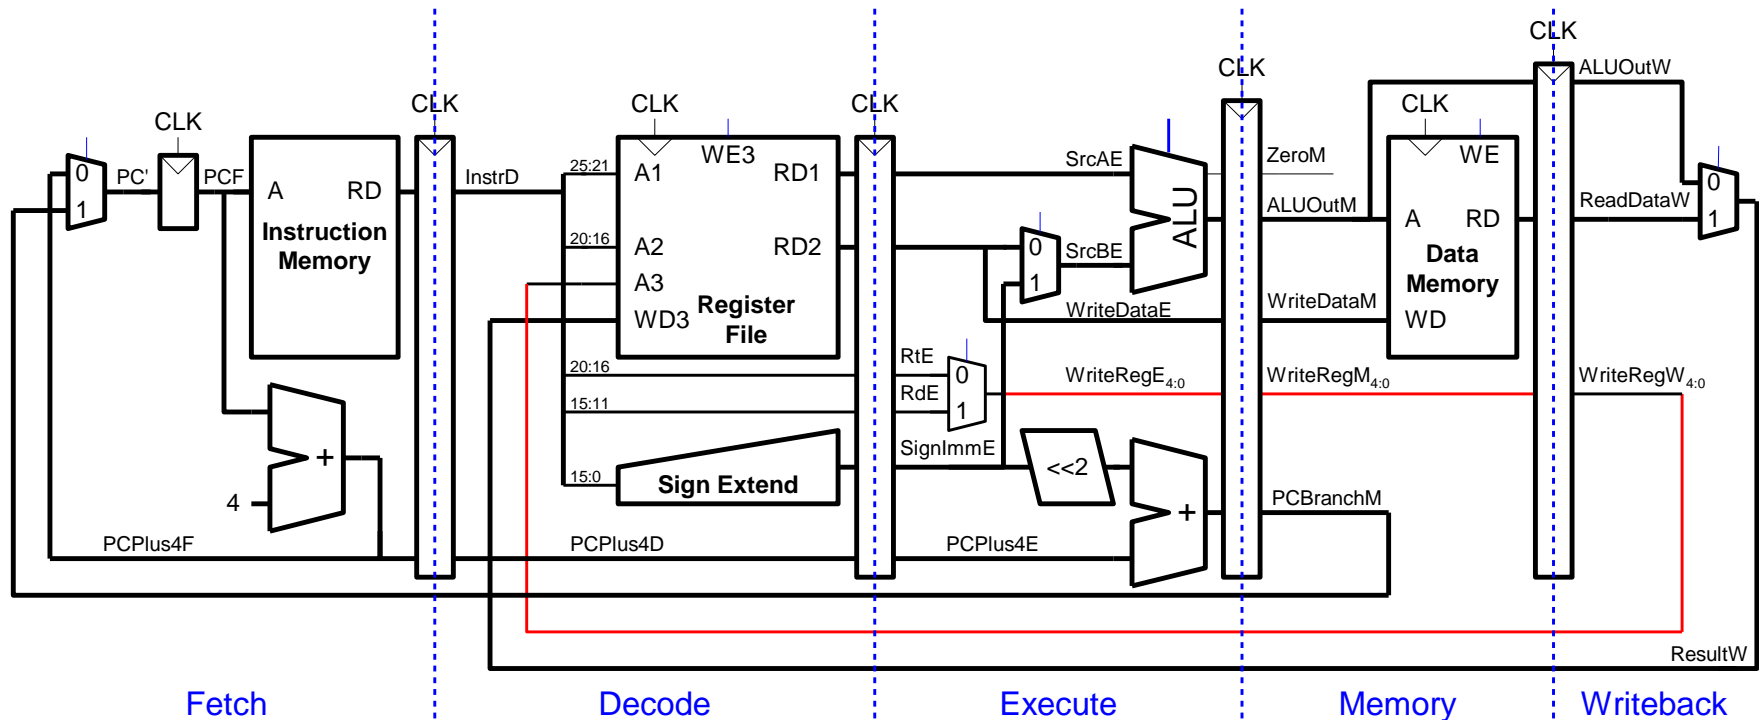


Single-Cycle and Pipelined Datapath

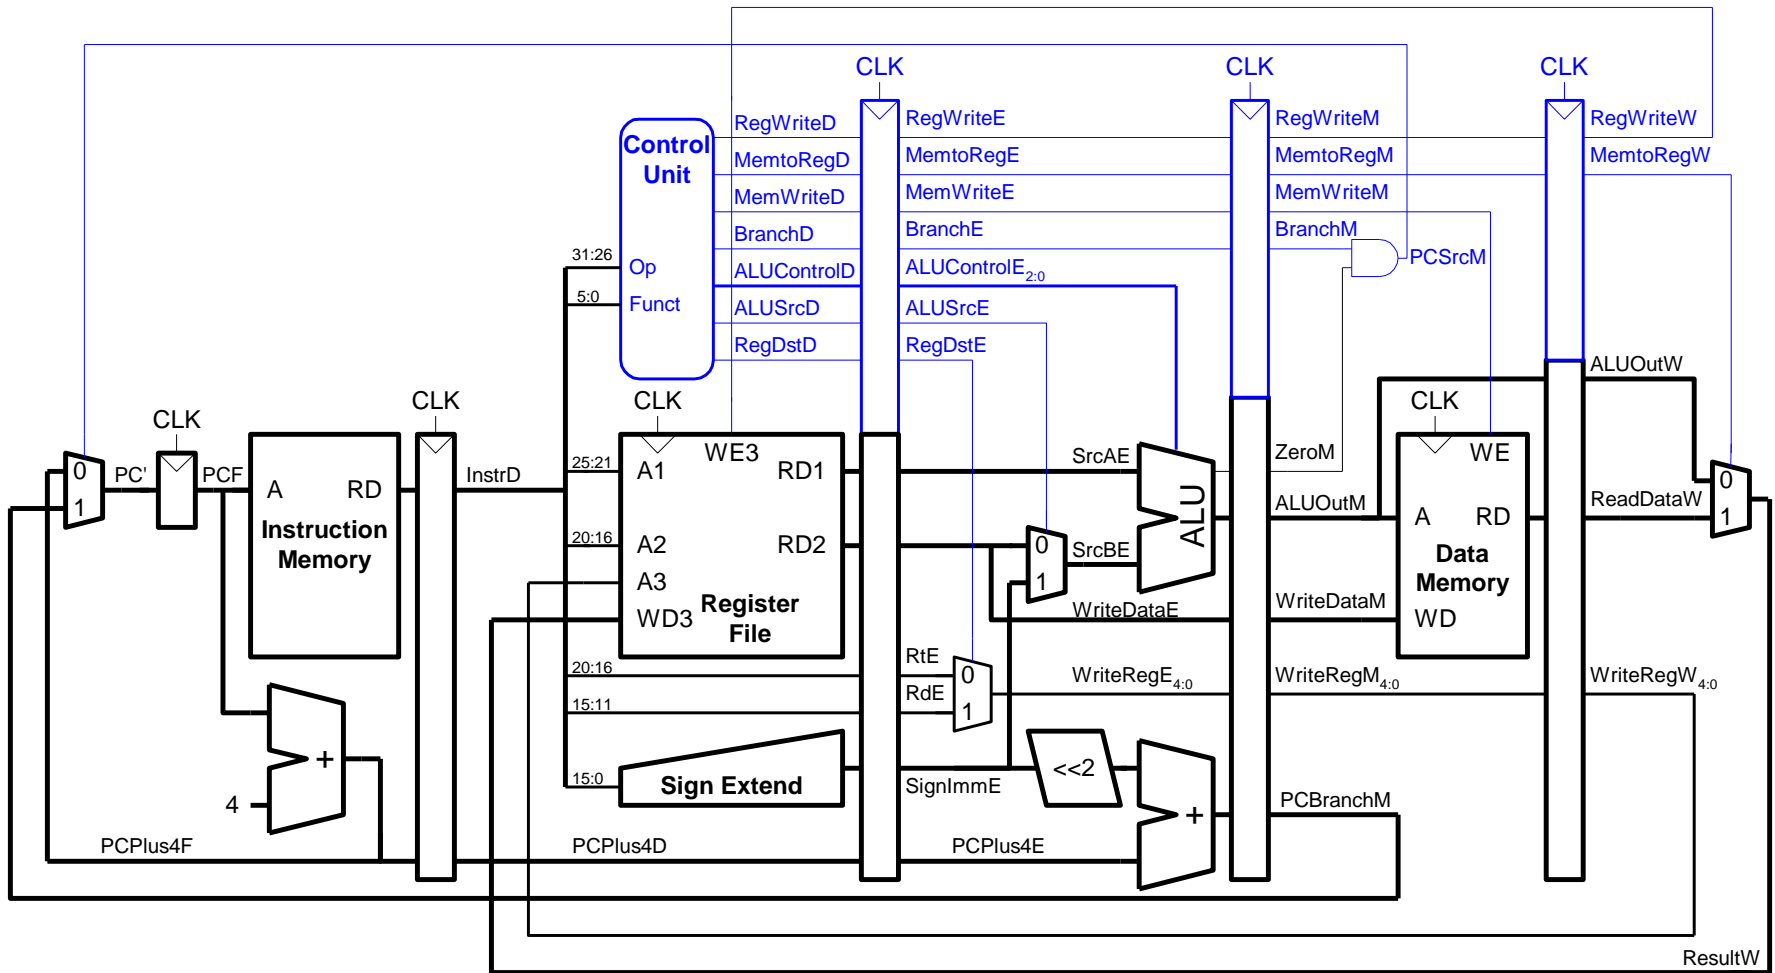


Corrected Pipelined Datapath

- WriteReg must arrive at the same time as Result



Pipelined Control



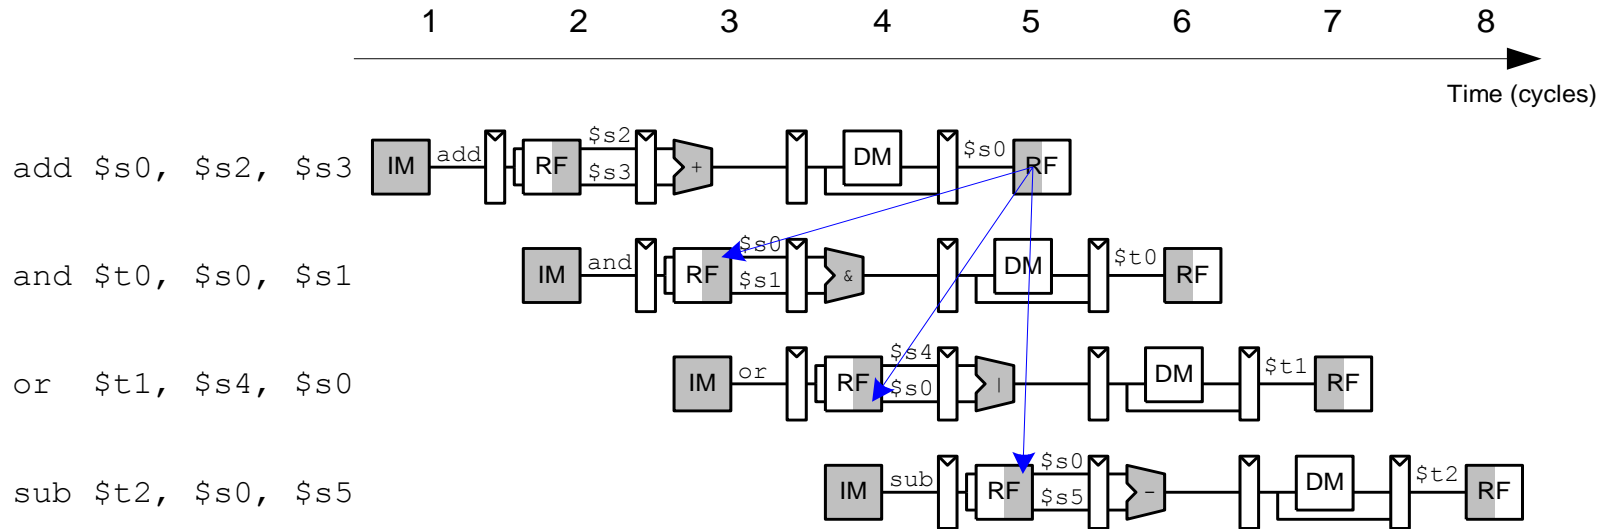
Same control unit as single-cycle processor

Control delayed to proper pipeline stage

Pipeline Hazard

- Occurs when an instruction depends on results from previous instruction that hasn't completed.
- Types of hazards:
 - **Data hazard:** register value not written back to register file yet
 - **Control hazard:** next instruction not decided yet (caused by branches)

Data Hazard

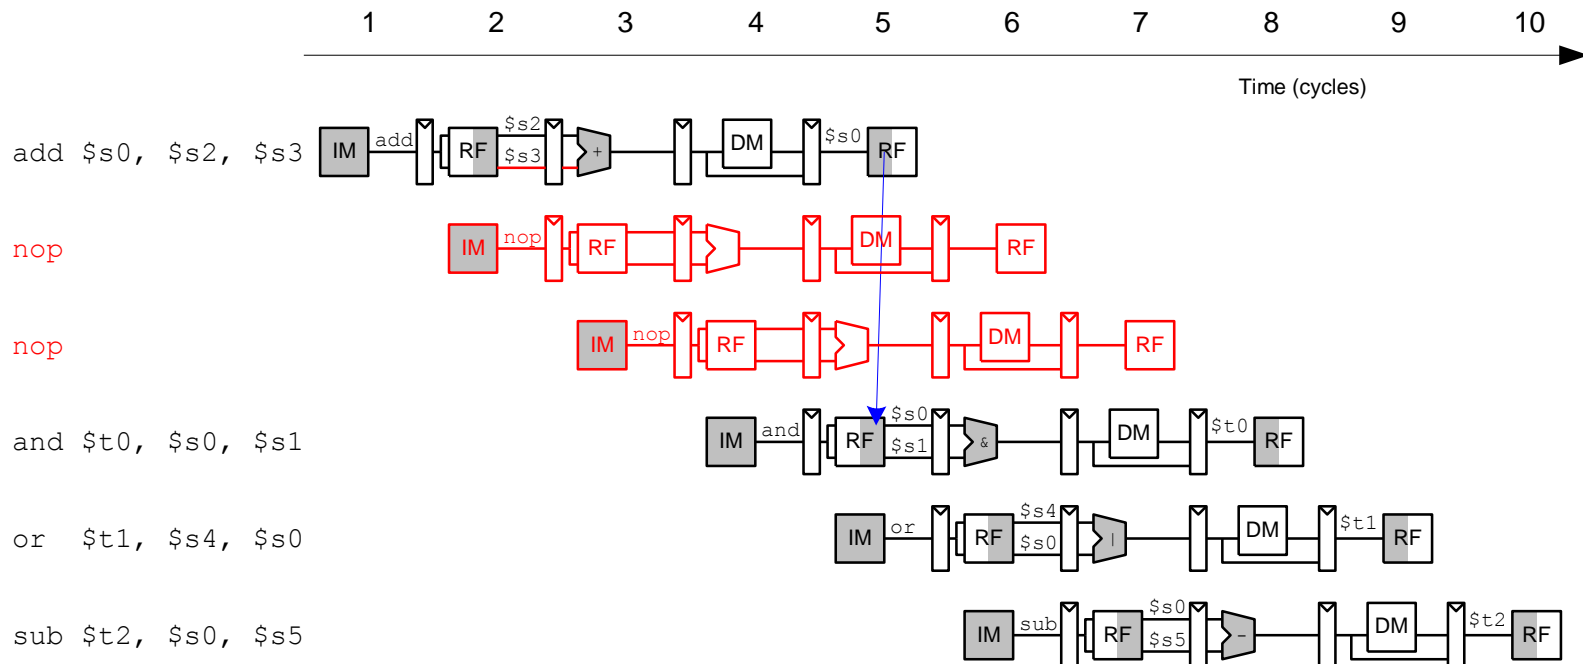


Handling Data Hazards

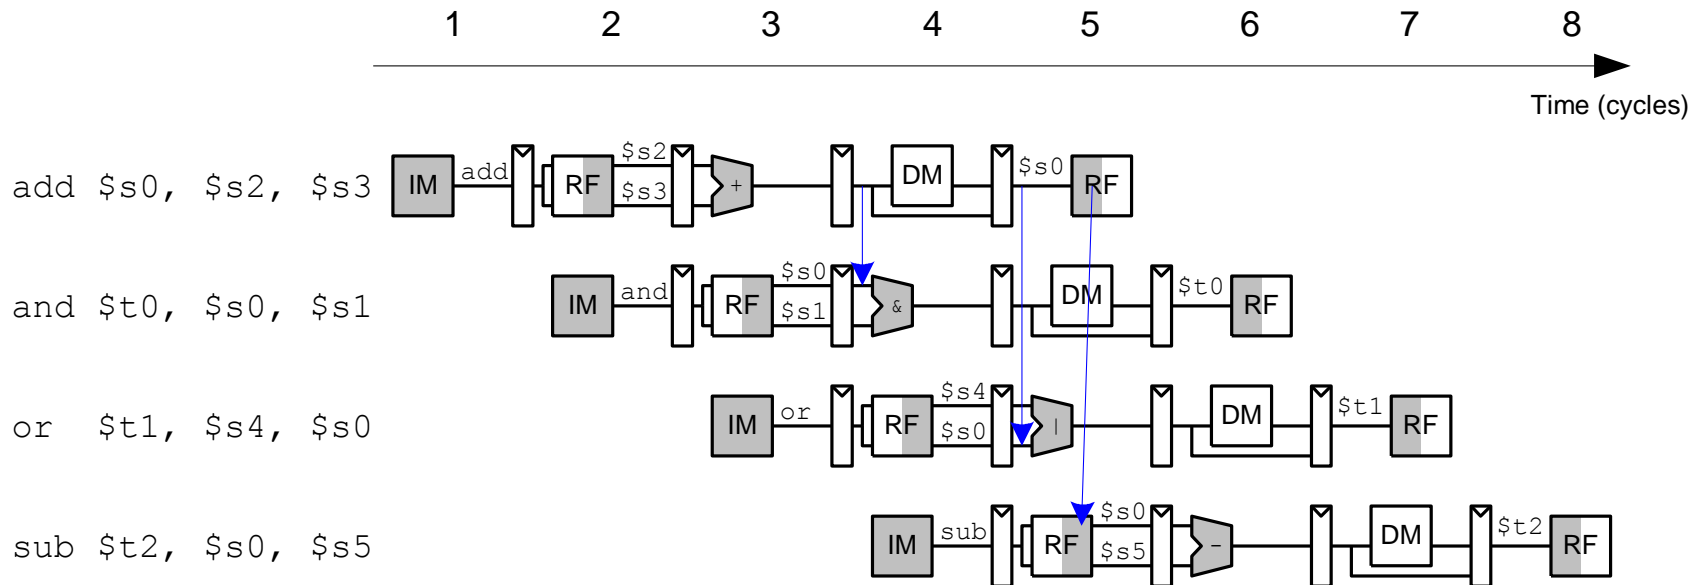
- Insert nops in code at compile time
- Rearrange code at compile time
- Forward data at run time
- Stall the processor at run time

Compile-Time Hazard Elimination

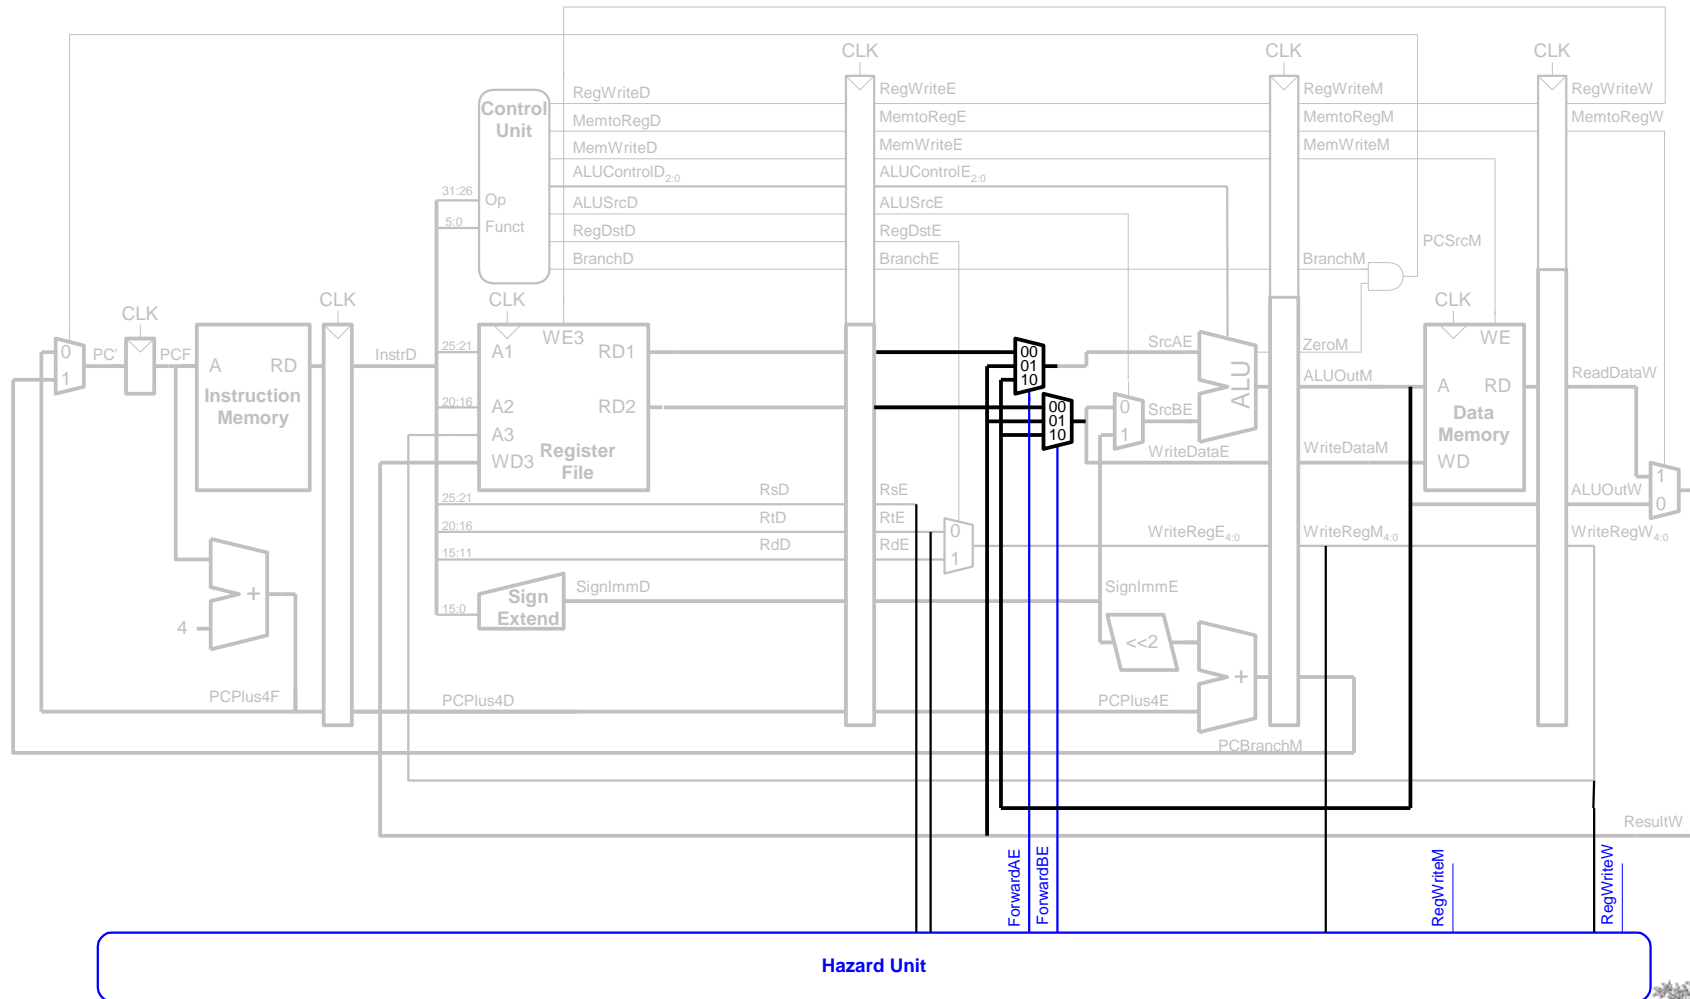
- Insert enough nops for result to be ready
- Or move independent useful instructions forward



Data Forwarding



Data Forwarding



Data Forwarding

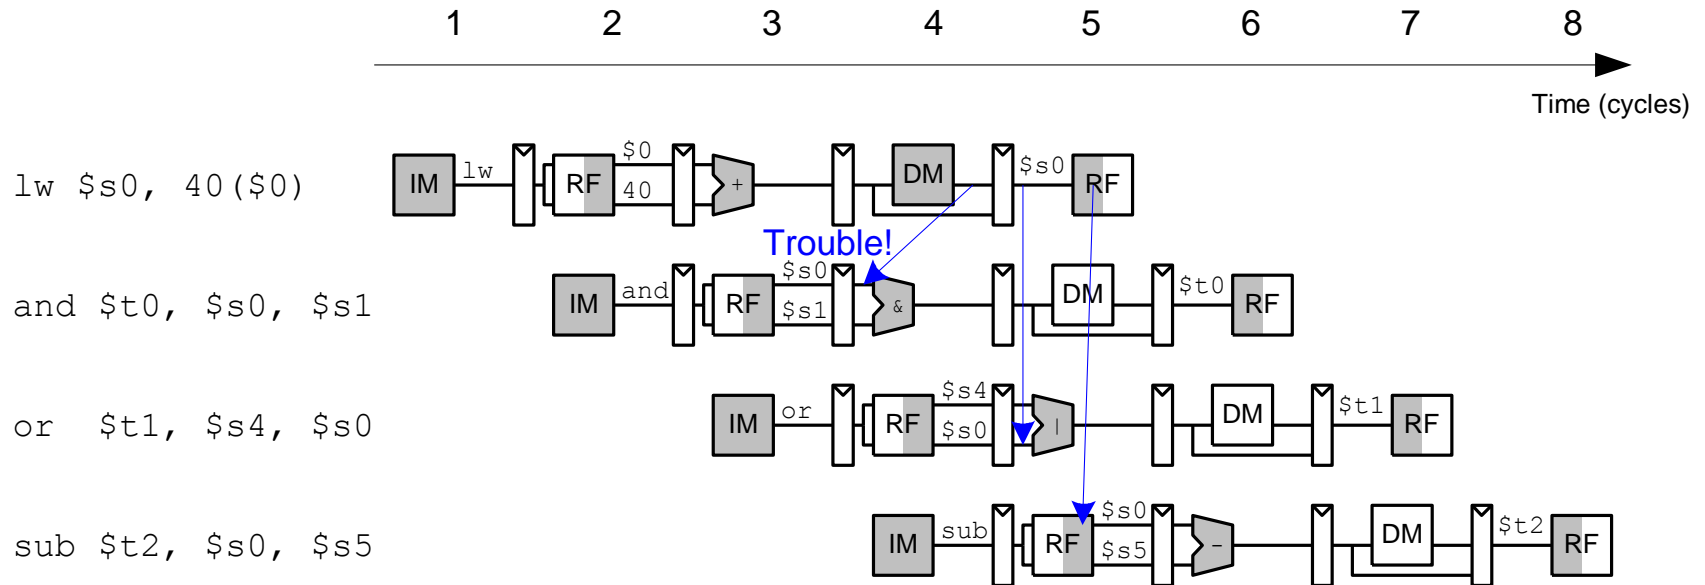
- Forward to Execute stage from either:
 - Memory stage or
 - Writeback stage

- Forwarding logic for *ForwardAE*:

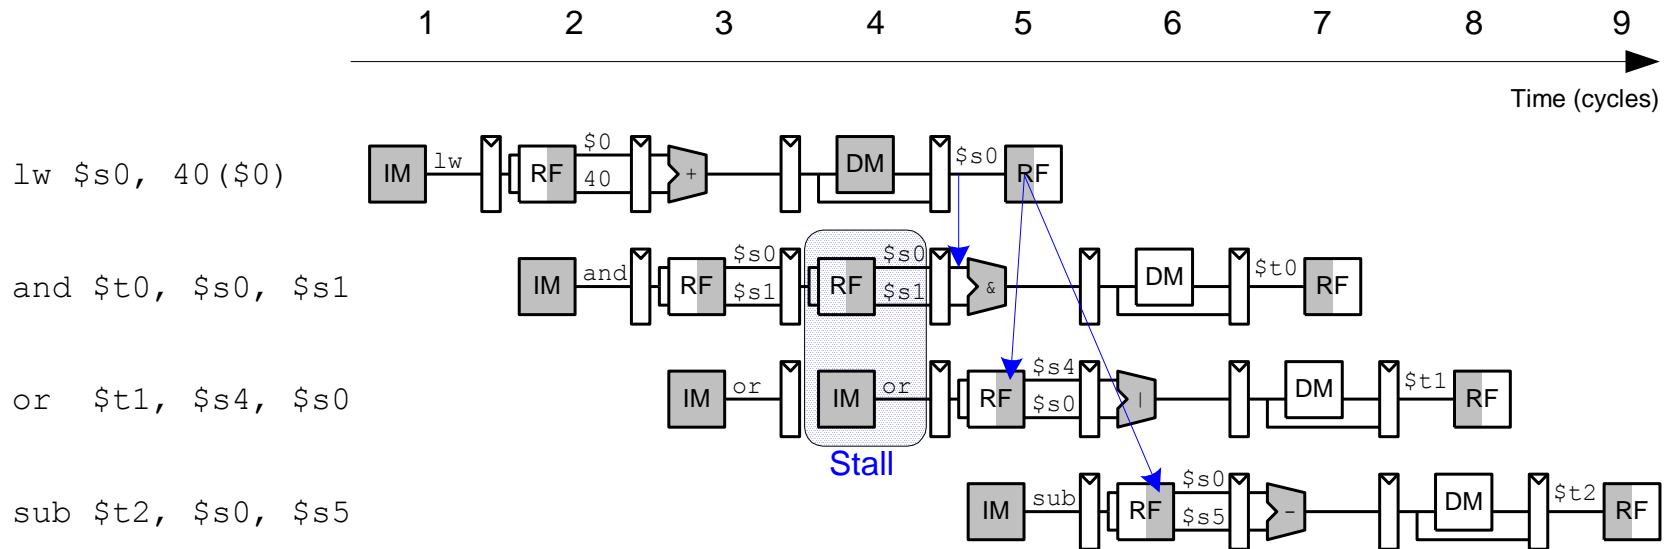
```
if      ((rsE != 0) AND (rsE == WriteRegM) AND RegWriteM)
then    ForwardAE = 10
else if ((rsE != 0) AND (rsE == WriteRegW) AND RegWriteW)
then    ForwardAE = 01
else    ForwardAE = 00
```

- Forwarding logic for *ForwardBE* same, but replace *rsE* with *rtE*

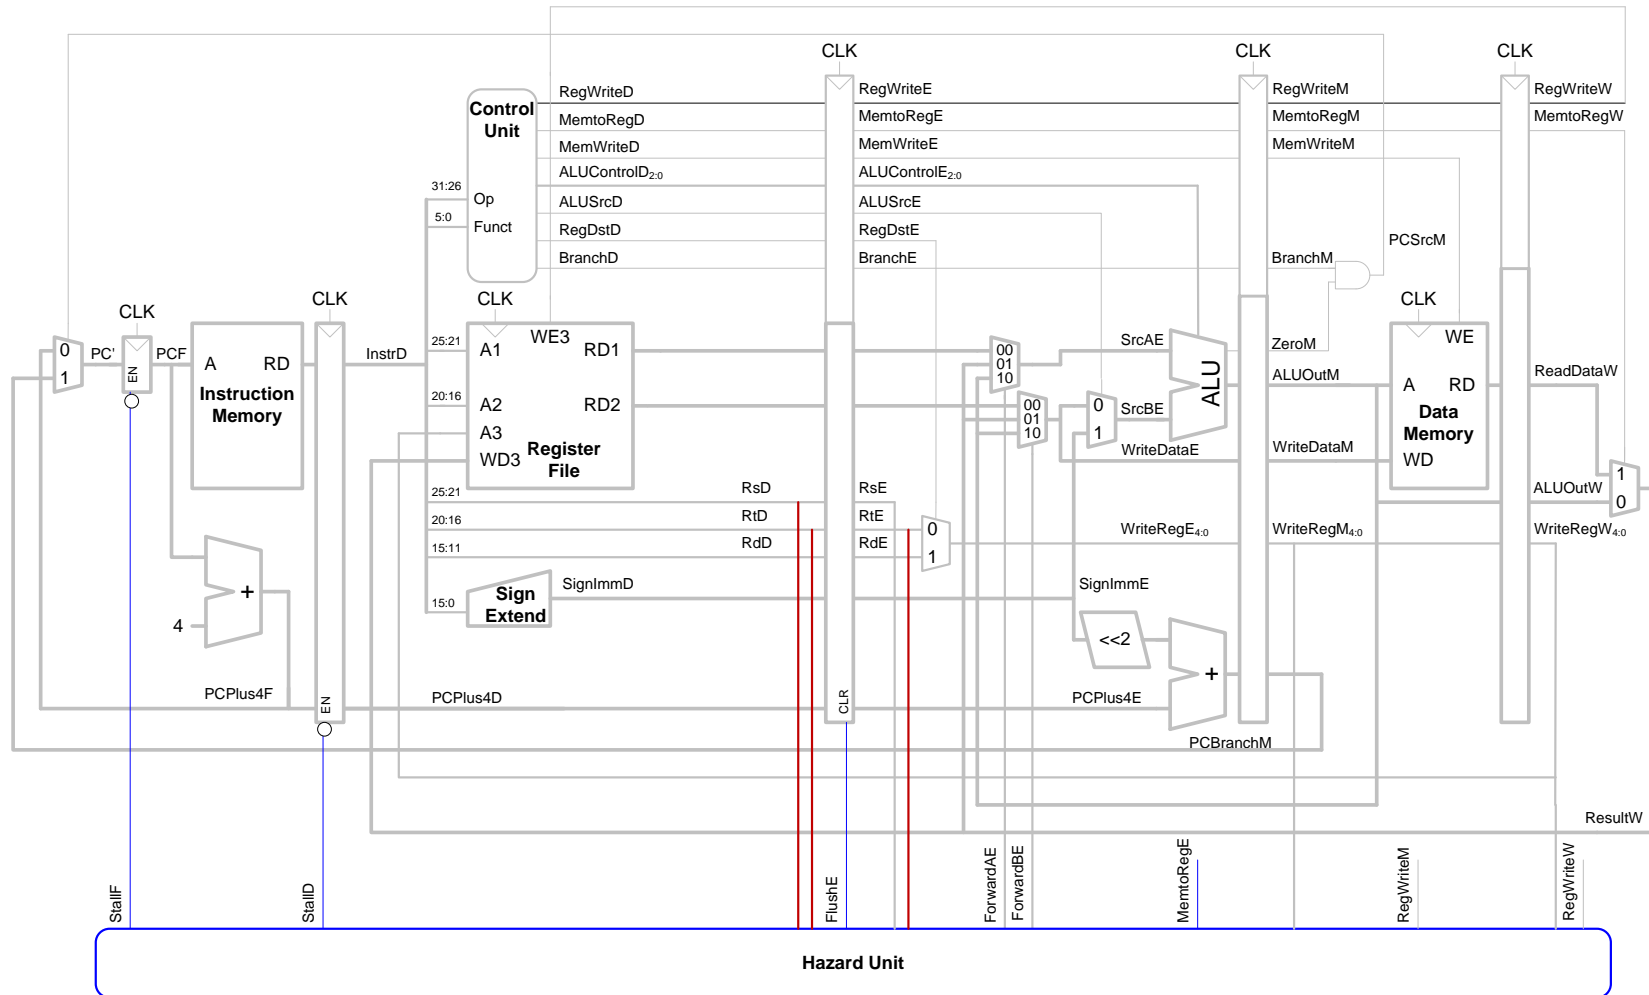
Stalling



Stalling



Stalling Hardware



Stalling Hardware

- Stalling logic:

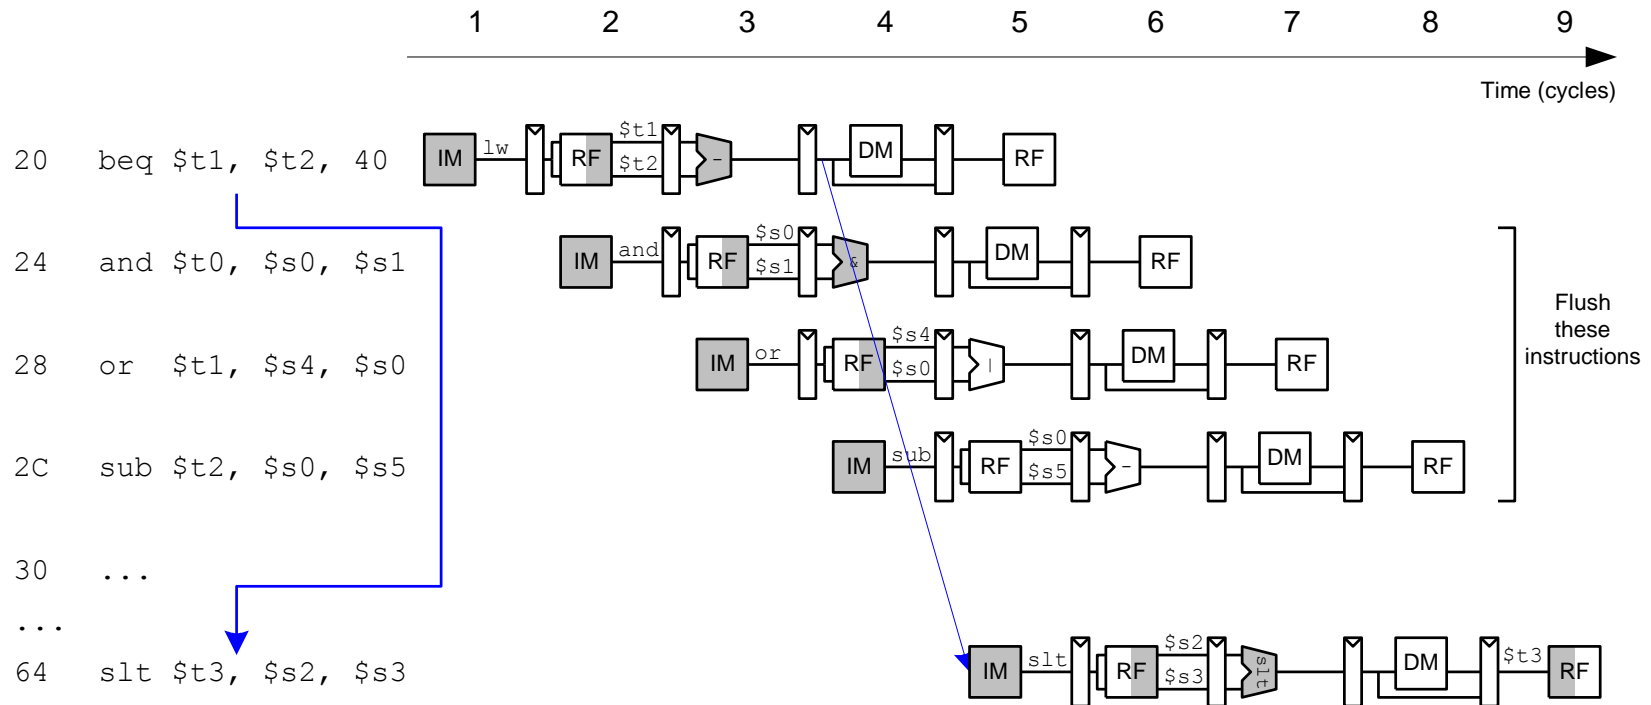
$lwstall = ((rsD == rtE) \text{ OR } (rtD == rtE)) \text{ AND } MemtoRegE$

$StallF = StallD = FlushE = lwstall$

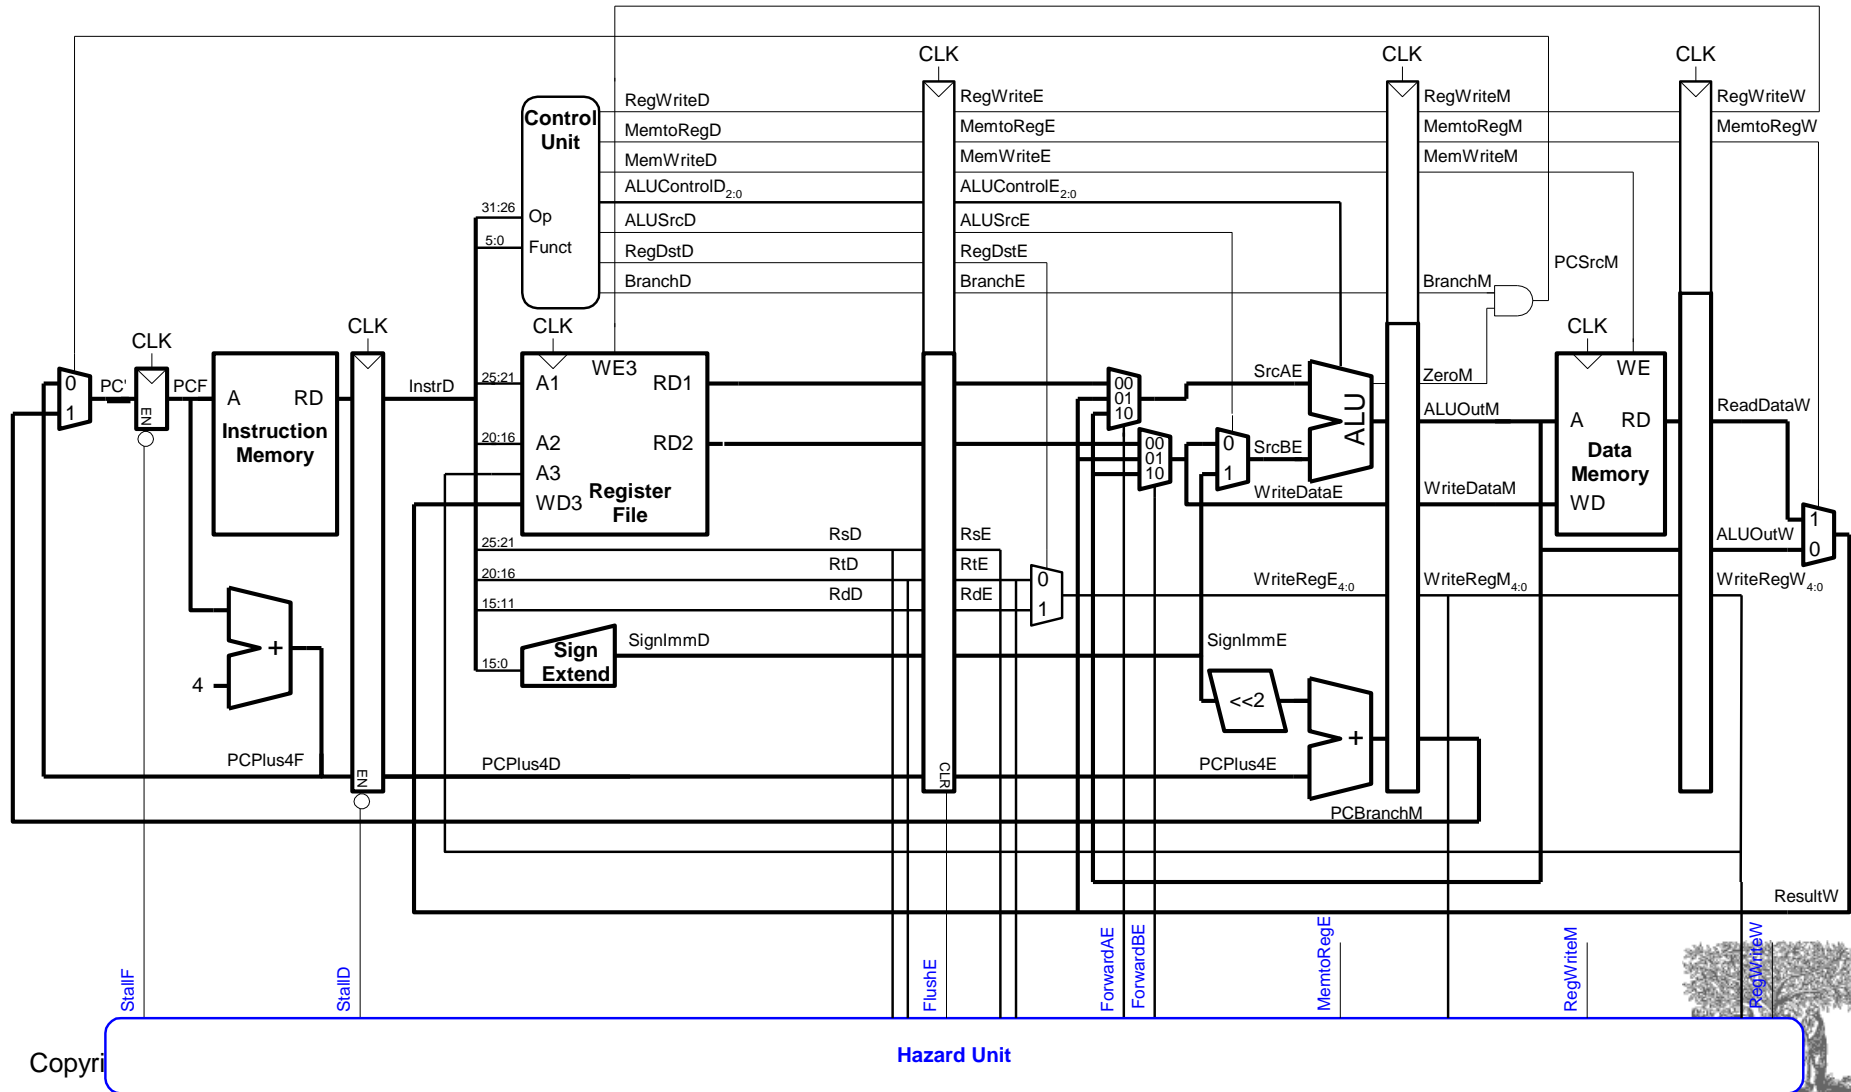
Control Hazards

- beq:
 - branch is not determined until the fourth stage of the pipeline
 - Instructions after the branch are fetched before branch occurs
 - These instructions must be flushed if the branch happens
- Branch misprediction penalty
 - number of instruction flushed when branch is taken
 - May be reduced by determining branch earlier

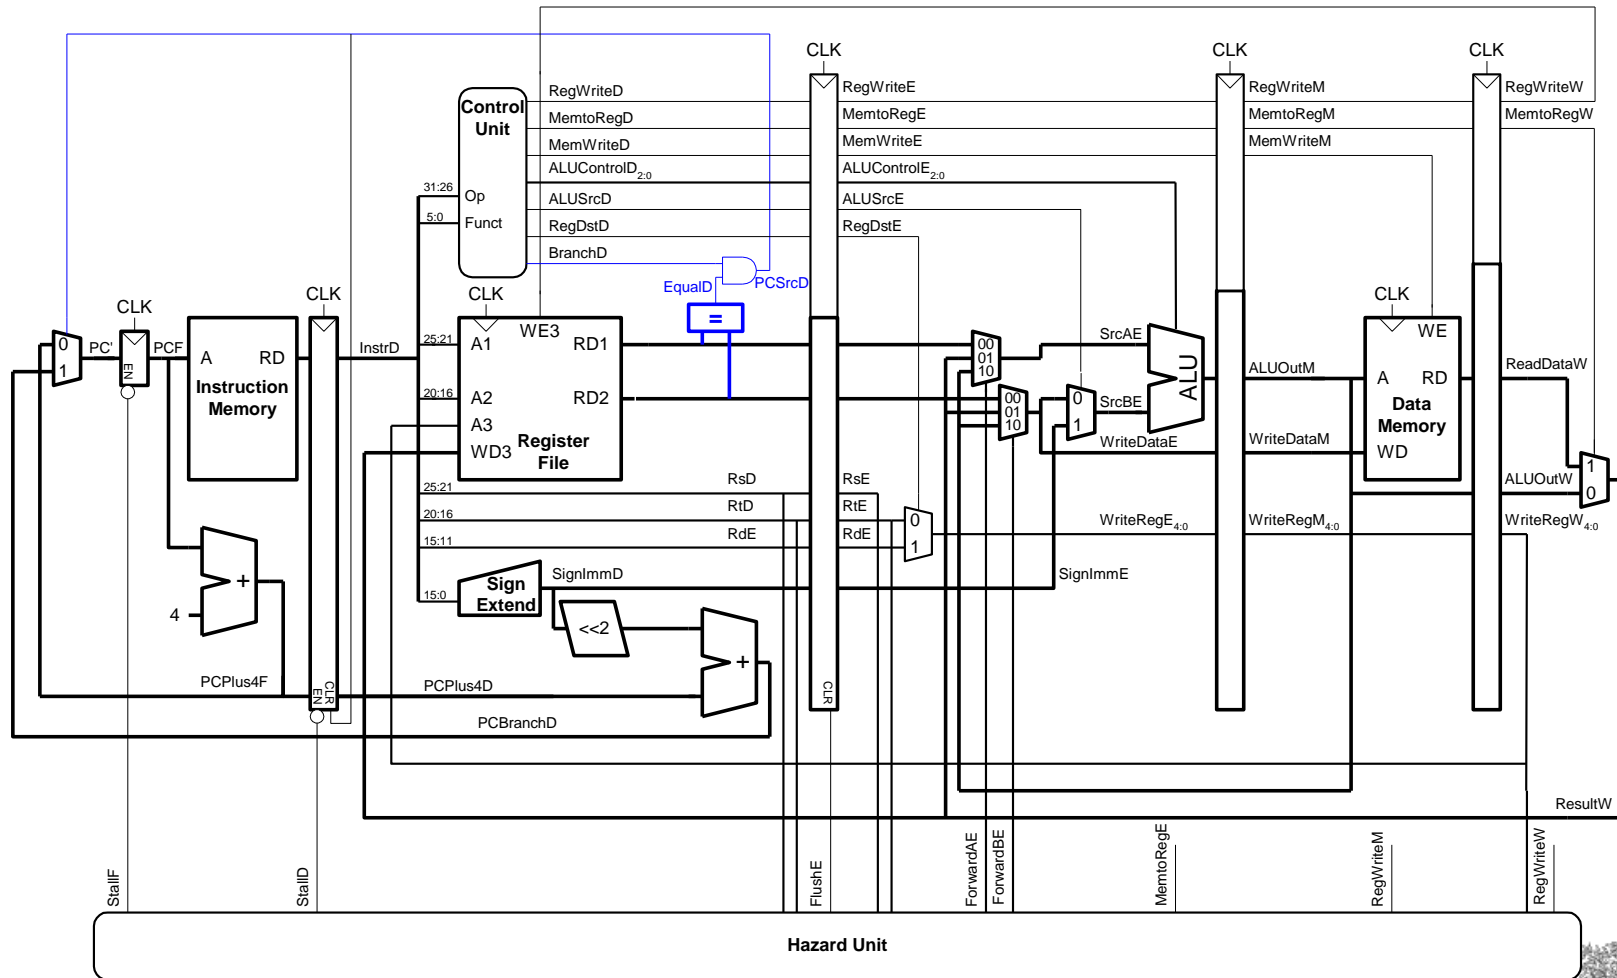
Control Hazards



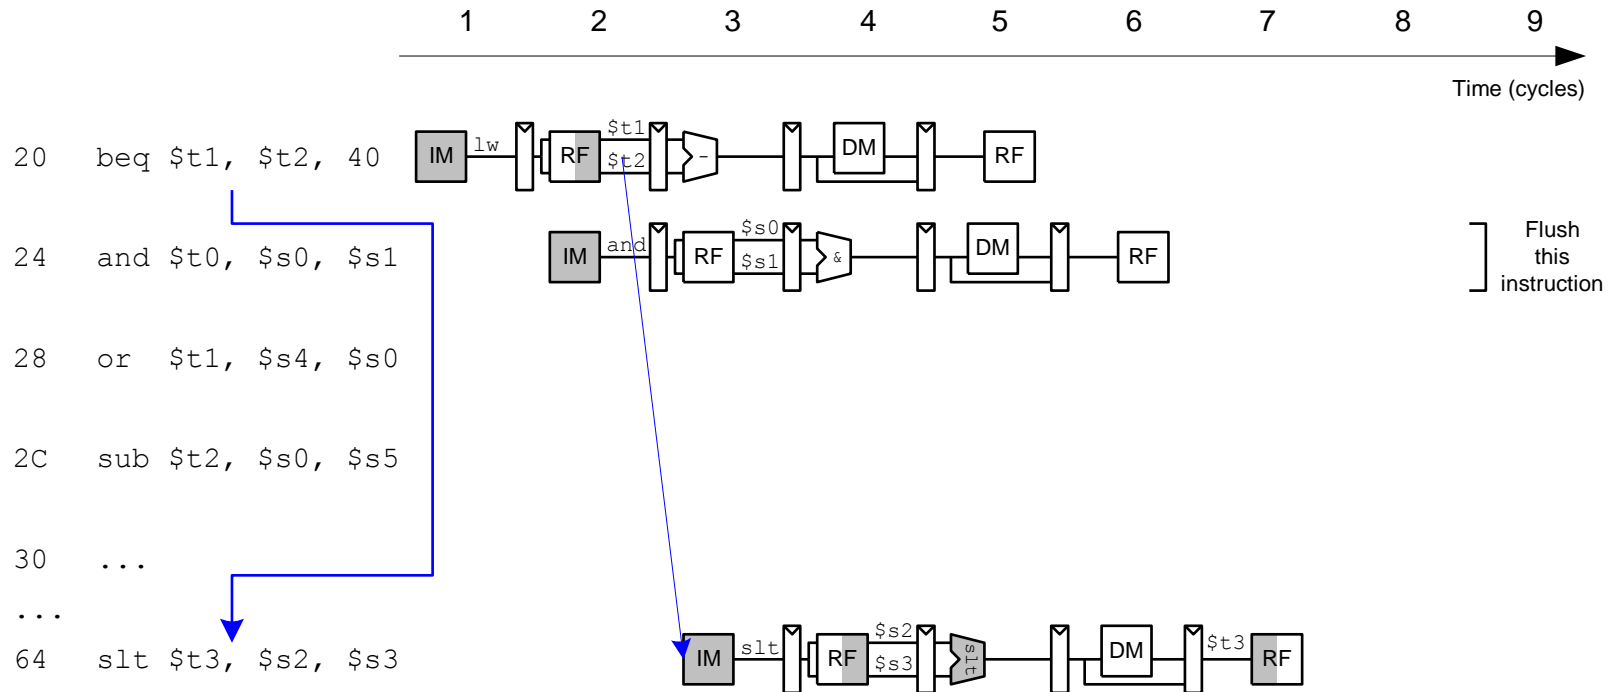
Control Hazards: Original Pipeline



Control Hazards: Early Branch Resolution

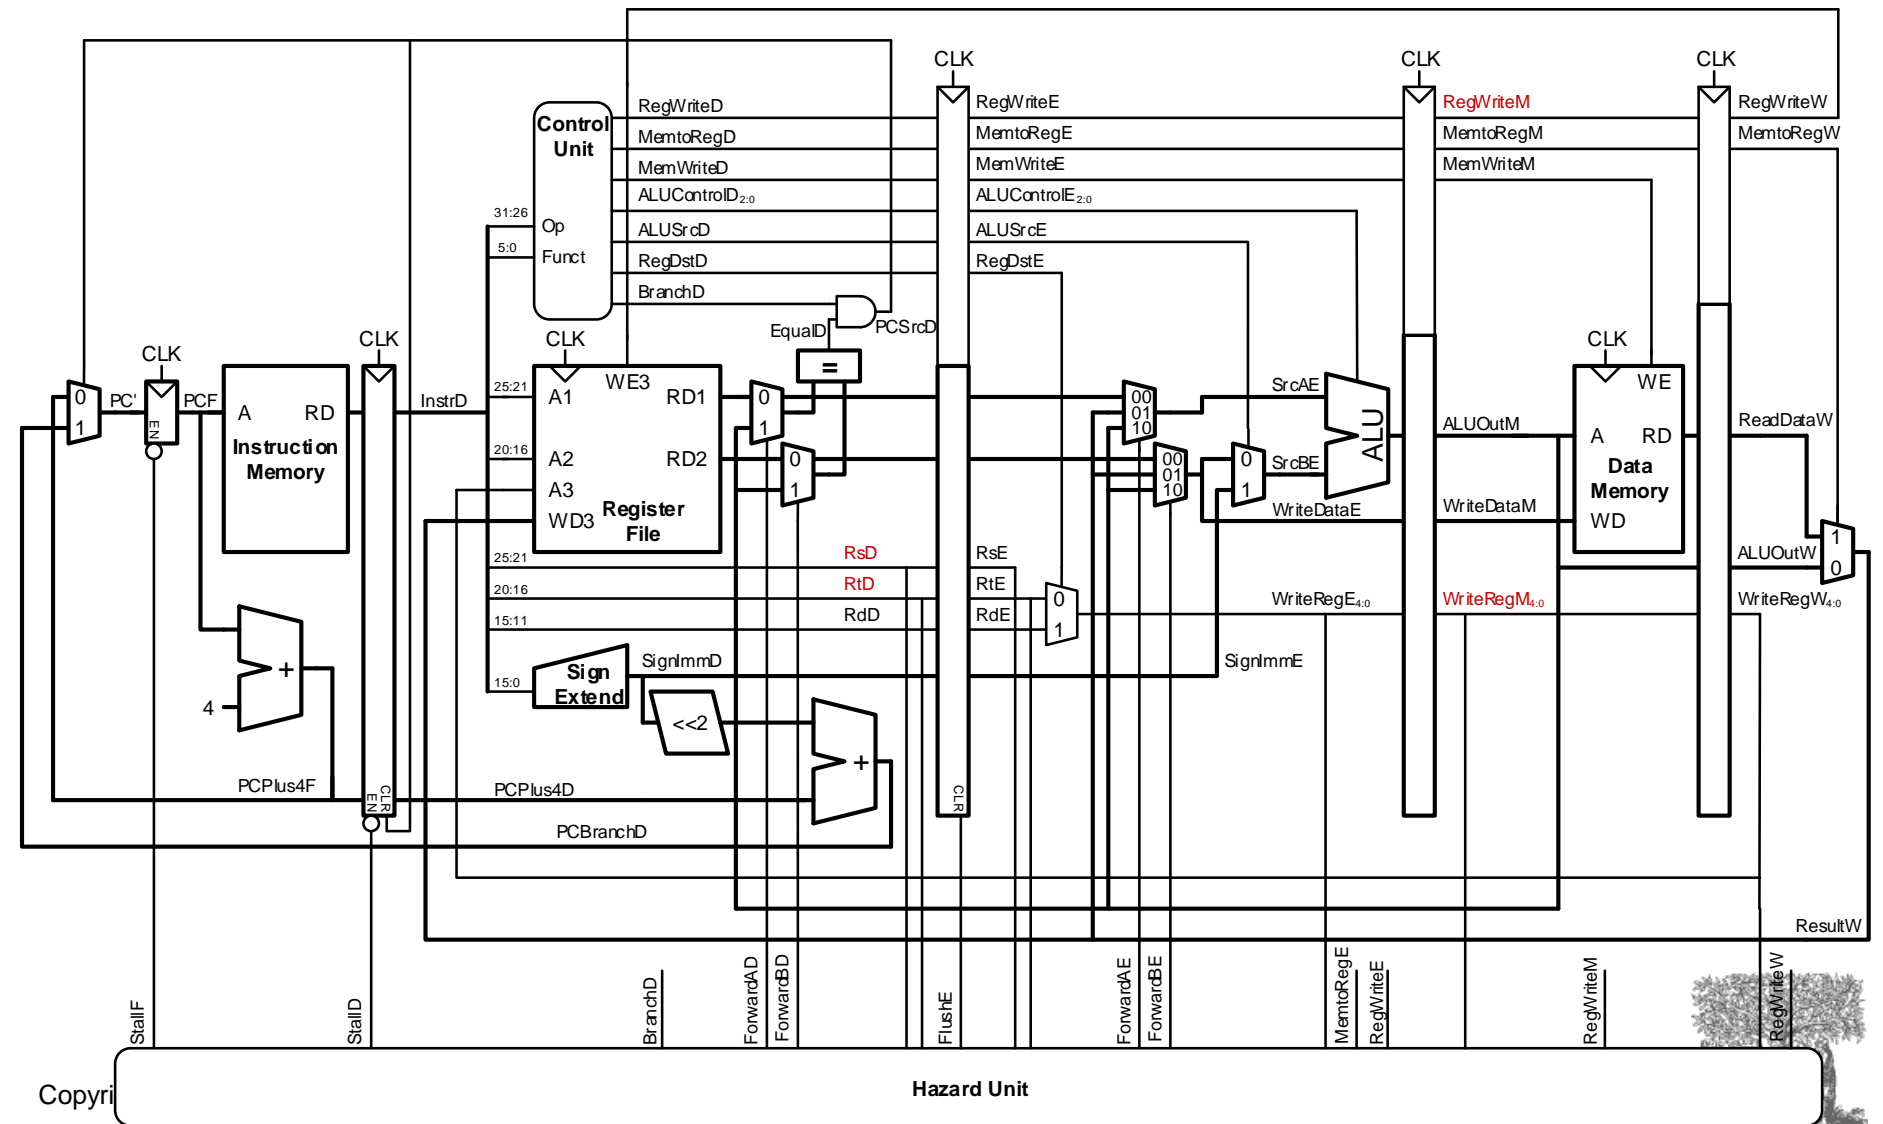


Control Hazards with Early Branch Resolution



$$ForwardAD = (rsD \neq 0) \text{ AND } (rsD == WriteRegM) \text{ AND } RegWriteM$$

$$ForwardBD = (rtD \neq 0) \text{ AND } (rtD == WriteRegM) \text{ AND } RegWriteM$$



Control Forwarding and Stalling Hardware

- Forwarding logic:

$ForwardAD = (rsD \neq 0) \text{ AND } (rsD == WriteRegM) \text{ AND } RegWriteM$

$ForwardBD = (rtD \neq 0) \text{ AND } (rtD == WriteRegM) \text{ AND } RegWriteM$

- Stalling logic:

$branchstall = BranchD \text{ AND } RegWriteE \text{ AND}$

$(WriteRegE == rsD \text{ OR } WriteRegE == rtD)$

OR

$BranchD \text{ AND } MemtoRegM \text{ AND}$

$(WriteRegM == rsD \text{ OR } WriteRegM == rtD)$

$StallF = StallD = FlushE = lwstall \text{ OR } branchstall$

Branch Prediction

- Guess whether branch will be taken
 - Backward branches are usually taken (loops)
 - Perhaps consider history of whether branch was previously taken to improve the guess
- Good prediction reduces the fraction of branches requiring a flush

Pipelined Performance Example

- Ideally $CPI = 1$
- But need to handle stalling (caused by loads and branches)
- SPECINT2000 benchmark:
 - 25% loads
 - 10% stores
 - 11% branches
 - 2% jumps
 - 52% R-type
- Suppose:
 - 40% of loads used by next instruction
 - 25% of branches mispredicted
- **What is the average CPI?**

Pipelined Performance Example

- SPECINT2000 benchmark:
 - 25% loads
 - 10% stores
 - 11% branches
 - 2% jumps
 - 52% R-type
- Suppose:
 - 40% of loads used by next instruction
 - 25% of branches mispredicted
 - All jumps flush next instruction
- **What is the average CPI?**
 - Load/Branch CPI = 1 when no stalling, 2 when stalling. Thus,
 - $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$
 - $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$
 - Thus,

$$\text{Average CPI} = (0.25)(1.4) + (0.1)(1) + (0.11)(1.25) + (0.02)(2) + (0.52)(1) \\ = 1.15$$

Pipelined Performance

- Pipelined processor critical path:

$$T_c = \max \{$$
$$t_{pcq} + t_{mem} + t_{setup}$$
$$2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup})$$
$$t_{pcq} + t_{mux} + t_{mux} + t_{ALU} + t_{setup}$$
$$t_{pcq} + t_{memwrite} + t_{setup}$$
$$2(t_{pcq} + t_{mux} + t_{RFwrite}) \}$$

Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	30
Register setup	t_{setup}	20
Multiplexer	t_{mux}	25
ALU	t_{ALU}	200
Memory read	t_{mem}	250
Register file read	t_{RFread}	150
Register file setup	$t_{RFsetup}$	20
Equality comparator	t_{eq}	40
AND gate	t_{AND}	15
Memory write	$T_{memwrite}$	220
Register file write	$t_{RFwrite}$	100 ps

$$\begin{aligned}
 T_c &= 2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup}) \\
 &= 2[150 + 25 + 40 + 15 + 25 + 20] \text{ ps} = \mathbf{550 \text{ ps}}
 \end{aligned}$$

Pipelined Performance Example

- For a program with 100 billion instructions executing on a pipelined MIPS processor,
- $CPI = 1.15$
- $T_c = 550$ ps

$$\begin{aligned}\text{Execution Time} &= (\# \text{ instructions}) \times CPI \times T_c \\ &= (100 \times 10^9)(1.15)(550 \times 10^{-12}) \\ &= 63 \text{ seconds}\end{aligned}$$

Processor	Execution Time (seconds)	Speedup (single-cycle is baseline)
Single-cycle	95	1
Multicycle	133	0.71
Pipelined	63	1.51

Review: Exceptions

- Unscheduled procedure call to the *exception handler*
- Casued by:
 - Hardware, also called an *interrupt*, e.g. keyboard
 - Software, also called *traps*, e.g. undefined instruction
- When exception occurs, the processor:
 - Records the cause of the exception (Cause register)
 - Jumps to the exception handler at instruction address 0x80000180
 - Returns to program (EPC register)

Example Exception

sequential circuits.

Can we design a spiff

Figure 2.11 shows a inputs, A and B, and on box indicates that it is this case, the function is

KeyAccess



The network KeyServer, which is required by KeyServer controlled programs, cannot grant you permission to run this program. If you think you have received this message in error, please contact your KeyServer Administrator.

Visio.exe - Application Error



The exception unknown software exception (0xc06d007e) occurred in the application at location 0x7c81eb33.

OK

words, we say the output Y is a function of the two inputs A and B where the function performed is A OR B.

The *implementation* of the combinational circuit is independent of its functionality. Figure 2.1 and Figure 2.2 show two possible implementa-

Exception Registers

- Not part of the register file.
 - Cause
 - Records the cause of the exception
 - Coprocessor 0 register 13
 - EPC (Exception PC)
 - Records the PC where the exception occurred
 - Coprocessor 0 register 14
- Move from Coprocessor 0
 - `mfc0 $t0, Cause`
 - Moves the contents of Cause into `$t0`

mfc0

010000	00000	\$t0 (8)	Cause (13)	000000000000
31:26	25:21	20:16	15:11	10:0

Exception Causes


Exception	Cause
Hardware Interrupt	0x00000000
System Call	0x00000020
Breakpoint / Divide by 0	0x00000024
Undefined Instruction	0x00000028
Arithmetic Overflow	0x00000030

We extend the multicycle MIPS processor to handle the last two types of exceptions.

© 2004 Blackwell Publishing Ltd
Journal of Internal Medicine 255: 103–110



Exception Hardware: mfc0



Control FSM with Exceptions

